

计算机组成与系统结构

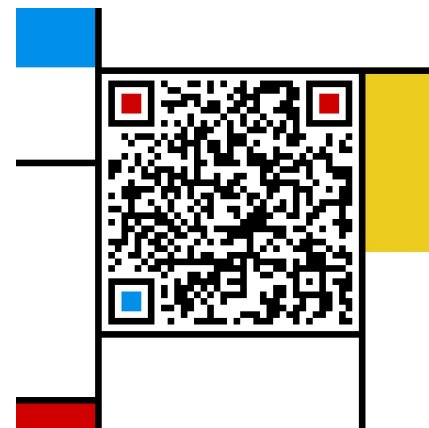
Computer Organization & System Architecture

Huang Kejie (黄科杰) 百人计划研究员

Office: 玉泉校区老生仪楼 304

Email address: huangkejie@zju.edu.cn

HP: 17706443800



Review

- We have designed a complete datapath
 - Capable of executing all RISC-V instructions in one cycle each
 - Not all units (hardware) used by all instructions
- 5 Phases of execution
 - IF, ID, EX, MEM, WB
 - Not all instructions are active in all phases
- Controller specifies how to execute instructions
 - New instructions can be added with just control?

Last Time in Lecture 5

- Microcoding, an effective technique to manage control unit complexity, invented in era when logic (tubes), main memory (magnetic core), and ROM (diodes) used different technologies
- Difference between ROM and RAM speed motivated additional complex instructions
- Technology advances leading to fast SRAM made technology assumptions invalid
- Complex instructions sets impede parallel and pipelined implementations
- Load/store, register-rich ISAs (pioneered by Cray, popularized by RISC) perform better in new VLSI technology

Analyzing Microcoded Machines

- John Cocke and group at IBM
 - Working on a simple pipelined processor, 801, and advanced compilers inside IBM
 - Ported experimental PL.8 compiler to IBM 370, and only used simple register-register and load/store instructions similar to 801
 - Code ran faster than other existing compilers that used all 370 instructions! (up to 6MIPS whereas 2MIPS considered good before)
- Emer, Clark, at DEC
 - Measured VAX-11/780 using external hardware
 - Found it was actually a 0.5MIPS machine, although usually assumed to be a 1MIPS machine
 - Found 20% of VAX instructions responsible for 60% of microcode, but only account for 0.2% of execution time!
- VAX8800
 - Control Store: 16K*147b RAM, Unified Cache: 64K*8b RAM
 - 4.5x more microstore RAM than cache RAM!

Reconsidering Microcode Machine (Nane 68000 example)

RISC!


Exploits recurring control
signal patterns in μ code,
e.g.,

ALU0 $A \leftarrow \text{Reg}[\text{rs1}]$

...

ALUI0 $A \leftarrow \text{Reg}[\text{rs1}]$

...



- Motorola 68000 had 17-bit μ code containing either 10-bit μ jump or 9-bit nanoinstruction pointer
 - Nanoinstructions were 68 bits wide, decoded to give 196 control signals

Microprogramming is far from extinct

- Played a crucial role in micros of the Eighties
 - DEC uVAX, Motorola 68K series, Intel 286/386
- Plays an assisting role in most modern micros
 - e.g., AMD Zen, Intel Sky Lake, Intel Atom, IBM PowerPC, ...
 - Most instructions executed directly, i.e., with hard-wired control
 - Infrequently-used and/or complicated instructions invoke microcode
- Patchable microcode common for post-fabrication bug fixes, e.g. Intel processors load μ code patches at bootup
 - Intel had to scramble to resurrect microcode tools and find original microcode engineers to patch Meltdown/Spectre security vulnerabilities


Recap: Complete RV32I ISA

| | | | | | | |
|---------------------|-----|-----|-----|-------------|---------|-------|
| imm[31:12] | | | | rd | 0110111 | LUI |
| imm[31:12] | | | | rd | 0010111 | AUIPC |
| imm[20:10:11:19:12] | | | | rd | 1101111 | JAL |
| imm[11:0] | | rs1 | 000 | rd | 1100111 | JALR |
| imm[12:10:5] | rs2 | rs1 | 000 | imm[4:1:11] | 1100011 | BEQ |
| imm[12:10:5] | rs2 | rs1 | 001 | imm[4:1:11] | 1100011 | BNE |
| imm[12:10:5] | rs2 | rs1 | 100 | imm[4:1:11] | 1100011 | BLT |
| imm[12:10:5] | rs2 | rs1 | 101 | imm[4:1:11] | 1100011 | BGE |
| imm[12:10:5] | rs2 | rs1 | 110 | imm[4:1:11] | 1100011 | BLTU |
| imm[12:10:5] | rs2 | rs1 | 111 | imm[4:1:11] | 1100011 | BGEU |
| imm[11:0] | | rs1 | 000 | rd | 0000011 | LB |
| imm[11:0] | | rs1 | 001 | rd | 0000011 | LH |
| imm[11:0] | | rs1 | 010 | rd | 0000011 | LW |
| imm[11:0] | | rs1 | 100 | rd | 0000011 | LBU |
| imm[11:0] | | rs1 | 101 | rd | 0000011 | LHU |
| imm[11:5] | rs2 | rs1 | 000 | imm[4:0] | 0100011 | SB |
| imm[11:5] | rs2 | rs1 | 001 | imm[4:0] | 0100011 | SH |
| imm[11:5] | rs2 | rs1 | 010 | imm[4:0] | 0100011 | SW |
| imm[11:0] | | rs1 | 000 | rd | 0010011 | ADDI |
| imm[11:0] | | rs1 | 010 | rd | 0010011 | SLTI |
| imm[11:0] | | rs1 | 011 | rd | 0010011 | SLTIU |
| imm[11:0] | | rs1 | 100 | rd | 0010011 | XORI |
| imm[11:0] | | rs1 | 110 | rd | 0010011 | ORI |
| imm[11:0] | | rs1 | 111 | rd | 0010011 | ANDI |

| | | | | | | | |
|---------------|------|-------|-------|-----|-------|---------|---------|
| 0000000 | | shamt | rs1 | 001 | rd | 0010011 | SLLI |
| 0000000 | | shamt | rs1 | 101 | rd | 0010011 | SRLI |
| 0100000 | | shamt | rs1 | 101 | rd | 0010011 | SRAI |
| 0000000 | | rs2 | rs1 | 000 | rd | 0110011 | ADD |
| 0100000 | | rs2 | rs1 | 000 | rd | 0110011 | SUB |
| 0000000 | | rs2 | rs1 | 001 | rd | 0110011 | SLL |
| 0000000 | | rs2 | rs1 | 010 | rd | 0110011 | SLT |
| 0000000 | | rs2 | rs1 | 011 | rd | 0110011 | SLTU |
| 0000000 | | rs2 | rs1 | 100 | rd | 0110011 | XOR |
| 0000000 | | rs2 | rs1 | 101 | rd | 0110011 | SRL |
| 0100000 | | rs2 | rs1 | 101 | rd | 0110011 | SRA |
| 0000000 | | rs2 | rs1 | 110 | rd | 0110011 | OR |
| 0000000 | | rs2 | rs1 | 111 | rd | 0110011 | AND |
| 0000 | pred | succ | 00000 | 000 | 00000 | 0001111 | FENCE |
| 0000 | 0000 | 0000 | 00000 | 001 | 00000 | 0001111 | FENCE.I |
| 0000000000000 | | | 00000 | 000 | 00000 | 1110011 | ECALL |
| 0000000000001 | | | 00000 | 000 | 00000 | 1110011 | EBREAK |
| csr | | | rs1 | 001 | rd | 1110011 | CSRRW |
| csr | | | rs1 | 010 | rd | 1110011 | CSRRS |
| csr | | | rs1 | 011 | rd | 1110011 | CSRRC |
| csr | | | zimm | 101 | rd | 1110011 | CSRRWI |
| csr | | | zimm | 110 | rd | 1110011 | CSRRSI |
| csr | | | zimm | 111 | rd | 1110011 | CSRRCI |

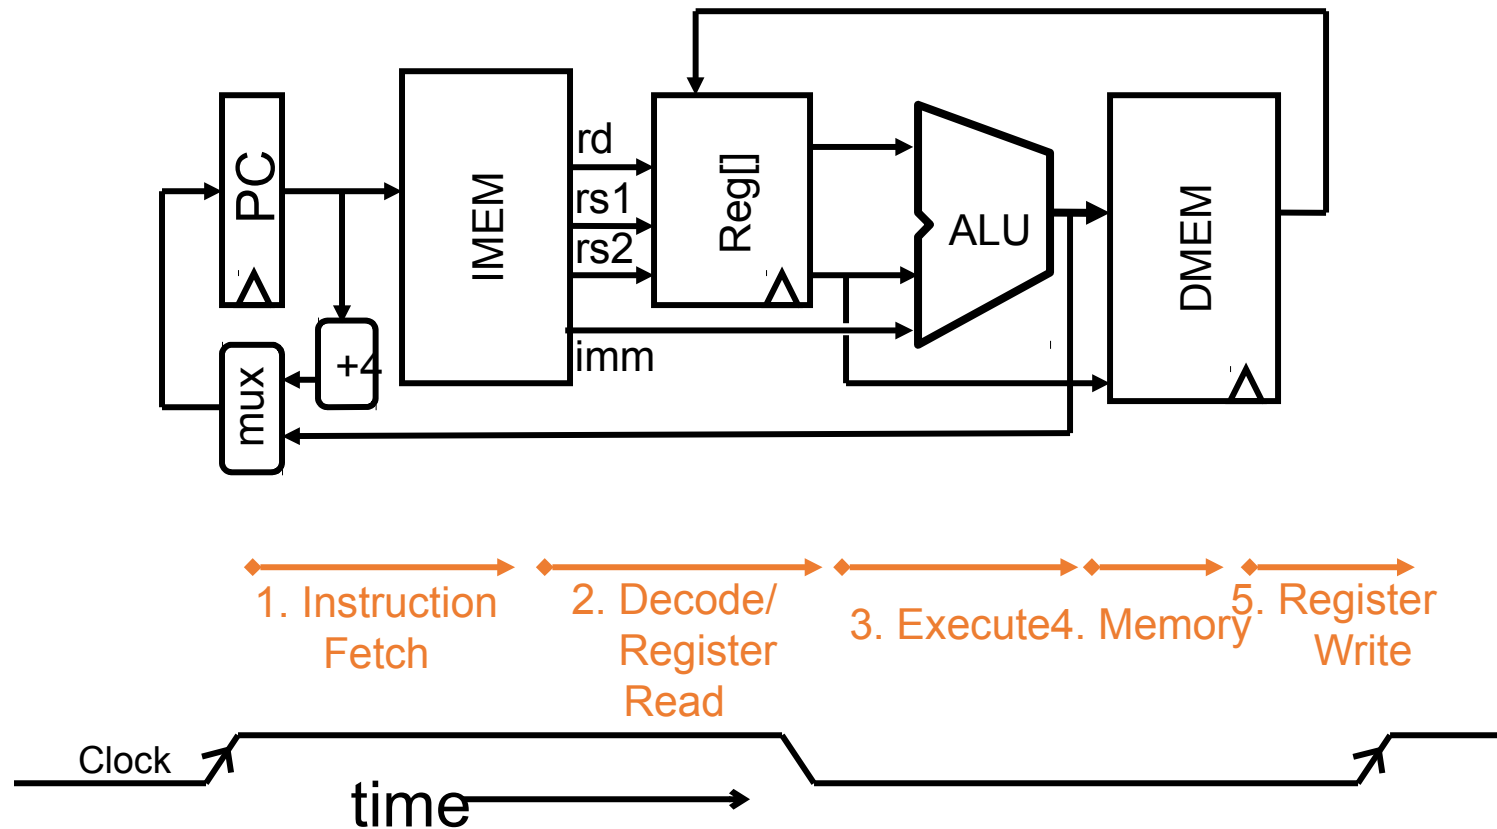
- RV32I has 47 instructions
- 37 instructions are enough to run any C program

One-Instruction-Per-Cycle RISC-V Machine

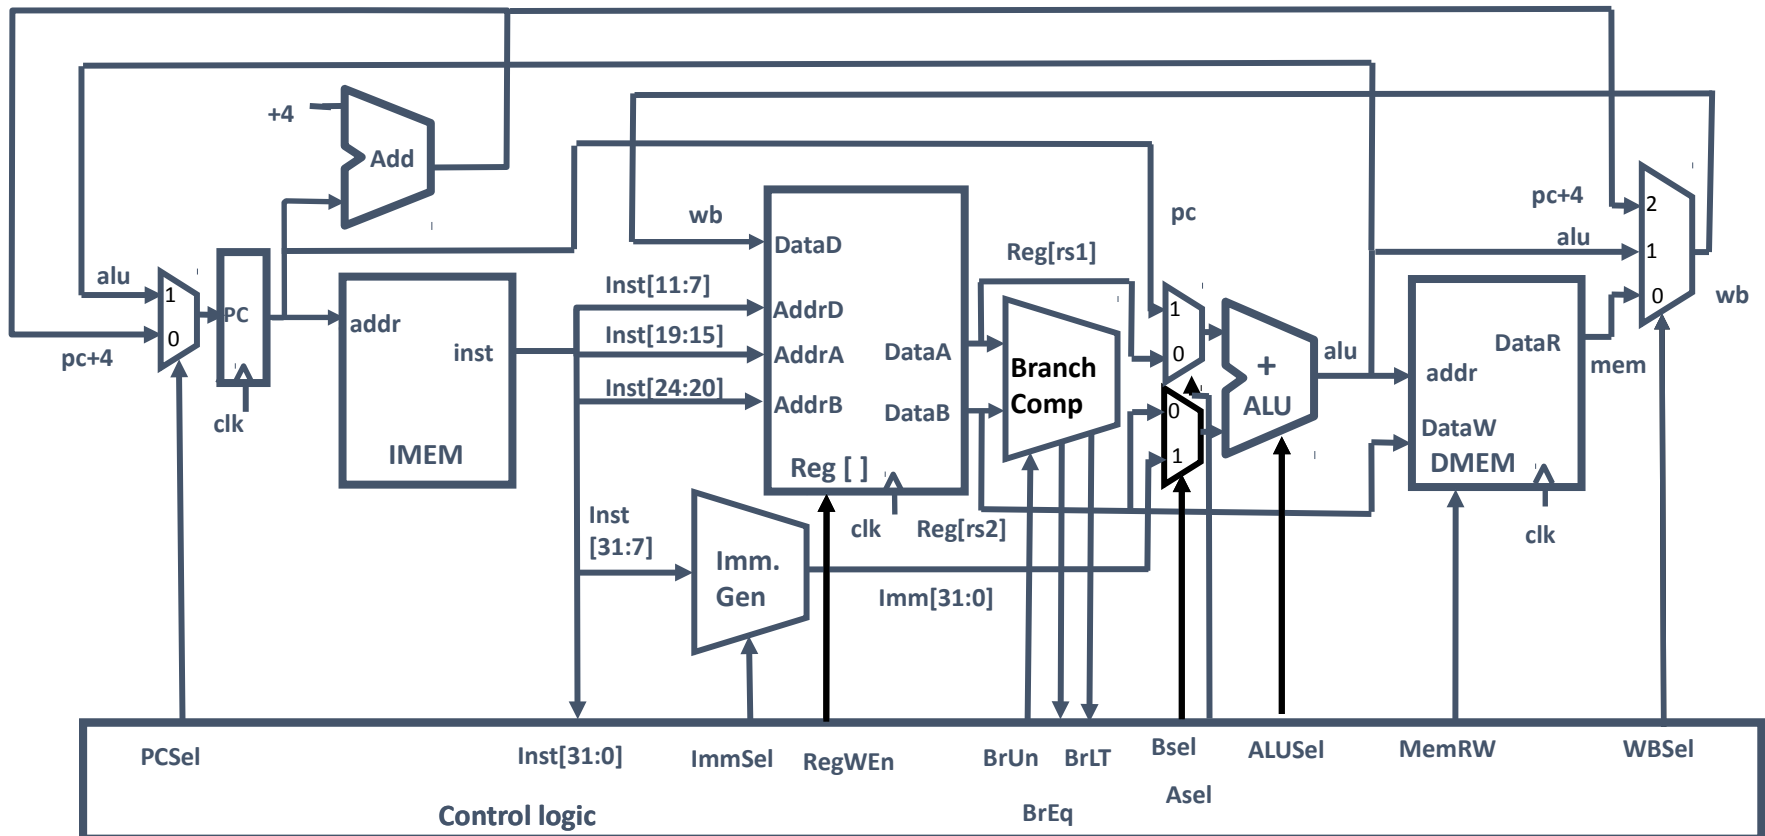


- On every tick of the clock, the computer executes one instruction
- Current state outputs drive the inputs to the combinational logic, whose outputs settles at the values of the state before the next clock edge
- At the rising clock edge, all the state elements are updated with the combinational logic outputs, and execution moves to the next clock cycle

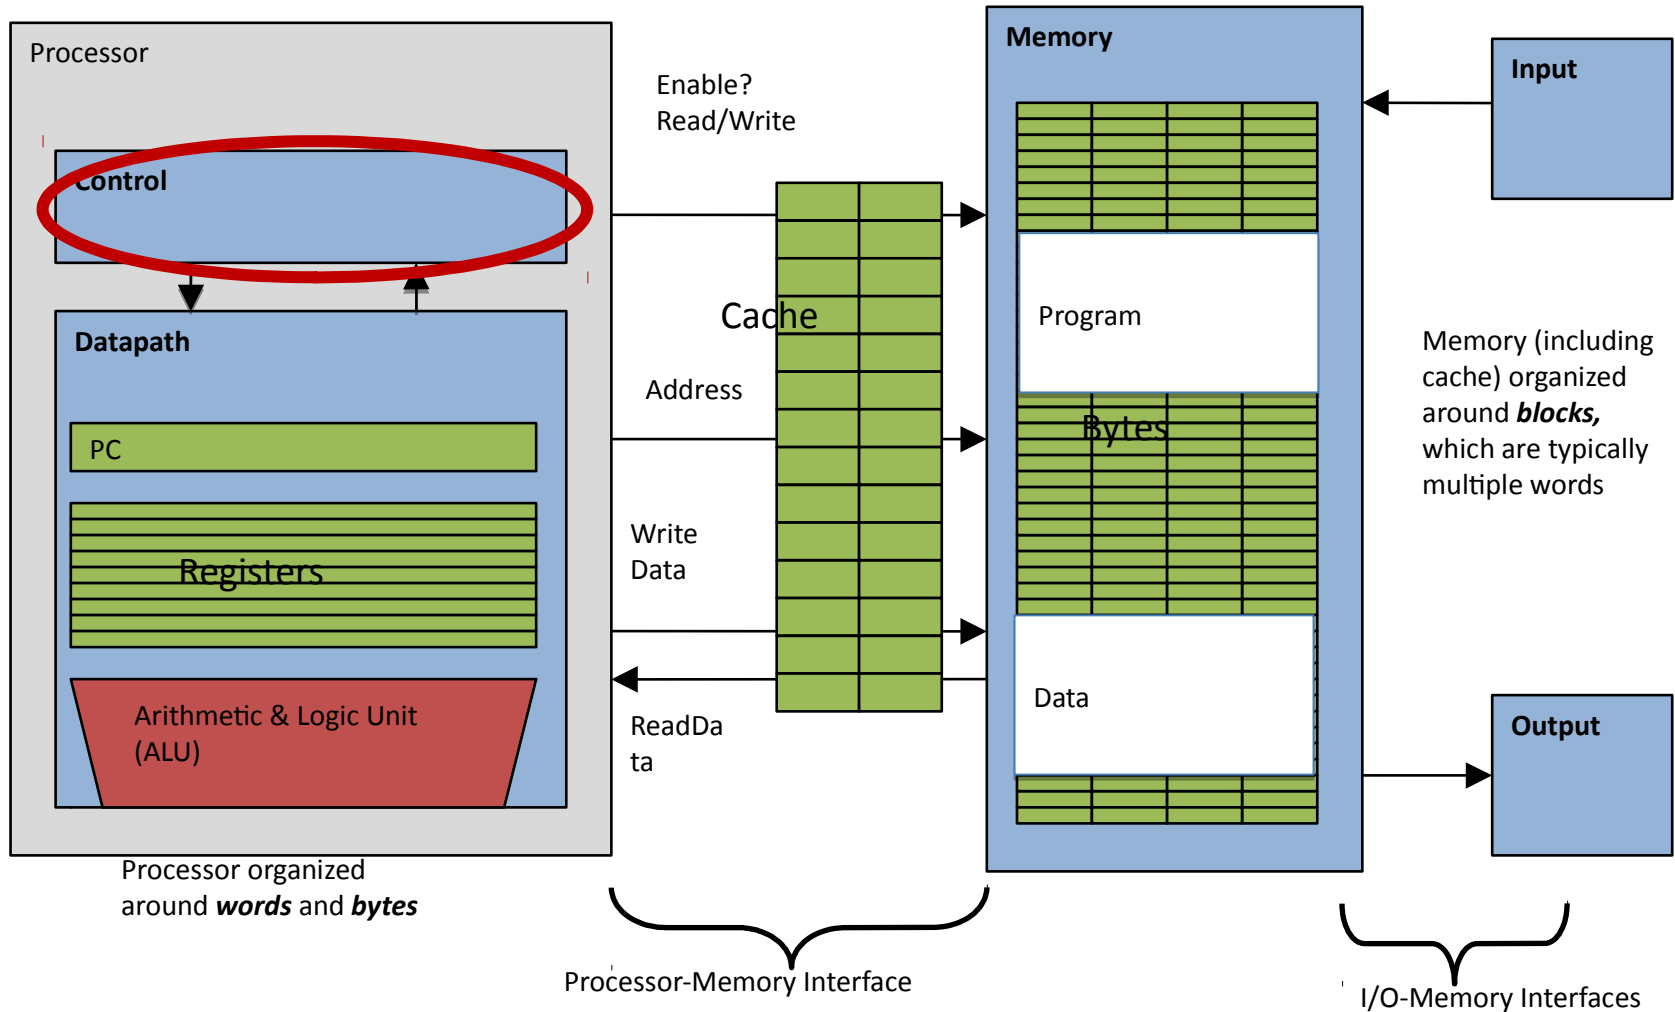
Basic Phases of Instruction Execution



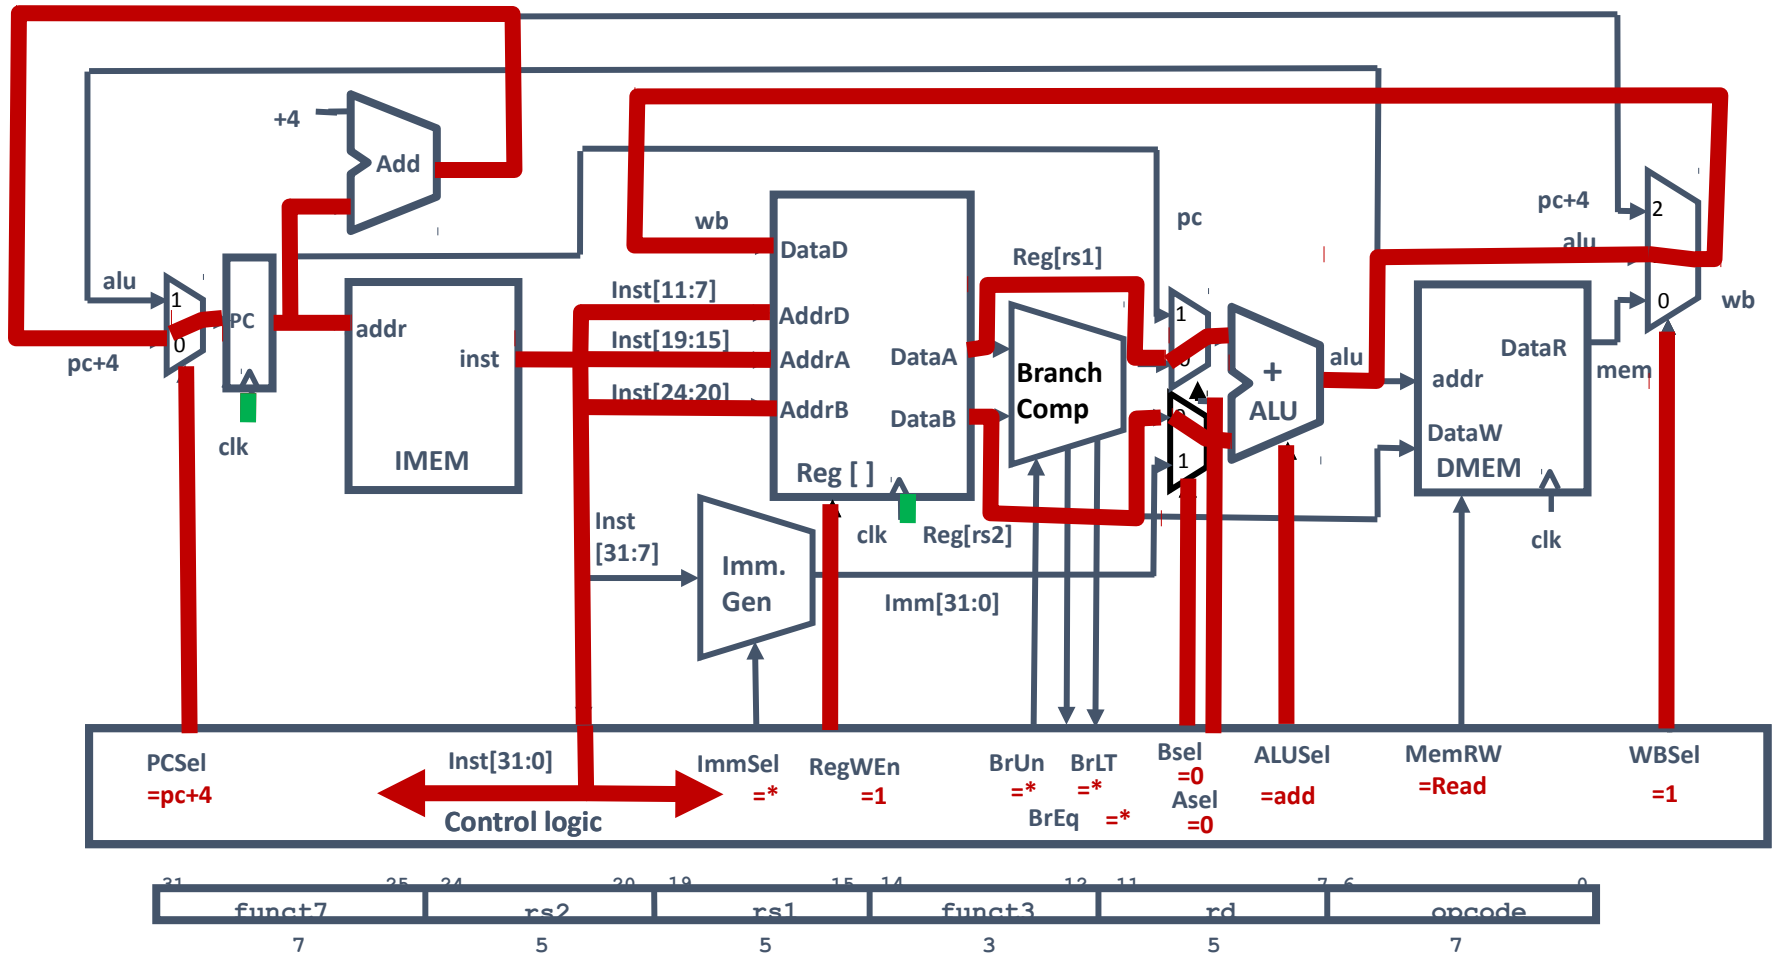
Complete Single-Cycle RV32I Datapath!



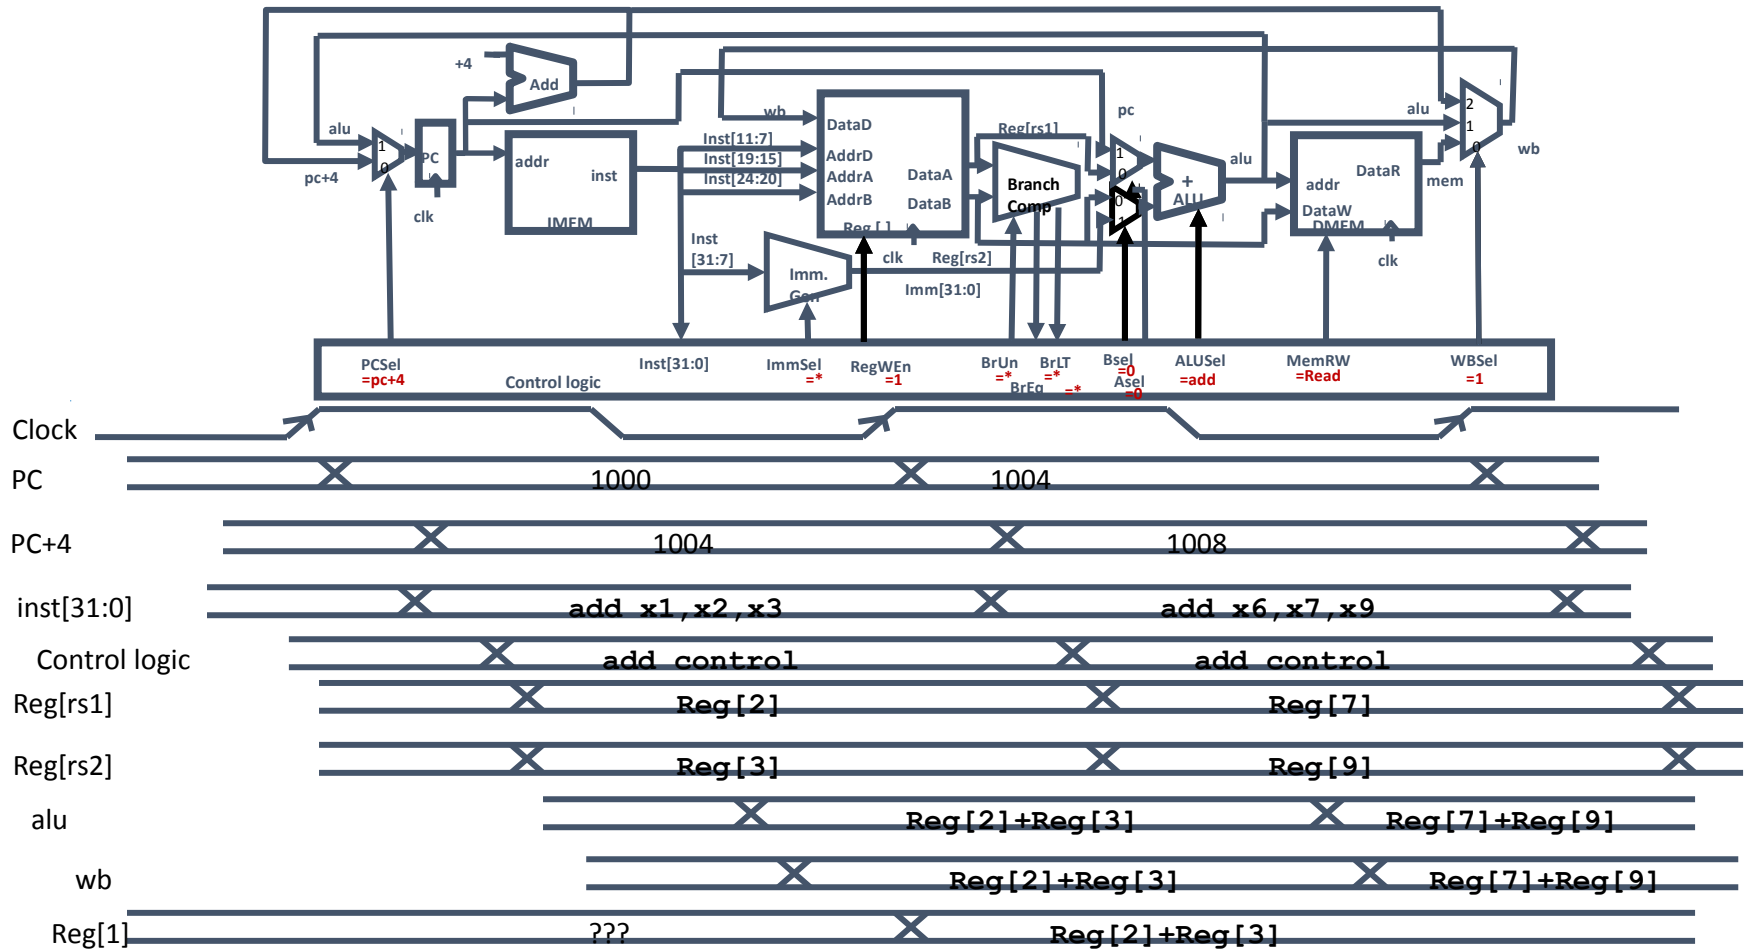
Our Single-Core Computer (w/o FPU, SIMD)



Example: add



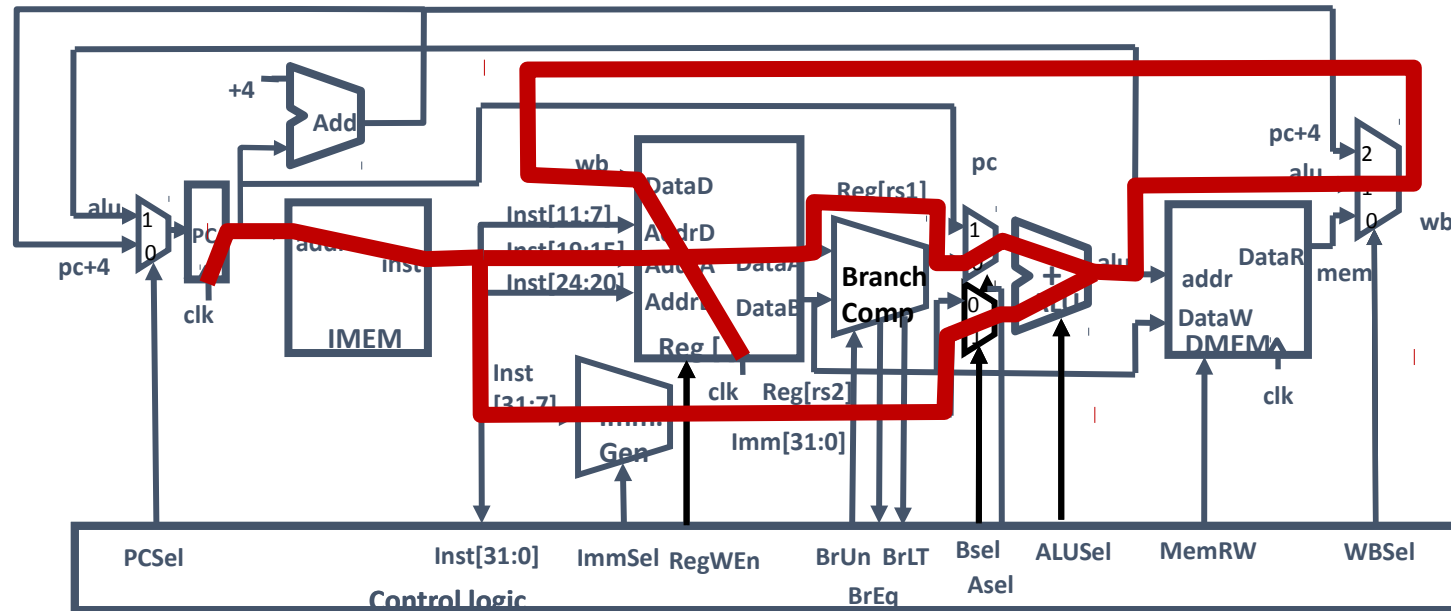
add Execution



Peer Instruction(s): Critical Path

Critical path for an addi

$R[rd] = R[rs1] + imm$



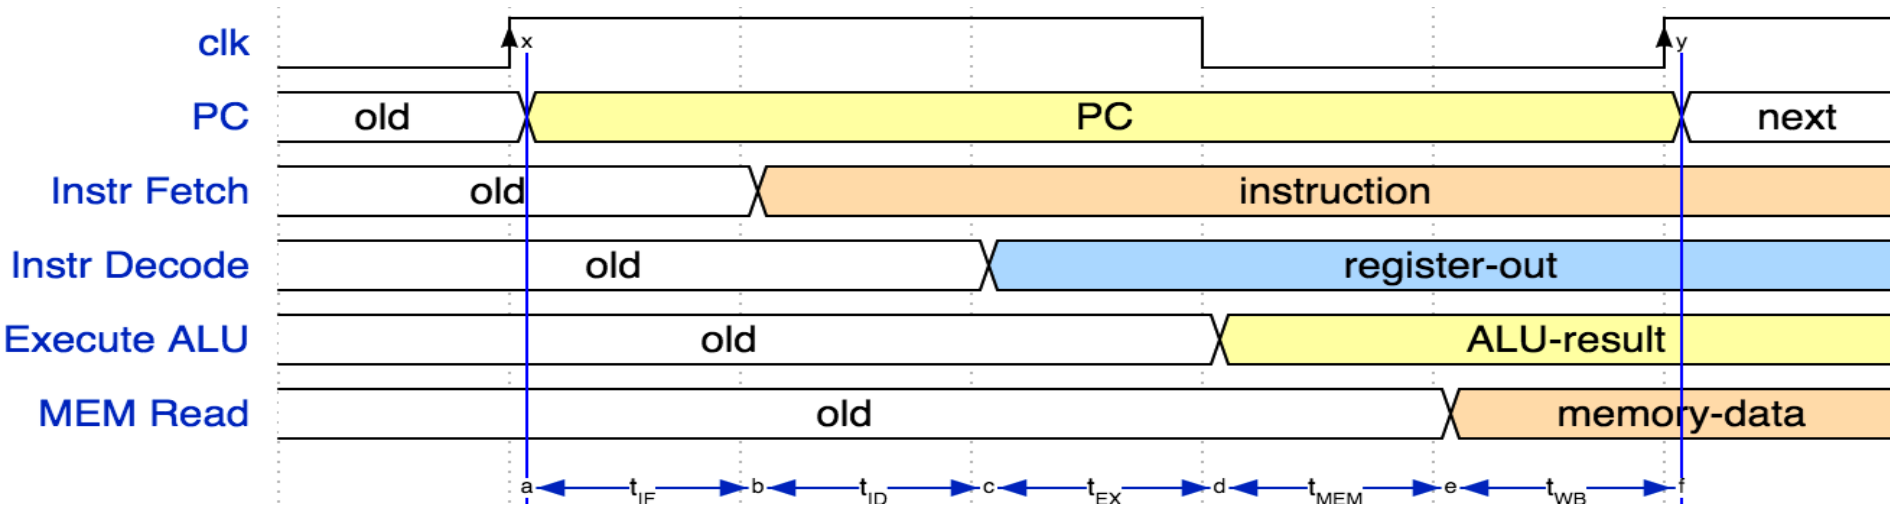
$$1) t_{\text{clk-q}} + t_{\text{Add}} + t_{\text{IMEM}} + t_{\text{Reg}} + t_{\text{BComp}} + t_{\text{ALU}} + t_{\text{DMEM}} + t_{\text{mux}} + t_{\text{Setup}}$$

$$2) t_{\text{clk-q}} + t_{\text{IMEM}} + \max\{t_{\text{Reg}}, t_{\text{Imm}}\} + t_{\text{ALU}} + 2t_{\text{mux}} + t_{\text{Setup}}$$

$$3) t_{\text{clk-q}} + t_{\text{IMEM}} + \max\{t_{\text{Reg}}, t_{\text{Imm}}\} + t_{\text{ALU}} + 3t_{\text{mux}} + t_{\text{DMEM}} + t_{\text{Setup}}$$

4) None of the above

Instruction Timing



| IF | ID | EX | MEM | WB | Total |
|--------|----------|--------|--------|--------|--------|
| I-MEM | Reg Read | ALU | D-MEM | Reg W | |
| 200 ps | 100 ps | 200 ps | 200 ps | 100 ps | 800 ps |

Instruction Timing

| Instr | IF = 200ps | ID = 100ps | ALU = 200ps | MEM=200ps | WB = 100ps | Total |
|-------|------------|------------|-------------|-----------|------------|-------|
| add | X | X | X | | X | 600ps |
| beq | X | X | X | | | 500ps |
| jal | X | X | X | | | 500ps |
| lw | X | X | X | X | X | 800ps |
| sw | X | X | X | X | | 700ps |

- Maximum clock frequency
 - $f_{\max} = 1/800\text{ps} = 1.25 \text{ GHz}$
- Most blocks idle most of the time
 - E.g. $f_{\max, \text{ALU}} = 1/200\text{ps} = 5 \text{ GHz!}$
 - How can we keep ALU busy all the time?
 - 5 billion adds/sec, rather than just 1.25 billion?
 - Idea: Factories use three employee shifts - equipment is always busy!

Performance Measures

- “Our” RISC-V executes instructions at 1.25 GHz
 - 1 instruction every 800 ps
- Can we improve its performance?
 - What do we mean with this statement?
 - Not so obvious:
 - Quicker response time, so one job finishes faster?
 - More jobs per unit time (e.g. web server returning pages)?
 - Longer battery life?

Transportation Analogy



| | Sports Car | Bus |
|--------------------|------------|--------|
| Passenger Capacity | 2 | 50 |
| Travel Speed | 200 mph | 50 mph |
| Gas Mileage | 5 mpg | 2 mpg |

50 Mile trip:

| | Sports Car | Bus |
|-------------------------|------------|-------------|
| Travel Time | 15 min | 60 min |
| Time for 100 passengers | 750 min | 120 min |
| Gallons per passenger | 5 gallons | 0.5 gallons |

Computer Analogy

| Transportation | Computer |
|-------------------------|--|
| Trip Time | Program execution time: e.g. time to update display |
| Time for 100 passengers | Throughput: e.g. number of server requests handled per hour |
| Gallons per passenger | Energy per task*: e.g. how many movies you can watch per battery charge or energy bill for datacenter |

* Note: power is not a good measure, since low-power CPU might run for a long time to complete one task consuming more energy than faster computer running at higher power for a shorter time

“Iron Law” of Processor Performance

$$\frac{\text{Time}}{\text{Program}} = \frac{\text{Instructions}}{\text{Program}} * \frac{\text{Cycles}}{\text{Instruction}} * \frac{\text{Time}}{\text{Cycle}}$$

Instructions per Program

- Determined by
 - Task
 - Algorithm, e.g. $O(N^2)$ vs $O(N)$
 - Programming language
 - Compiler
 - Instruction Set Architecture (ISA)

(Average) Clock cycles per Instruction

- Determined by
 - ISA
 - Processor implementation (or *microarchitecture*)
 - E.g. for “our” single-cycle RISC-V design, CPI = 1
 - Complex instructions (e.g. **strcpy**), CPI $\gg 1$
 - Superscalar processors, CPI < 1 (next lecture)

Time per Cycle ($1/\text{Frequency}$)

- Determined by
 - Processor microarchitecture (determines critical path through logic gates)
 - Technology (e.g. 14nm versus 28nm)
 - Power budget (lower voltages reduce transistor speed)

Speed Tradeoff Example

- For some task (e.g. image compression) ...

| | Processor A | Processor B |
|----------------|-------------|-------------|
| # Instructions | 1 Million | 1.5 Million |
| Average CPI | 2.5 | 1 |
| Clock rate f | 2.5 GHz | 2 GHz |
| Execution time | 1 ms | 0.75 ms |

Processor B is faster for this task, despite executing more instructions and having a lower clock rate!

Energy per Task

$$\frac{\text{Energy}}{\text{Program}} = \frac{\text{Instructions}}{\text{Program}} \frac{\text{Energy}}{\text{Instruction}}$$
$$\frac{\text{Energy}}{\text{Program}} \propto \frac{\text{Instructions}}{\text{Program}} * C V^2$$

“Capacitance” depends on technology,
processor features
e.g. # of cores

Supply voltage,
e.g. 1V

Want to reduce capacitance and voltage to reduce energy/task

Energy Tradeoff Example

- “Next-generation” processor
 - C (Moore’s Law): -15 %
 - Supply voltage, V_{sup} : -15 %
 - Energy consumption: $1 - (1-0.85)^3 = -39 \%$
- Significantly improved energy efficiency thanks to
 - Moore’s Law **AND**
 - Reduced supply voltage

Energy “Iron Law”

$$\text{Performance} = \frac{\text{Power}}{\text{Energy Efficiency}}$$


(Tasks/Second) (Joules/Second) (Tasks/Joule)

- Energy efficiency (e.g., instructions/Joule) is key metric in all computing devices
- For power-constrained systems (e.g., 20MW datacenter), need better energy efficiency to get more performance at same power
- For energy-constrained systems (e.g., 1W phone), need better energy efficiency to prolong battery life

End of Scaling

- In recent years, industry has not been able to reduce supply voltage much, as reducing it further would mean increasing “leakage power” where transistor switches don’t fully turn off (more like dimmer switch than on-off switch)
- Also, size of transistors and hence capacitance, not shrinking as much as before between transistor generations
- Power becomes a growing concern – the “power wall”
- Cost-effective air-cooled chip limit around ~150W

End of Scaling



Pipelining

- A familiar example:
 - Getting a university degree




Year 1



Year 2



Year 3

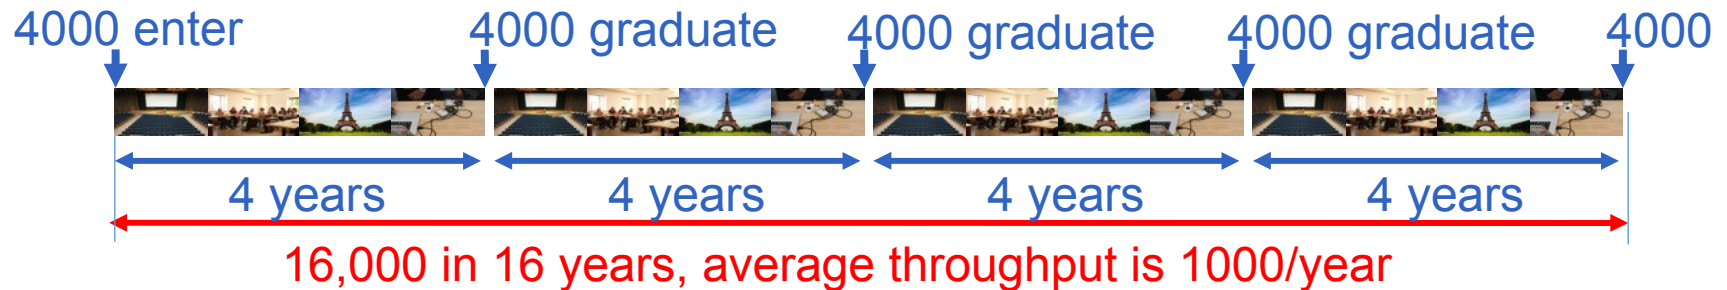


Year 4

- Shortage of Computer scientists (your startup is growing):
 - How long does it take to educate 16,000?

Computer Scientist Education

- Option 1: *serial*



- Option 2: *pipelining*




- 16,000 in 7 years
- Steady state throughput is 4000/year
- Resources used efficiently
- **4-fold improvement over serial education**

Latency versus Throughput

- Latency
 - Time from entering college to graduation
 - Serial 4 years
 - Pipelining 4 years
- Throughput
 - Average number of students graduating each year
 - Serial 1000
 - Pipelining 4000
- Pipelining
 - Increases throughput (4x in this example)
 - But does nothing to latency
 - sometimes worse (additional overhead e.g. for shift transition)

Simultaneous versus Sequential

- What happens *sequentially*?
- What happens *simultaneously*?



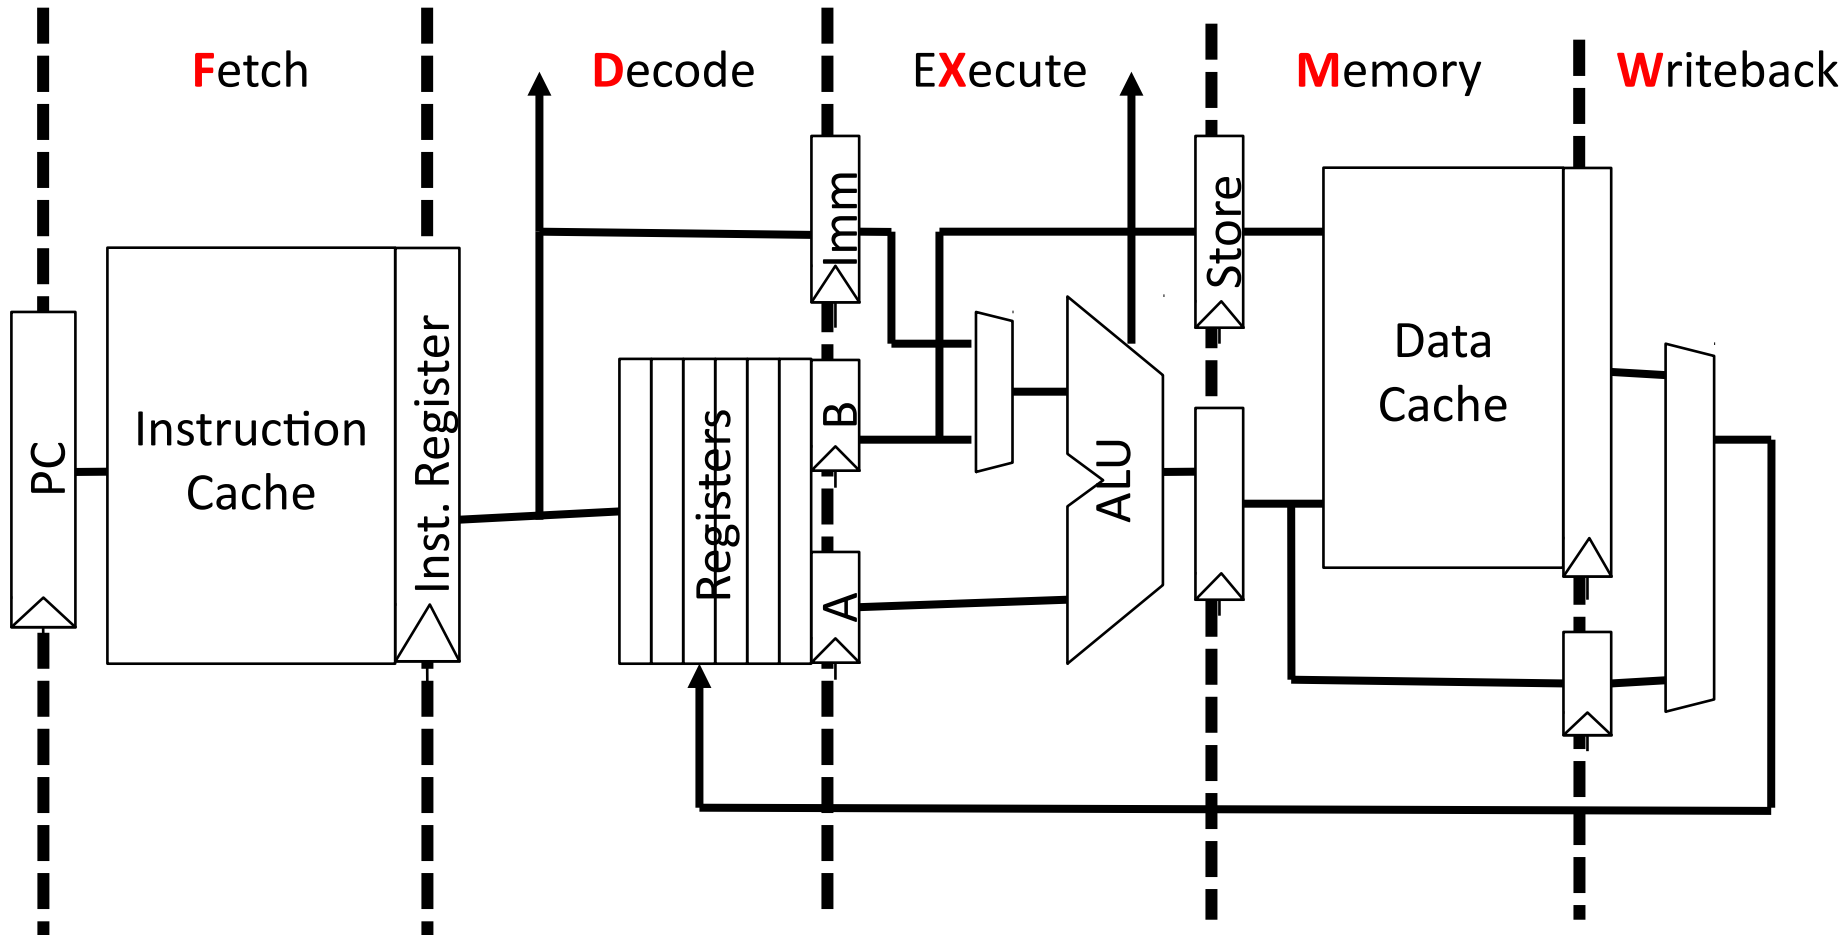
“Iron Law” of Processor Performance

$$\frac{\text{Time}}{\text{Program}} = \frac{\text{Instructions}}{\text{Program}} * \frac{\text{Cycles}}{\text{Instruction}} * \frac{\text{Time}}{\text{Cycle}}$$

- Instructions per program depends on source code, compiler technology, and ISA
- Cycles per instructions (CPI) depends on ISA and μ architecture
- Time per cycle depends upon the μ architecture and base technology

| Microarchitecture | CPI | cycle time |
|--------------------------|-----|------------|
| Microcoded | >1 | short |
| Single-cycle unpipelined | 1 | long |
| Pipelined | 1 | short |


Classic 5-Stage RISC Pipeline



*This version designed for regfiles/memories
with synchronous reads and writes.*

CPI Examples

Microcoded machine




3 instructions, 22 cycles, $CPI=7.33$

Unpipelined machine



3 instructions, 3 cycles, $CPI=1$







Pipelined machine

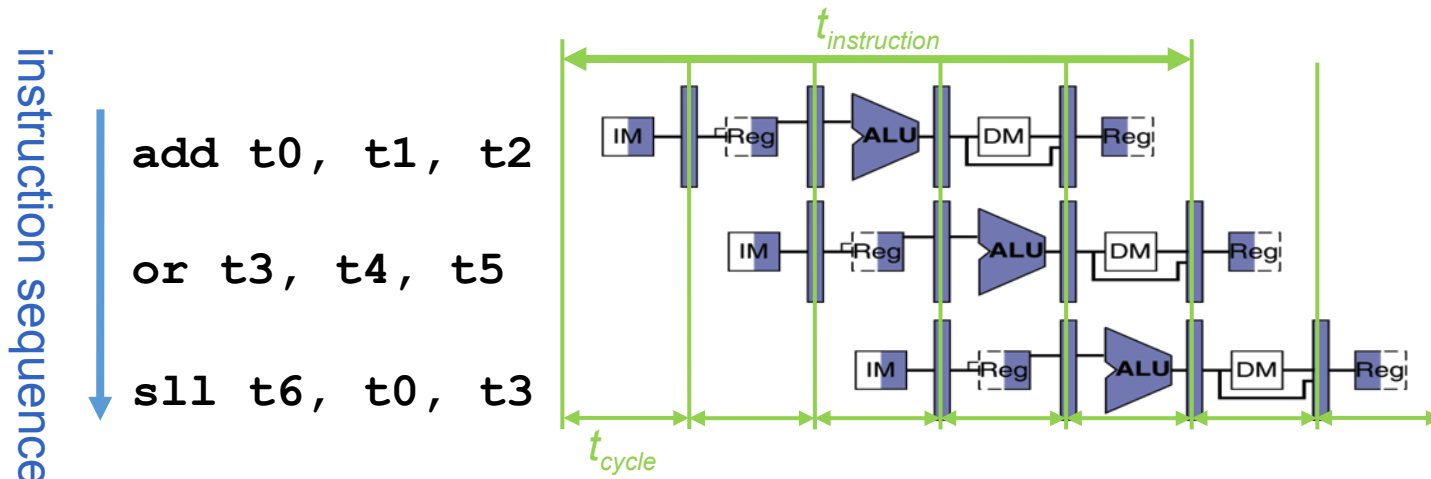


3 instructions, 3 cycles, $CPI=1$

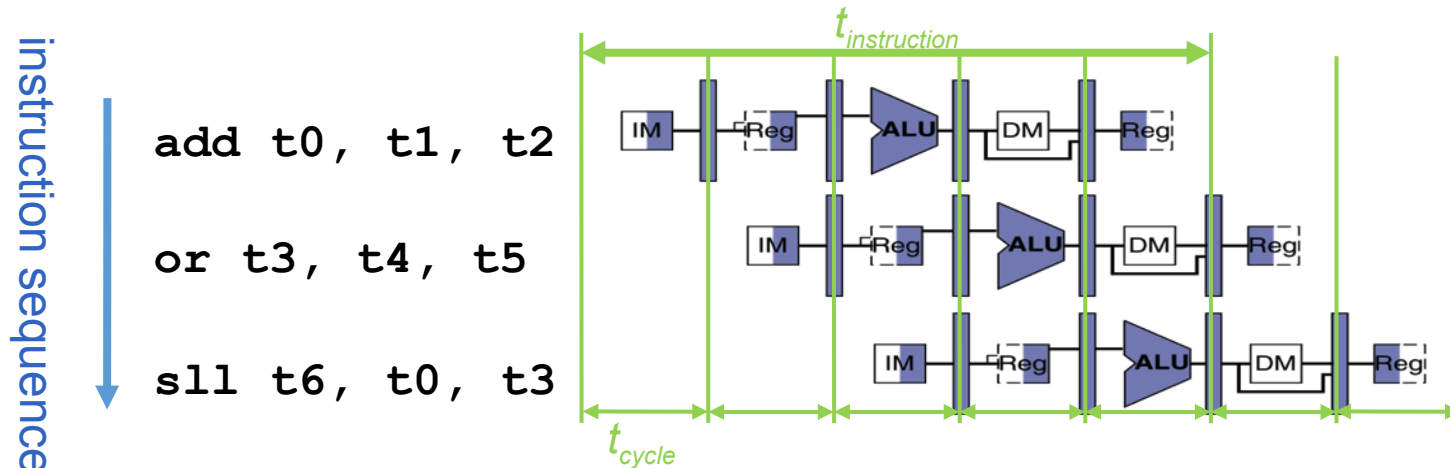
5-stage pipeline $CPI \neq 5!!!$

Pipelining with RISC-V

| Phase | Pictogram | t_{step} Serial | t_{cycle} Pipelined |
|-------------------|--|-------------------|-----------------------|
| Instruction Fetch |  | 200 ps | 200 ps |
| Reg Read |  | 100 ps | 200 ps |
| ALU |  | 200 ps | 200 ps |
| Memory |  | 200 ps | 200 ps |
| Register Write |  | 100 ps | 200 ps |
| $t_{instruction}$ |  | 800 ps | 1000 ps |



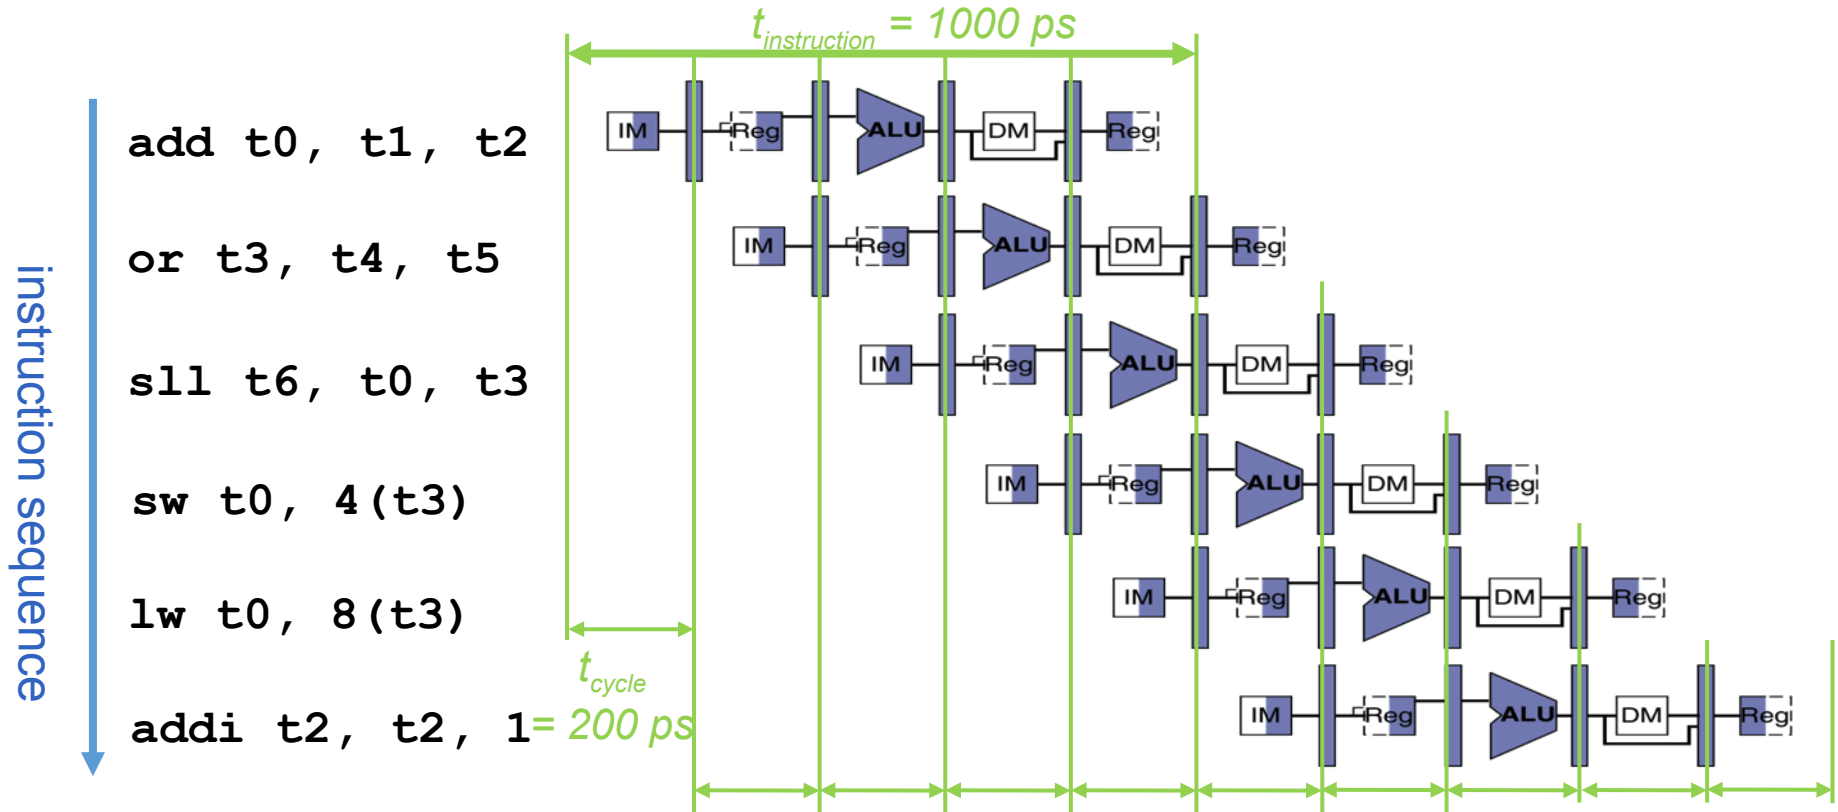
Pipelining with RISC-V



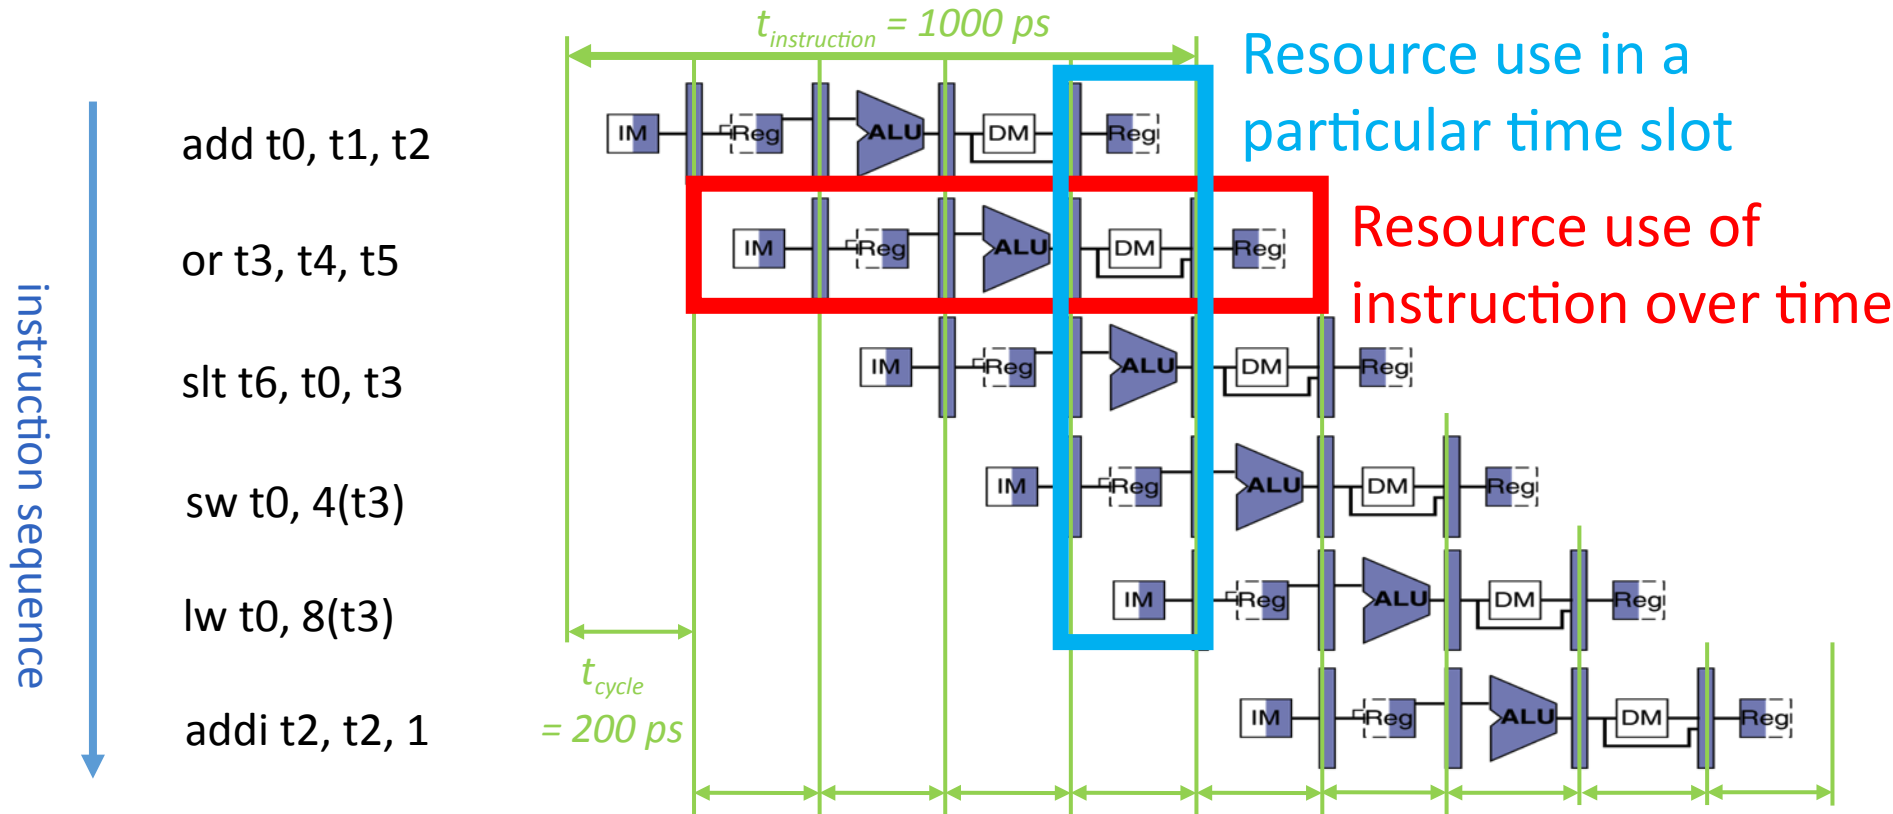
| | Single Cycle | Pipelining |
|-------------------------------------|---------------------------------------|------------------------------------|
| Timing | $t_{step} = 100 \dots 200 \text{ ps}$ | $t_{cycle} = 200 \text{ ps}$ |
| | Register access only 100 ps | All cycles same length |
| Instruction time, $t_{instruction}$ | $= t_{cycle} = 800 \text{ ps}$ | 1000 ps |
| Clock rate, f_s | $1/800 \text{ ps} = 1.25 \text{ GHz}$ | $1/200 \text{ ps} = 5 \text{ GHz}$ |
| Relative speed | 1 x | 4 x |

Sequential vs Simultaneous

What happens sequentially, what happens simultaneously?



RISC-V Pipeline



And in Conclusion, ...

- Controller
 - Tells universal datapath how to execute each instruction
- Instruction timing
 - Set by instruction complexity, architecture, technology
 - Pipelining increases clock frequency, “instructions per second”
 - But does not reduce time to complete instruction
- Performance measures
 - Different measures depending on objective
 - Response time
 - Jobs / second
 - Energy per task