

计算机组成与系统结构

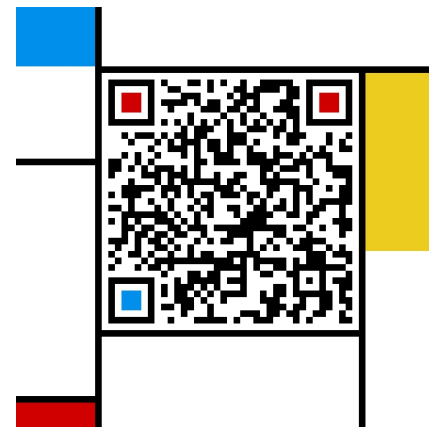
Computer Organization & System Architecture

Huang Kejie (黄科杰) 百人计划研究员

Office: 玉泉校区老生仪楼 304

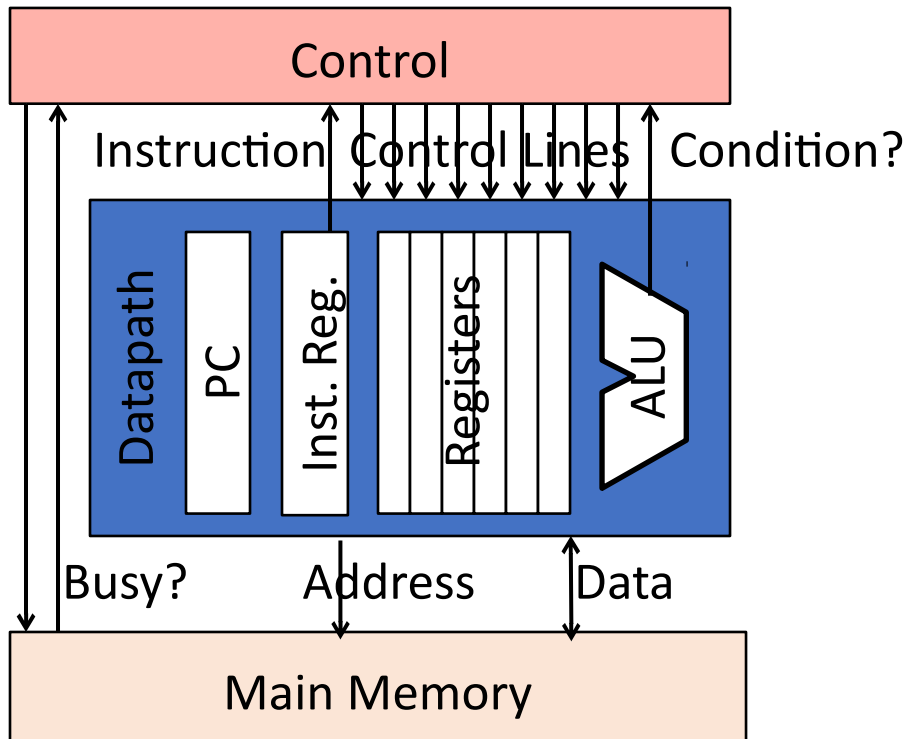
Email address: huangkejie@zju.edu.cn

HP: 17706443800



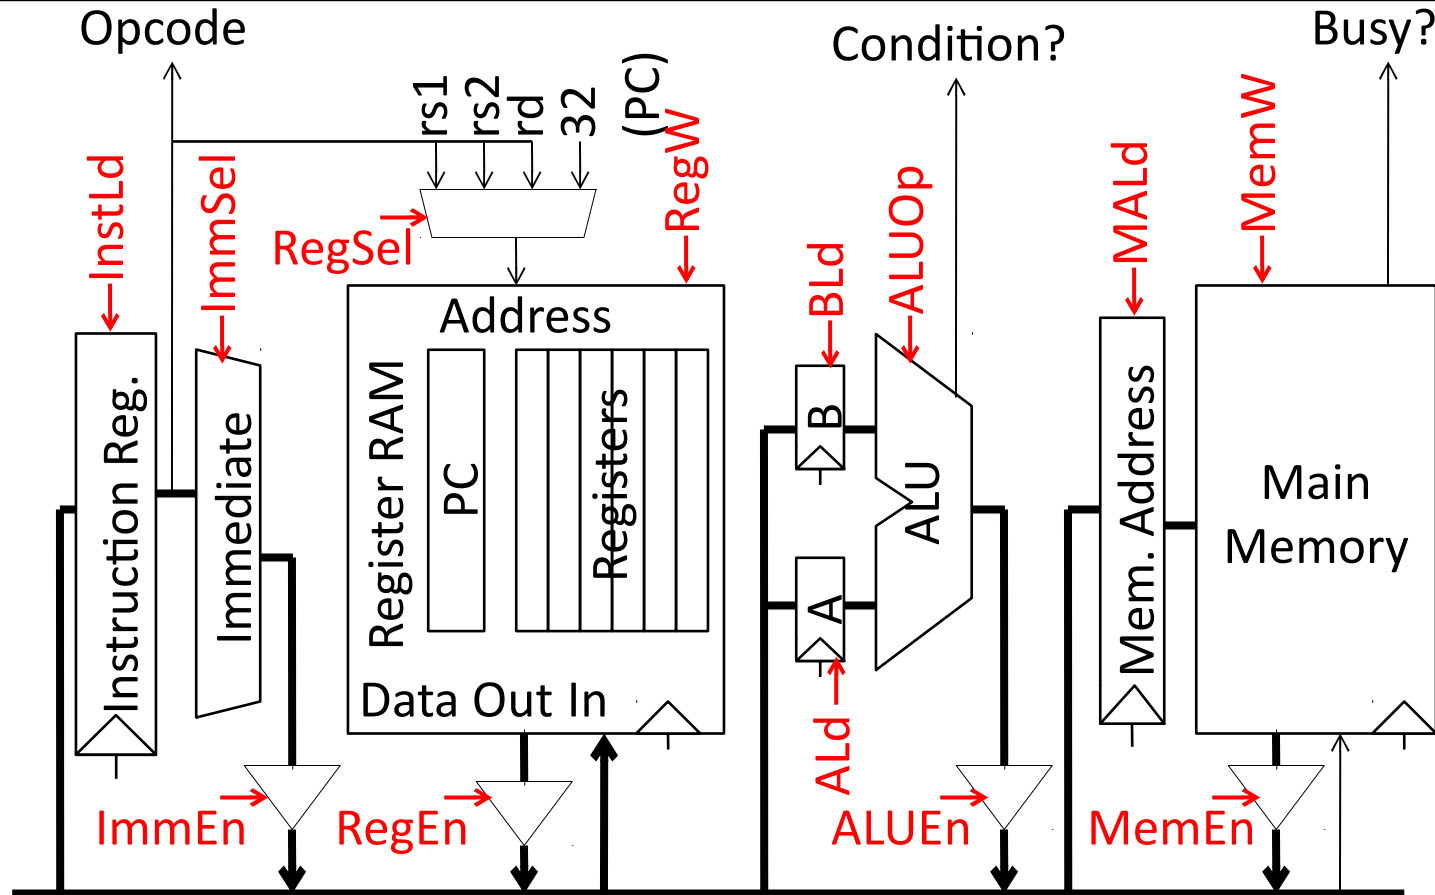
Control versus Datapath

- Processor designs can be split between datapath, where numbers are stored and arithmetic operations computed, and control, which sequences operations on datapath



- Biggest challenge for early computer designers was getting control circuitry correct
- Maurice Wilkes invented the idea of microprogramming to design the control unit of a processor for EDSAC-II, 1958
 - Foreshadowed by Babbage's "Barrel" and mechanisms in earlier programmable calculators

Single-Bus Datapath for Microcoded RISC-V



- Microinstructions written as register transfers:
 - $MA := PC$ means $RegSel = PC$; $RegW = 0$; $RegEn = 1$; $MALd = 1$
 - $B := Reg[rs2]$ means $RegSel = rs2$; $RegW = 0$; $RegEn = 1$; $BLd = 1$
 - $Reg[rd] := A + B$ means $ALUop = Add$; $ALUEn = 1$; $RegSel = rd$; $RegW = 1$

Control Logic Truth Table (incomplete)

Inst[31:0]	BrEq	BrLT	PCSel	ImmSel	BrUn	ASel	BSel	ALUSel	MemRW	RegWEn	WBSel
add	*	*	+4	*	*	Reg	Reg	Add	Read	1	ALU
sub	*	*	+4	*	*	Reg	Reg	Sub	Read	1	ALU
(R-R Op)	*	*	+4	*	*	Reg	Reg	(Op)	Read	1	ALU
addi	*	*	+4	I	*	Reg	Imm	Add	Read	1	ALU
lw	*	*	+4	I	*	Reg	Imm	Add	Read	1	Mem
sw	*	*	+4	S	*	Reg	Imm	Add	Write	0	*
beq	0	*	+4	B	*	PC	Imm	Add	Read	0	*
beq	1	*	ALU	B	*	PC	Imm	Add	Read	0	*
bne	0	*	ALU	B	*	PC	Imm	Add	Read	0	*
bne	1	*	+4	B	*	PC	Imm	Add	Read	0	*
blt	*	1	ALU	B	0	PC	Imm	Add	Read	0	*
bltu	*	1	ALU	B	1	PC	Imm	Add	Read	0	*
jalr	*	*	ALU	I	*	Reg	Imm	Add	Read	1	PC+4
jal	*	*	ALU	J	*	PC	Imm	Add	Read	1	PC+4
auipc	*	*	+4	U	*	PC	Imm	Add	Read	1	ALU

RISC-V Instruction Execution Phases

- Instruction Fetch
- Instruction Decode
- Register Fetch
- ALU Operations
- *Optional* Memory Operations
- *Optional* Register Writeback
- Calculate Next Instruction Address

Microcode Sketches (1)

Instruction Fetch: MA, A:=PC

PC:=A+4

wait for memory

IR:=Mem

dispatch on opcode

ALU: A:=Reg[rs1]

B:=Reg[rs2]

Reg[rd]:=ALUOp(A,B)

goto instruction fetch

ALUI: A:=Reg[rs1]

B:=ImmI //Sign-extend 12b immediate

Reg[rd]:=ALUOp(A,B)

goto instruction fetch

Microcode Sketches (2)

LW: A:=Reg[rs1]

B:=ImmI //Sign-extend 12b immediate

MA:=A+B

wait for memory

Reg[rd]:=Mem

goto instruction fetch

JAL: Reg[rd]:=A // Store return address

A:=A-4 // Recover original PC

B:=ImmJ // Jump-style immediate

PC:=A+B

goto instruction fetch

Branch: A:=Reg[rs1]

B:=Reg[rs2]

if (!ALUOp(A,B)) *goto instruction fetch* //Not taken

A:=PC //Microcode fall through if branch taken

A:=A-4

B:=ImmB// Branch-style immediate

PC:=A+B

goto instruction fetch

Control Realization Options

- ROM
 - “Read-Only Memory”
 - Regular structure
 - Can be easily reprogrammed
 - fix errors
 - add instructions
 - Popular when designing control logic manually
- Combinatorial Logic
 - Today, chip designers use logic synthesis tools to convert truth tables to networks of gates

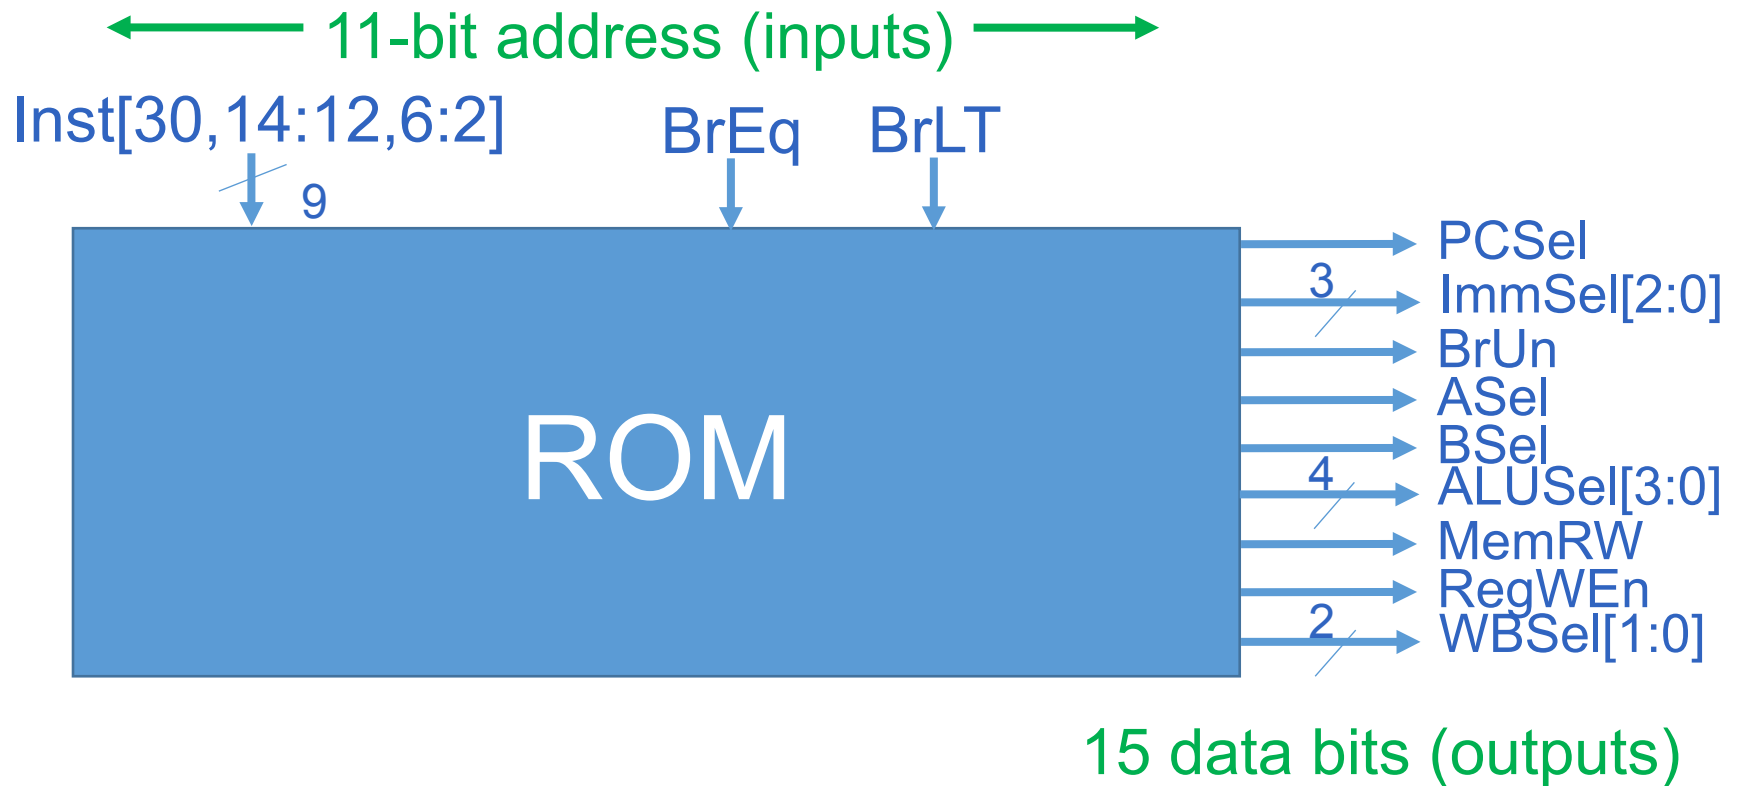
RV32I, a nine-bit ISA!

imm[31:12]				rd	011011	LUI
imm[31:12]				rd	001011	AUIPC
imm[20:10:11:19:12]				rd	110111	JAL
imm[11:0]				rd	110011	JALR
imm[12:10:5]	rs2	rs1	000	imm[4:1:11]	110001	BEQ
imm[12:10:5]	rs2	rs1	001	imm[4:1:11]	110001	BNE
imm[12:10:5]	rs2	rs1	100	imm[4:1:11]	110001	BLT
imm[12:10:5]	rs2	rs1	101	imm[4:1:11]	110001	BGE
imm[12:10:5]	rs2	rs1	110	imm[4:1:11]	110001	BLTU
imm[12:10:5]	rs2	rs1	111	imm[4:1:11]	110001	BGEU
imm[11:0]				rd	000001	LB
imm[11:0]				rd	000001	LH
imm[11:0]				rd	000001	LW
imm[11:0]				rd	000001	LBU
imm[11:0]				rd	000001	LHU
imm[11:5]	rs2	rs1	000	imm[4:0]	010001	SB
imm[11:5]	rs2	rs1	001	imm[4:0]	010001	SH
imm[11:5]	rs2	rs1	010	imm[4:0]	010001	SW
imm[11:0]				rd	001001	ADDI
imm[11:0]				rd	001001	SLTI
imm[11:0]				rd	001001	SLTIU
imm[11:0]				rd	001001	XORI
imm[11:0]				rd	001001	ORI
imm[11:0]				rd	001001	ANDI

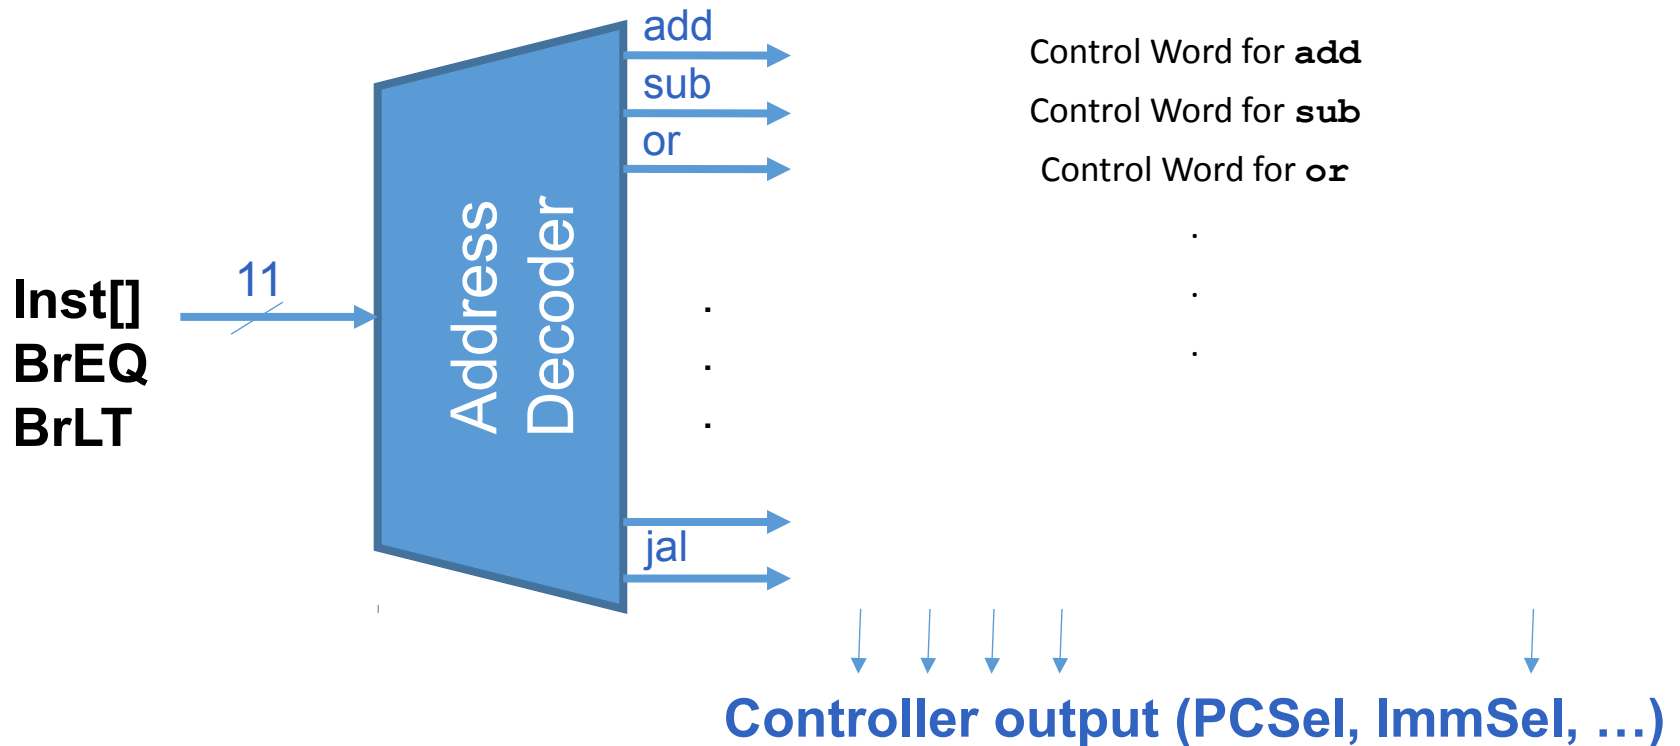
inst[30]			inst[14:12]		inst[6:2]		
0000000	shamt	rs1	001	rd	0010011	SLLI	
0000000	shamt	rs1	101	rd	0010011	SRLI	
0100000	shamt	rs1	101	rd	0010011	SRAI	
0000000	rs2	rs1	000	rd	0110011	ADD	
0100000	rs2	rs1	000	rd	0110011	SUB	
0000000	rs2	rs1	001	rd	0110011	SLL	
0000000	rs2	rs1	010	rd	0110011	SLT	
0000000	rs2	rs1	011	rd	0110011	SLTU	
0000000	rs2	rs1	100	rd	0110011	XOR	
0000000	rs2	rs1	101	rd	0110011	SRL	
0100000	rs2	rs1	101	rd	0110011	SRA	
0000000	rs2	rs1	110	rd	0110011	OR	
0000000	rs2	rs1	111	rd	0110011	AND	
0000	pred	succ	00000	000	00000	0001111	FENCE
0000	0000	0000	00000	001	00000	0001111	FENCE.I
0000000000000			00000	000	00000	1110011	ECALL
0000000000001			00000	000	00000	1110011	EBREAK
csr		rs1	001	rd	1110011	CSRRW	
csr		rs1	010	rd	1110011	CSRRS	
csr		rs1	011	rd	1110011	CSRRC	
csr		zimm	101	rd	1110011	CSRRWI	
csr		zimm	110	rd	1110011	CSRRSI	
csr		zimm	111	rd	1110011	CSRRCI	

Instruction type encoded using only 9 bits
inst[30],inst[14:12], inst[6:2]

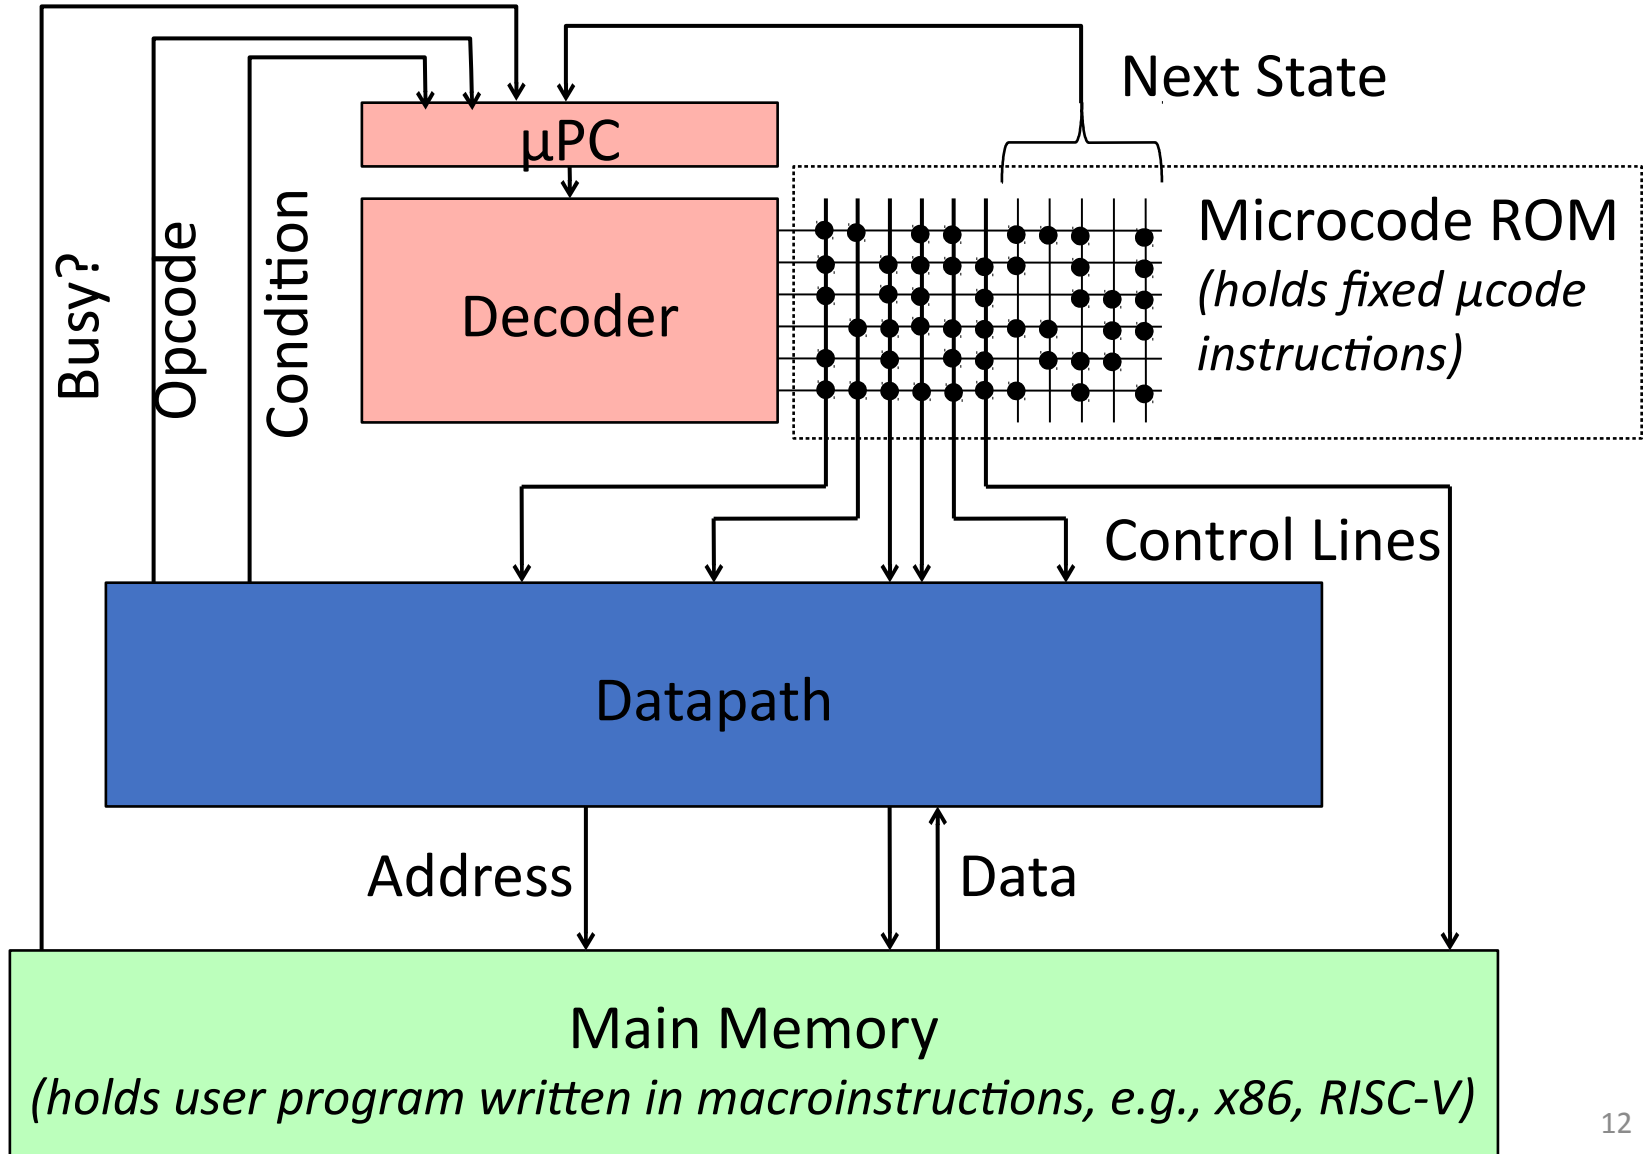
ROM-based Control



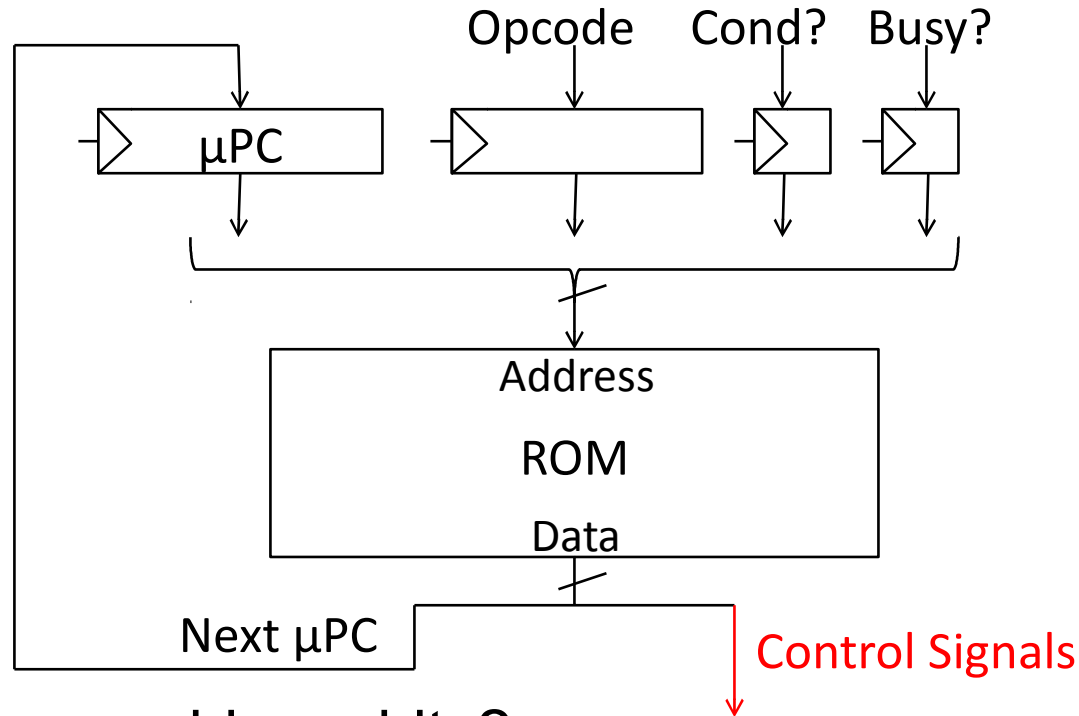
ROM Controller Implementation



Microcoded CPU



Pure ROM Implementation



- How many address bits?
 $|\mu\text{address}| = |\mu\text{PC}| + |\text{opcode}| + 1 + 1$
- How many data bits?
 $|\text{data}| = |\mu\text{PC}| + |\text{control signals}| = |\mu\text{PC}| + 18$
- Total ROM size = $2^{|\mu\text{address}|} \times |\text{data}|$

Pure ROM Contents

<u>Address</u>				<u> Data</u>		
<u>μPC</u>	<u>Opcode Cond? Busy?</u>			<u> Control Lines</u>	<u>Next μPC</u>	
fetch0	X	X	X	MA,A:=PC	fetch1	
fetch1	X	X	1		fetch1	
fetch1	X	X	0	IR:=Mem	fetch2	
fetch2	ALU	X	X	PC:=A+4	ALU0	
fetch2	ALUI		X	X PC:=A+4	ALUI0	
fetch2	LW	X	X	PC:=A+4	LW0	
....						
ALU0	X	X	X	A:=Reg[rs1]	ALU1	
ALU1	X	X	X	B:=Reg[rs2]	ALU2	
ALU2	X	X	X	Reg[rd]:=ALUOp(A,B)	fetch0	

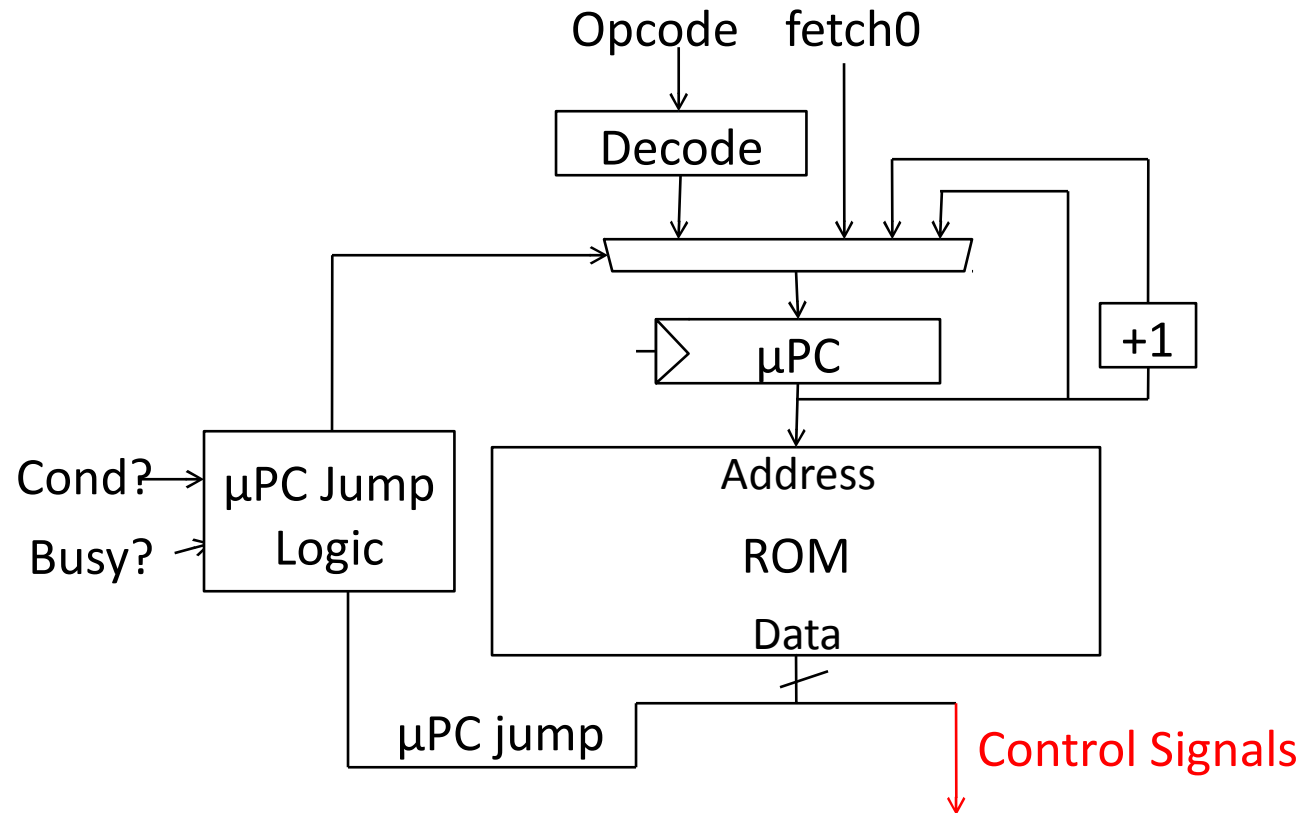
Single-Bus Microcode RISC-V ROM Size

- Instruction fetch sequence 3 common steps
- ~12 instruction groups
- Each group takes ~5 steps (1 for dispatch)
- Total steps $3 + 12 \times 5 = 63$, needs 6 bits for μ PC
- Opcode is 5 bits, ~18 control signals
- Total size = $2^{(6+5+2)} \times (6+18) = 2^{13} \times 24 = \sim 25\text{KiB!}$

Reducing Control Store Size

- Reduce ROM height (#address bits)
 - Use external logic to combine input signals
 - Reduce #states by grouping opcodes
- Reduce ROM width (#data bits)
 - Restrict μ PC encoding (next, dispatch, wait on memory,...)
 - Encode control signals (vertical μ coding, nanocoding)

Single-Bus RISC-V Microcode Engine



μPC jump = next | spin | fetch | dispatch | ftrue | ffalse

μPC Jump Types

- *next* increments μPC
- *spin* waits for memory
- *fetch* jumps to start of instruction fetch
- *dispatch* jumps to start of decoded opcode group
- *ftrue/ffalse* jumps to fetch if Cond? true/false

Encoded ROM Contents

<u>Address</u>	<u>Data</u>	
<u>μPC</u>	<u>Control Lines</u>	<u>Next μPC</u>
fetch0	MA,A:=PC	next
fetch1	IR:=Mem	spin
fetch2	PC:=A+4	dispatch
ALU0	A:=Reg[rs1]	next
ALU1	B:=Reg[rs2]	next
ALU2	Reg[rd]:=ALUOp(A,B)	fetch
Branch0	A:=Reg[rs1]	next
Branch1	B:=Reg[rs2]	next
Branch2	A:=PC	ffalse
Branch3	A:=A-4	next
Branch4	B:=ImmB	next
Branch5	PC:=A+B	fetch

Implementing Complex Instructions

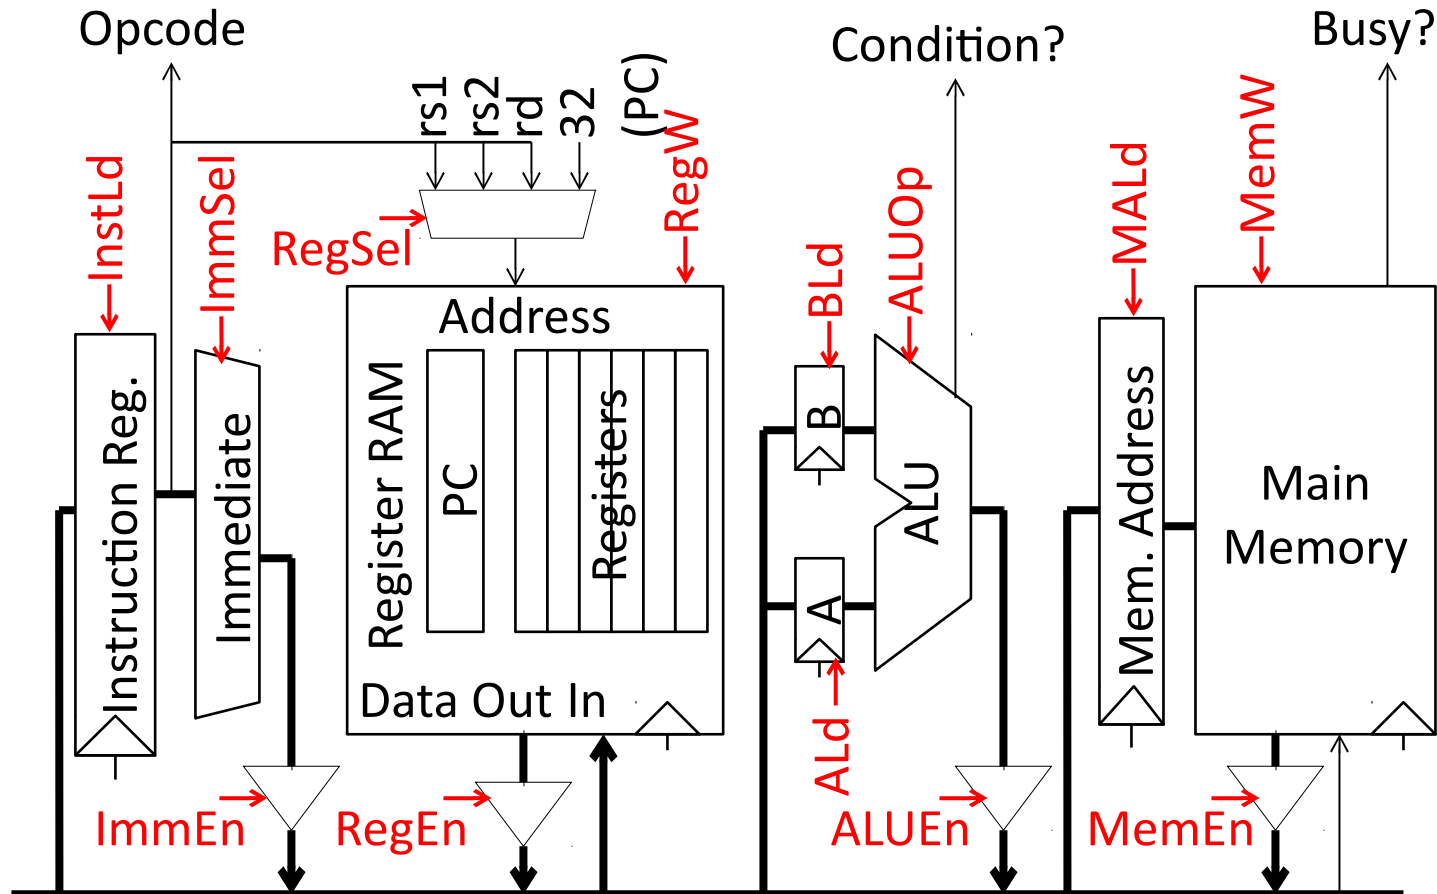
Memory-memory add: $M[rd] = M[rs1] + M[rs2]$

<u>Address</u>	<u>Data</u>
<u>μPC</u>	<u>Control Lines</u> <u>Next μPC</u>
MMA0	MA:=Reg[rs1] next
MMA1	A:=Mem spin
MMA2	MA:=Reg[rs2] next
MMA3	B:=Mem spin
MMA4	MA:=Reg[rd] next
MMA5	Mem:=ALUOp(A,B) spin
MMA6	fetch

Complex instructions usually do not require datapath modifications, only extra space for control program


Very difficult to implement these instructions using a hardwired controller without substantial datapath modifications

Single-Bus Datapath for Microcoded RISC-V



- Datapath unchanged for complex instructions!

Horizontal vs Vertical μ Code




- Horizontal μ code has wider μ instructions
 - Multiple parallel operations per μ instruction
 - Fewer microcode steps per macroinstruction
 - Sparser encoding \rightarrow more bits
- Vertical μ code has narrower μ instructions
 - Typically a single datapath operation per μ instruction
 - separate μ instruction for branches
 - More microcode steps per macroinstruction
 - More compact \rightarrow less bits
- Nanocoding
 - Tries to combine best of horizontal and vertical μ code

Nanocoding

- Exploits recurring control signal patterns in μ code, e.g.,

- ALU0 $A \leftarrow \text{Reg}[\text{rs1}]$
- ...
- ALUI0 $A \leftarrow \text{Reg}[\text{rs1}]$
- ...




- Motorola 68000 had 17-bit μ code containing either 10-bit μ jump or 9-bit nanoinstruction pointer
 - Nanoinstructions were 68 bits wide, decoded to give 196 control signals

Microprogramming in IBM 360

	M30	M40	M50	M65
Datapath width (bits)	8	16	32	64
μ inst width (bits)	50	52	85	87
μ code size (K μ insts)	4	4	2.75	2.75
μ store technology	CCROS	TCROS	BCROS	BCROS
μ store cycle (ns)	750	625	500	200
memory cycle (ns)	1500	2500	2000	750
Rental fee (\$K/month)	4	7	15	35

- Only the fastest models (75 and 95) were hardwired

IBM Card-Capacitor Read-Only Storage



[IBM Journal, January 1961]

Microcode Emulation

- IBM initially miscalculated the importance of software compatibility with earlier models when introducing the 360 series
- Honeywell stole some IBM 1401 customers by offering translation software (“Liberator”) for Honeywell H200 series machine
- IBM retaliated with optional additional microcode for 360 series that could emulate IBM 1401 ISA, later extended for IBM 7000 series
 - one popular program on 1401 was a 650 simulator, so some customers ran many 650 programs on emulated 1401s
 - i.e., 650 simulated on 1401 emulated on 360

Microprogramming thrived in '60s and '70s

- Significantly faster ROMs than DRAMs were available
- For complex instruction sets, datapath and controller were cheaper and simpler
- New instructions , e.g., floating point, could be supported without datapath modifications
- Fixing bugs in the controller was easier
- ISA compatibility across various models could be achieved easily and cheaply

Except for the cheapest and fastest machines, all computers were microprogrammed

Microprogramming thrived in '60s and '70s

- Evolution bred more complex micro-machines
 - Complex instruction sets led to need for subroutine and call stacks in μ code
 - Need for fixing bugs in control programs was in conflict with read-only nature of μ ROM
 - \rightarrow Writable Control Store (WCS) (B1700, QMachine, Intel i432, ...)
- With the advent of VLSI technology assumptions about ROM & RAM speed became invalid \rightarrow more complexity
- Better compilers made complex instructions less important.
- Use of numerous micro-architectural innovations, e.g., pipelining, caches and buffers, made multiple-cycle execution of reg-reg instructions unattractive

Writable Control Store (WCS)

- Implement control store in RAM not ROM
 - MOS SRAM memories now almost as fast as control store (core memories/DRAMs were 2-10x slower)
 - Bug-free microprograms difficult to write
- User-WCS provided as option on several minicomputers
 - Allowed users to change microcode for each processor
- User-WCS failed
 - Little or no programming tools support
 - Difficult to fit software into small space
 - Microcode control tailored to original ISA, less useful for others
 - Large WCS part of processor state - expensive context switches
 - Protection difficult if user can change microcode
 - Virtual memory required restartable microcode

Analyzing Microcoded Machines

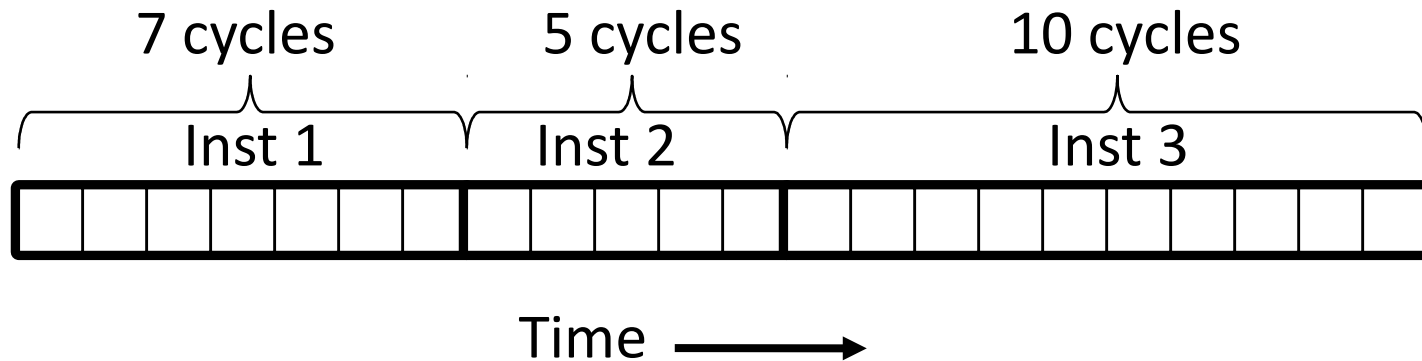
- John Cocke and group at IBM
 - Working on a simple pipelined processor, 801, and advanced compilers inside IBM
 - Ported experimental PL.8 compiler to IBM 370, and only used simple register-register and load/store instructions similar to 801
 - Code ran faster than other existing compilers that used all 370 instructions! (up to 6MIPS whereas 2MIPS considered good before)
- Emer, Clark, at DEC
 - Measured VAX-11/780 using external hardware
 - Found it was actually a 0.5MIPS machine, although usually assumed to be a 1MIPS machine
 - Found 20% of VAX instructions responsible for 60% of microcode, but only account for 0.2% of execution time!
 - VAX8800
- Control Store: 16K*147b RAM, Unified Cache: 64K*8b RAM
 - 4.5x more microstore RAM than cache RAM!

“Iron Law” of Processor Performance

$$\frac{\text{Time}}{\text{Program}} = \frac{\text{Instructions}}{\text{Program}} * \frac{\text{Cycles}}{\text{Instruction}} * \frac{\text{Time}}{\text{Cycle}}$$

- Instructions per program depends on source code, compiler technology, and ISA
- Cycles per instructions (CPI) depends on ISA and μ architecture
- Time per cycle depends upon the μ architecture and base technology

CPI for Microcoded Machine



Total clock cycles = $7+5+10 = 22$

Total instructions = 3

$CPI = 22/3 = 7.33$

CPI is always an average over a large number of instructions.

IC Technology Changes Tradeoffs

- Logic, RAM, ROM all implemented using MOS transistors
- Semiconductor RAM ~ same speed as ROM

Reconsidering Microcode Machine (Nane 68000 example)

RISC!

Exploits recurring control
signal patterns in μ code,
e.g.,


ALU0 $A \leftarrow \text{Reg}[\text{rs1}]$

...

ALUI0 $A \leftarrow \text{Reg}[\text{rs1}]$

...

- Motorola 68000 had 17-bit μ code containing either 10-bit μ jump or 9-bit nanoinstruction pointer
 - Nanoinstructions were 68 bits wide, decoded to give 196 control signals





From CISC to RISC

- Use fast RAM to build fast instruction cache of user-visible instructions, not fixed hardware microroutines
 - Contents of fast instruction memory change to fit application needs
 - Use simple ISA to enable hardwired pipelined implementation
- Most compiled code only used few CISC instructions
 - Simpler encoding allowed pipelined implementations
- Further benefit with integration
 - In early '80s, finally fit 32-bit datapath + small caches on single chip
 - No chip crossings in common case allows faster operation

Berkeley RISC Chips

RISC-I (1982) Contains 44,420 transistors, fabbed in 5 μm NMOS, with a die area of 77 mm^2 , ran at 1 MHz. This chip is probably the first VLSI RISC.



RISC-II (1983) contains 40,760 transistors, was fabbed in 3 μm NMOS, ran at 3 MHz, and the size is 60 mm^2 .

Stanford built some too...