

# 计算机组成与系统结构

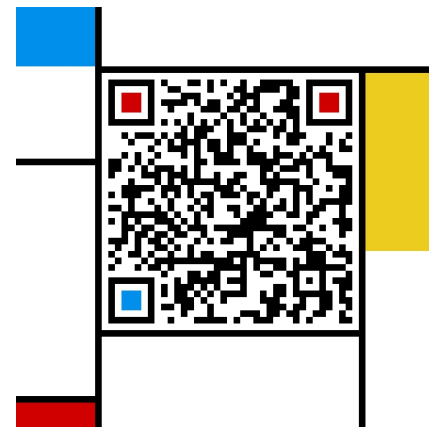
## Computer Organization & System Architecture

Huang Kejie (黄科杰) 百人计划研究员

Office: 玉泉校区老生仪楼 304

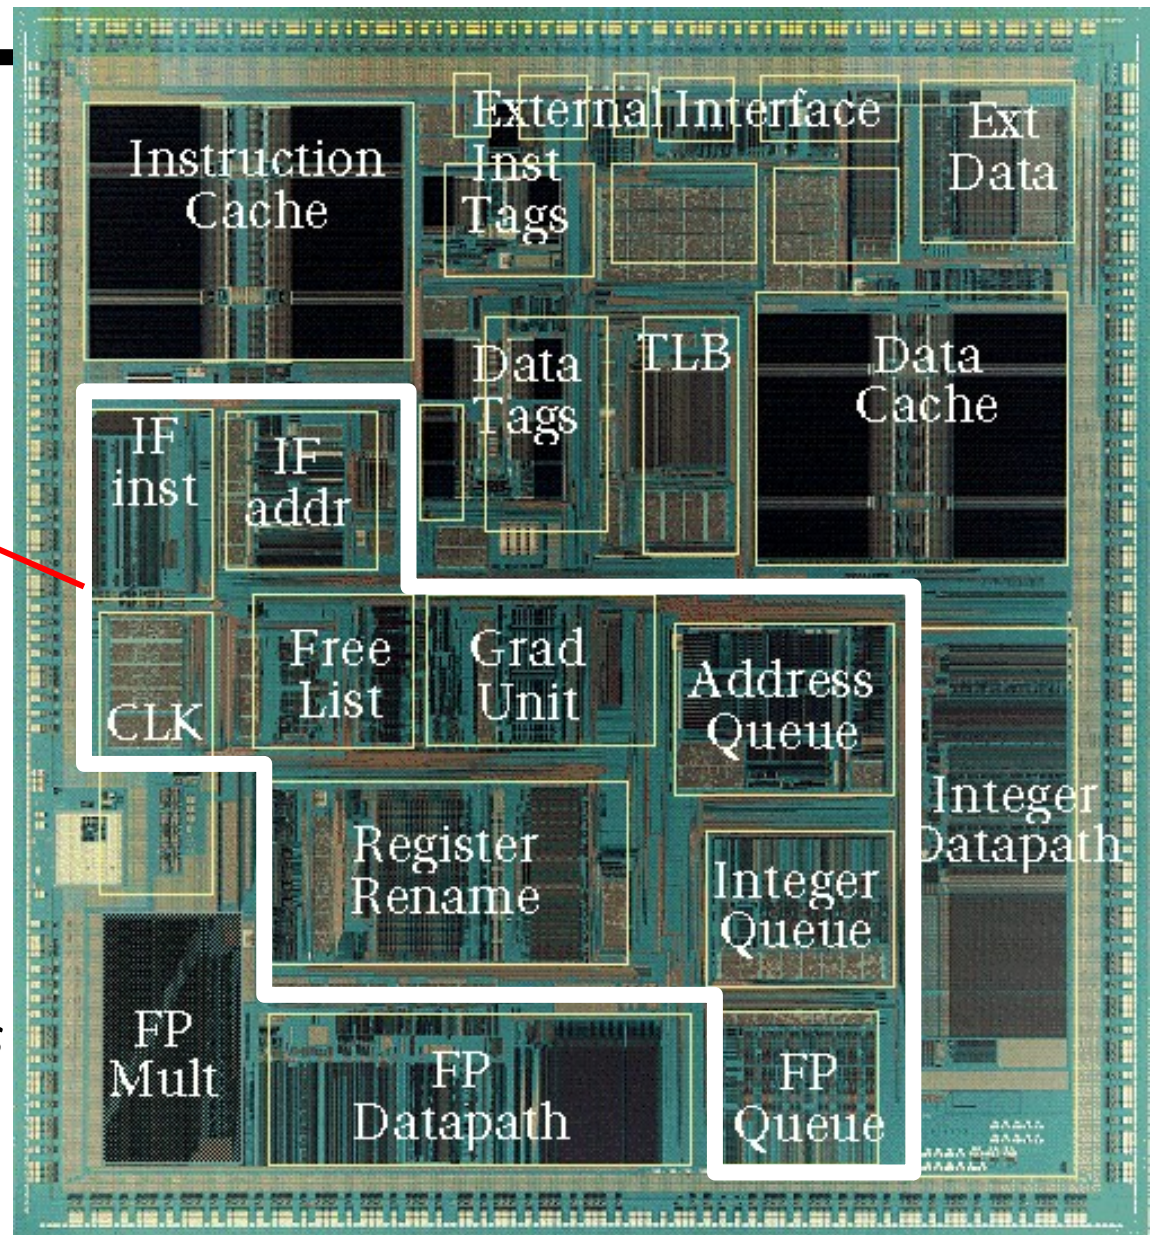
Email address: [huangkejie@zju.edu.cn](mailto:huangkejie@zju.edu.cn)

HP: 17706443800



# Out-of-Order Control Complexity: MIPS R10000

*Control  
Logic*



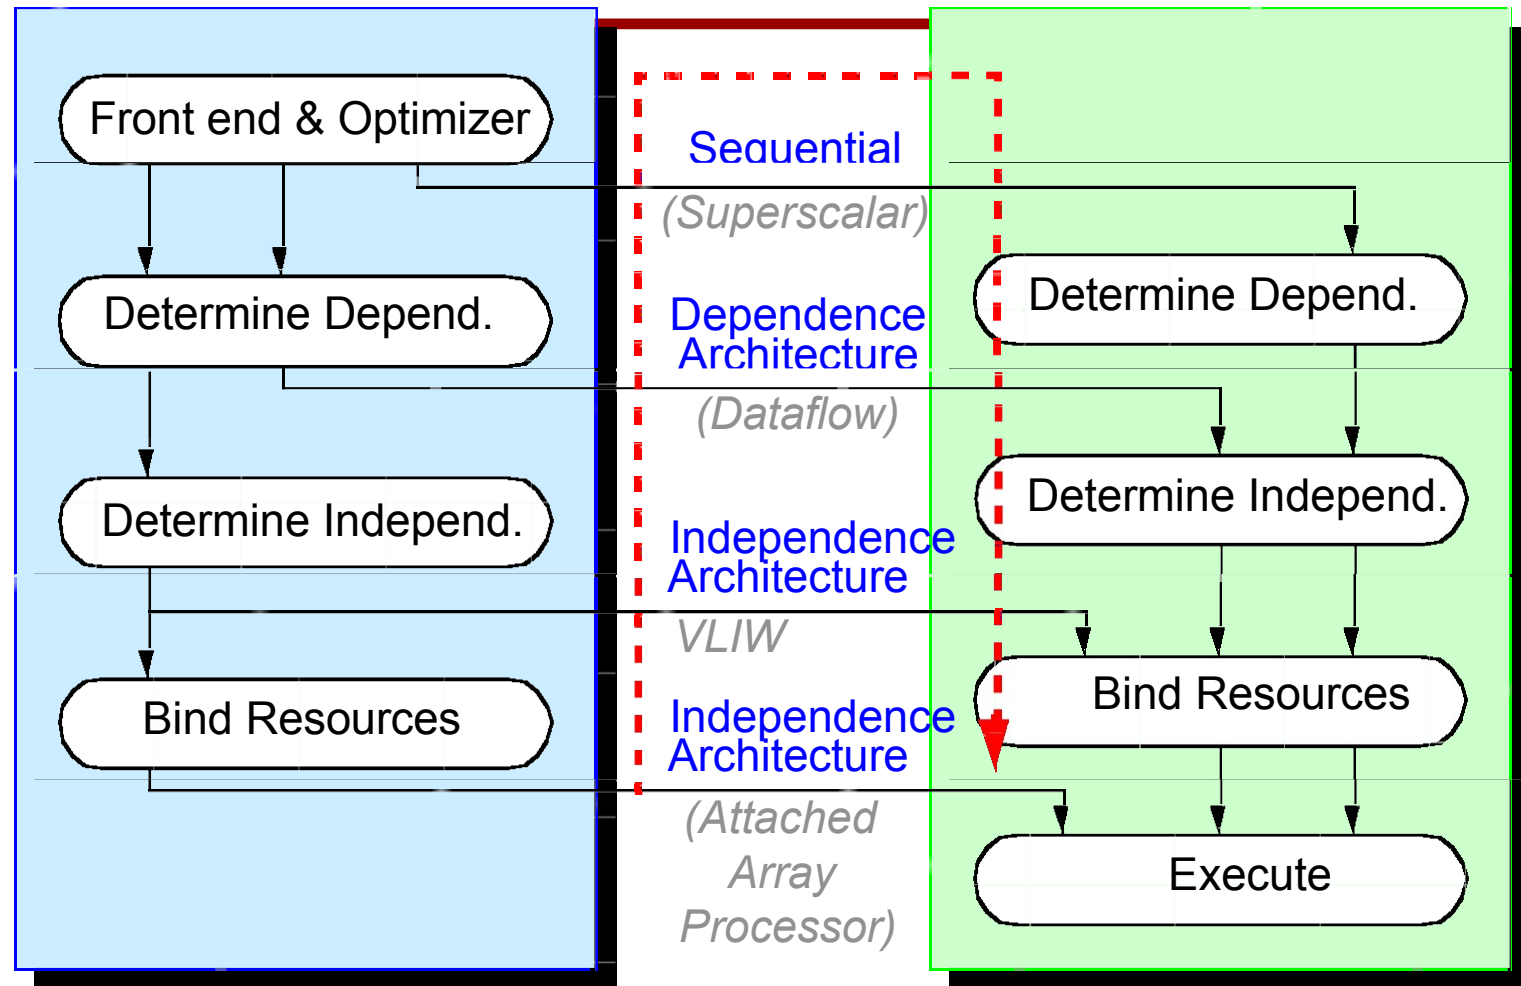
*[ SGI/MIPS Technologies  
Inc., 1995 ]*

# The Paradox of Superscalar Processing

---

- Compiler
  - Analyze sequential program to identify independent instructions
  - Produce sequential schedule with dependent instructions spaced apart
  - Map variables to a small set of registers by maximizing reuse
- Superscalar compiler
  - Analyze sequential schedule to identify independent instructions
  - Schedule instructions for parallel execution
  - Remap small register set into a large register set
- Idea behind VLIW
  - Design ISAs, compiler, and hardware the work synergistically for ILP

# HW/SW DesignSpace for ILP

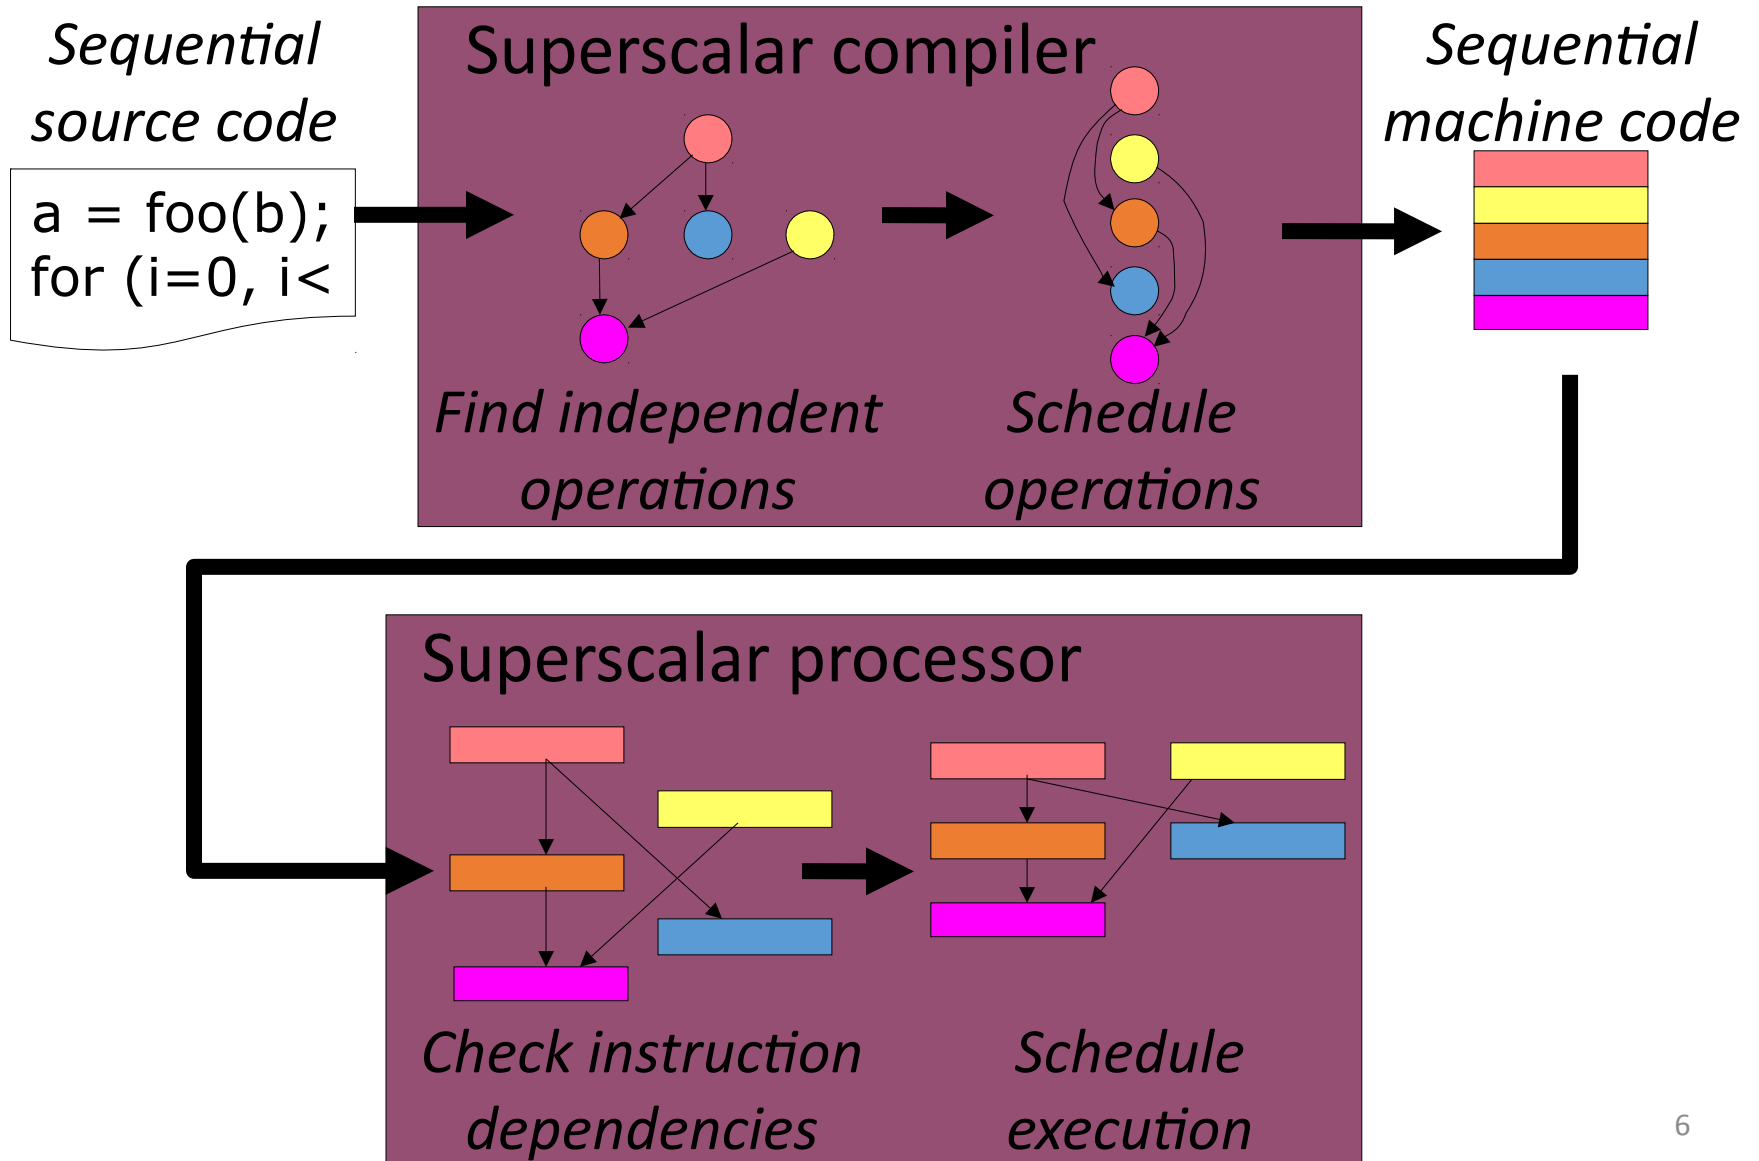


Compiler

Hardware

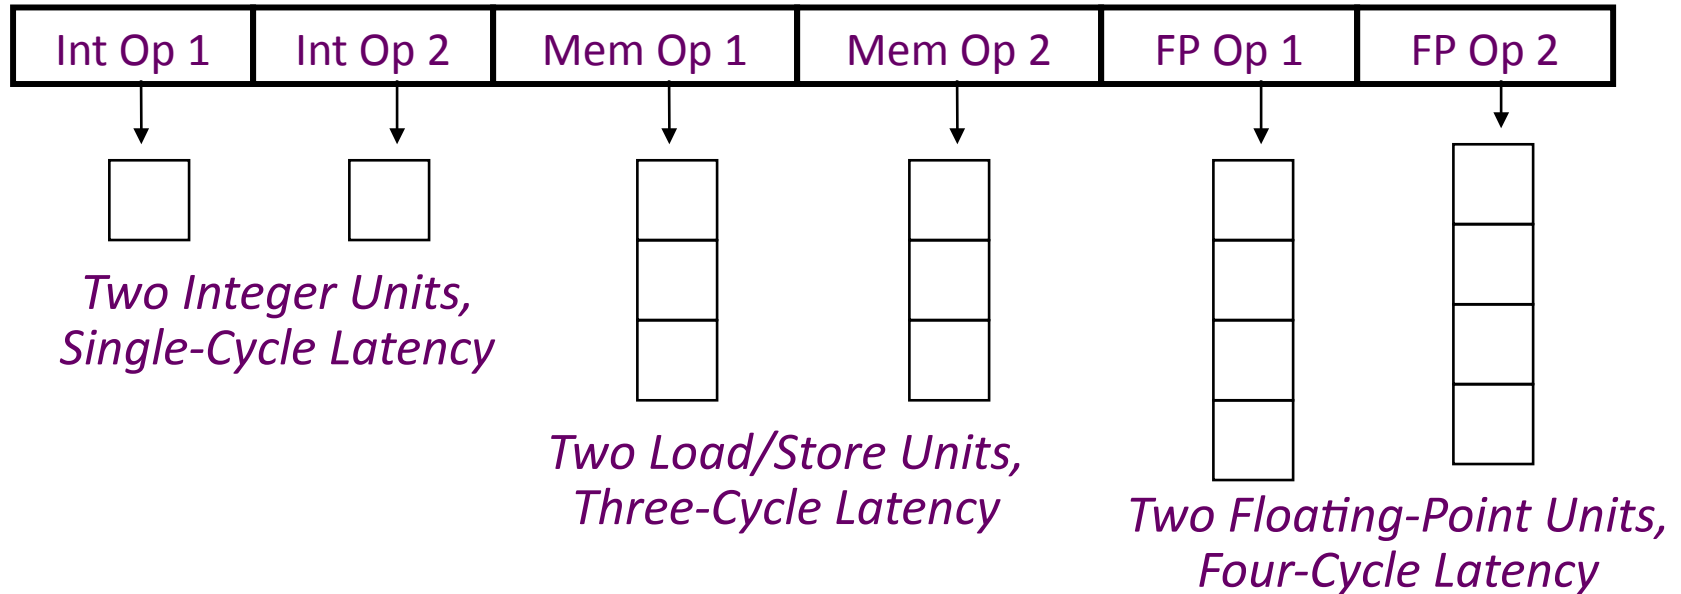
[B. Rau & J. Fisher, 1993]

# Sequential ISA Bottleneck



# VLIW: Very Long Instruction Word

---



- Multiple operations packed into one instruction
- Each operation slot is for a fixed function
- Constant operation latencies are specified
- Architecture requires guarantee of:
  - Parallelism within an instruction => no cross-operation RAW check
  - No data use before data ready => no data interlocks

# Early VLIW Machines

---

- FPS AP120B (1976)
  - scientific attached array processor
  - first commercial wide instruction machine
  - hand-coded vector math libraries using software pipelining and loop unrolling
- Multiflow Trace (1987)
  - commercialization of ideas from Fisher's Yale group including "trace scheduling"
  - available in configurations with 7, 14, or 28 operations/instruction
  - 28 operations packed into a 1024-bit instruction word
- Cydrome Cydra-5 (1987)
  - 7 operations encoded in 256-bit instruction word
  - rotating register file

# VLIW Compiler Responsibilities

---

- Schedule operations to maximize parallel execution
- Guarantees intra-instruction parallelism
- Schedule to avoid data hazards (no interlocks)
  - Typically separates operations with explicit NOPs




# Loop Execution

```
for (i=0; i<N; i++)  
    B[i] = A[i] + C;
```

↓ *Compile*

```
loop: fld f1, 0(x1)  
      add x1, 8  
      fadd f2, f0, f1  
      fsd f2, 0(x2)  
      add x2, 8  
      bne x1, x3,  
loop
```

loop



How many FP ops/cycle?

1 fadd / 8 cycles = 0.125

# Loop Unrolling

---

```
for (i=0; i<N; i++)  
    B[i] = A[i] + C;
```



Unroll inner loop to perform 4 iterations at once

```
for (i=0; i<N; i+=4)  
{  
    B[i]    = A[i] + C;  
    B[i+1] = A[i+1] + C;  
    B[i+2] = A[i+2] + C;  
    B[i+3] = A[i+3] + C;  
}
```

Need to handle values of N that are not multiples of unrolling factor with final cleanup loop

# Scheduling Loop Unrolled Code

*Unroll 4 ways*

```
loop: fld f1, 0(x1)
      fld f2, 8(x1)
      fld f3, 16(x1)
      fld f4, 24(x1)
      add x1, 32
      fadd f5, f0, f1
      fadd f6, f0, f2
      fadd f7, f0, f3
      fadd f8, f0, f4
      fsd f5, 0(x2)
      fsd f6, 8(x2)
      fsd f7, 16(x2)
      fsd f8, 24(x2)
      add x2, 32
      bne x1, x3, loop
```

loop:

*Schedule* →

Int1	Int2	M1	M2	FP+	FPx
		fld f1			
		fld f2			
		fld f3			
add x1		fld f4		fadd f5	
				fadd f6	
				fadd f7	
				fadd f8	
		fsd f5			
		fsd f6			
		fsd f7			
add x2	bne	fsd f8			

How many FLOPS/cycle?

4 fadds / 11 cycles = 0.36

# Software Pipelining

*Unroll 4 ways first*

```

loop: fld f1, 0(x1)
      fld f2, 8(x1)
      fld f3, 16(x1)
      fld f4, 24(x1)
      add x1, 32
      fadd f5, f0, f1
      fadd f6, f0, f2
      fadd f7, f0, f3
      fadd f8, f0, f4
      fsd f5, 0(x2)
      fsd f6, 8(x2)
      fsd f7, 16(x2)
      add x2, 32
      fsd f8, -8(x2)
      bne x1, x3, loop
    
```

prolog

iterate

loop:



epilog

Int1	Int 2	M1	M2	FP+	FPx
		fld f1			
		fld f2			
		fld f3			
add x1		fld f4			
		fld f1		fadd f5	
		fld f2		fadd f6	
		fld f3		fadd f7	
add x1		fld f4		fadd f8	
		fld f1	fsd f5	fadd f5	
		fld f2	fsd f6	fadd f6	
	add x2	fld f3	fsd f7	fadd f7	
add x1 bne		fld f4	fsd f8	fadd f8	
			fsd f5	fadd f5	
			fsd f6	fadd f6	
	add x2		fsd f7	fadd f7	
	bne		fsd f8	fadd f8	
			fsd f5		

How many FLOPS/cycle?

4 fadds / 4 cycles = 1


# Software Pipelining vs. Loop Unrolling



*Software pipelining pays startup/wind-down costs only once per loop, not once per iteration*

# What if there are no loops?


---



- Branches limit basic block size in control-flow intensive irregular code
- Difficult to find ILP in individual basic blocks


# Trace Scheduling [ Fisher,Ellis]

---



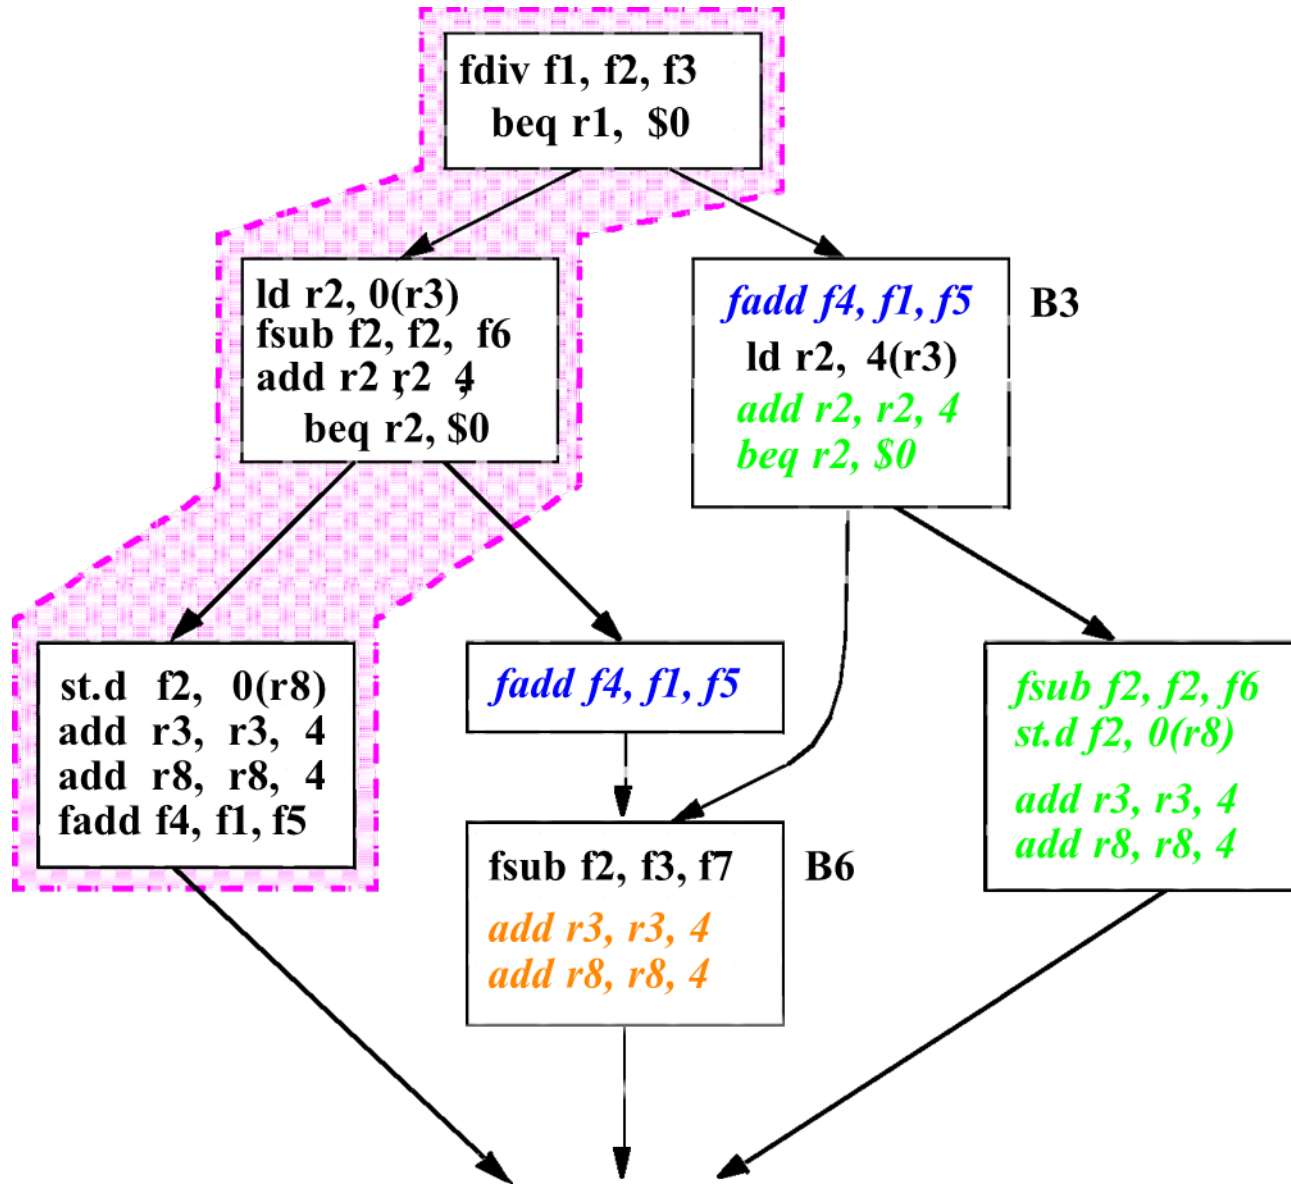
- Pick string of basic blocks, a trace, that represents most frequent branch path
- Use profiling feedback or compiler heuristics to find common branch paths
- Schedule whole “trace” at once
- Add fixup code to cope with branches jumping out of trace

# Trace Scheduling Example





# Compensation Code Illustration




# Problems with “Classic” VLIW

---

- Object-code compatibility
  - have to recompile all code for every machine, even for two machines in same generation
- Object code size
  - instruction padding wastes instruction memory/cache
  - loop unrolling/software pipelining replicates code
- Scheduling variable latency memory operations
  - caches and/or memory bank conflicts impose statically unpredictable variability
- Knowing branch probabilities
  - Profiling requires an significant extra step in build process
- Scheduling for statically unpredictable branches
  - optimal schedule varies with branch path

# VLIW Instruction Encoding

---



- Schemes to reduce effect of unused fields
  - Compressed format in memory, expand on I-cache refill
    - used in Multiflow Trace
    - introduces instruction addressing challenge
  - Mark parallel groups
    - used in TMS320C6x DSPs, Intel IA-64
  - Provide a single-op VLIW instruction
    - Cydra-5 UniOp instructions


# Intel Itanium, EPIC IA-64

---

- EPIC is the style of architecture (cf. CISC, RISC)
  - Explicitly Parallel Instruction Computing (really just VLIW)
- IA-64 is Intel's chosen ISA (cf. x86, MIPS)
  - IA-64 = Intel Architecture 64-bit
  - An object-code-compatible VLIW
- Merced was first Itanium implementation (cf. 8086)
  - First customer shipment expected 1997 (actually 2001)
  - McKinley, second implementation shipped in 2002
  - Recent version, Poulson, eight cores, 32nm, announced 2011


# Eight Core Itanium “Poulson” [Intel 2011]

---




- 8 cores
- 1-cycle 16KB L1 I&D caches
- 9-cycle 512KB L2 I-cache
- 8-cycle 256KB L2 D-cache
- 32 MB shared L3 cache
- 544mm<sup>2</sup> in 32nm CMOS
- Over 3 billion transistors
- Cores are 2-way multithreaded
- 6 instruction/cycle fetch
  - Two 128-bit bundles
- Up to 12 insts/cycle execute

# IA-64 Instruction Format



- Template bits describe grouping of these instructions with others in adjacent bundles
- Each group contains instructions that can execute in parallel



# IA-64 Registers

---

- 128 General Purpose 64-bit Integer Registers
- 128 General Purpose 64/80-bit Floating Point Registers
- 64 1-bit Predicate Registers
  
- GPRs “rotate” to reduce code size for software pipelined loops
  - Rotation is a simple form of register renaming allowing one instruction to address different physical registers on each iteration

# Rotating Register Files

---


Problems: Scheduled loops require lots of registers,  
Lots of duplicated code in prolog, epilog

Solution: Allocate new set of registers for each loop iteration



# Rotating Register File

---



Rotating Register Base (RRB) register points to base of current register set. Value added on to logical register specifier to give physical register number. Usually, split into rotating and non-rotating registers.

# Rotating Register File (Previous Loop Example)

Three cycle load latency  
encoded as difference of 3 in  
register specifier number (f4  
- f1 = 3)

Four cycle fadd latency  
encoded as difference of 4 in  
register specifier number (f9  
- f5 = 4)

ld f1, ()	fadd f5, f4, ...	sd f9, ()	bloop
-----------	------------------	-----------	-------

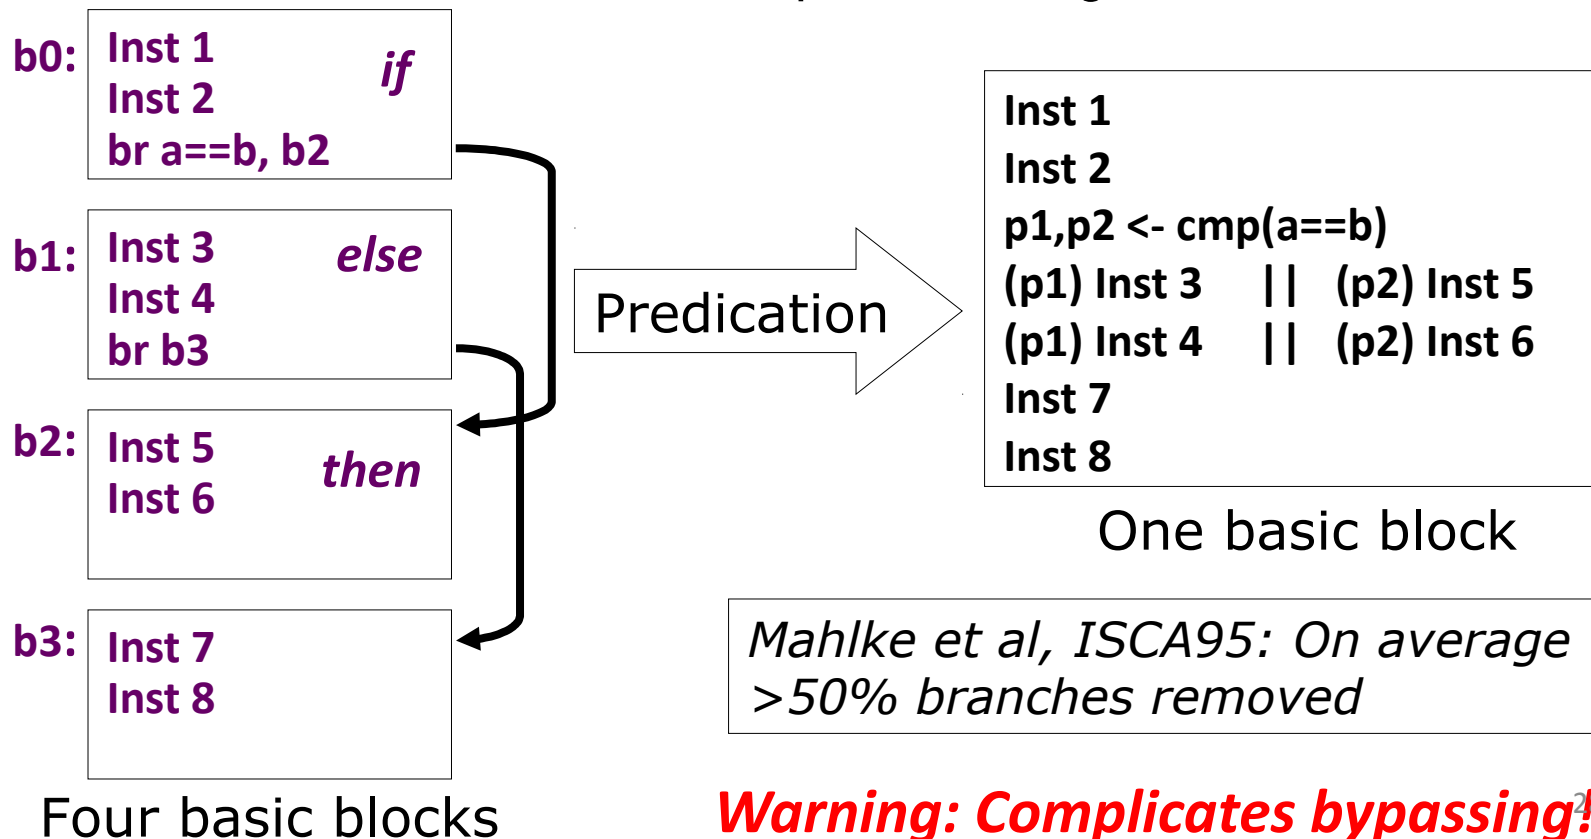
ld P9, ()	fadd P13, P12,	sd P17, ()	bloop	RRB=8
ld P8, ()	fadd P12, P11,	sd P16, ()	bloop	RRB=7
ld P7, ()	fadd P11, P10,	sd P15, ()	bloop	RRB=6
ld P6, ()	fadd P10, P9,	sd P14, ()	bloop	RRB=5
ld P5, ()	fadd P9, P8,	sd P13, ()	bloop	RRB=4
ld P4, ()	fadd P8, P7,	sd P12, ()	bloop	RRB=3
ld P3, ()	fadd P7, P6,	sd P11, ()	bloop	RRB=2
ld P2, ()	fadd P6, P5,	sd P10, ()	bloop	RRB=1

# IA-64 Predicated Execution

**Problem:** Mispredicted branches limit ILP

**Solution:** Eliminate hard to predict branches with predicated execution


- Almost all IA-64 instructions can be executed conditionally under predicate
- Instruction becomes NOP if predicate register false




# IA-64 Speculative Execution

**Problem:** Branches restrict compiler code motion

**Solution:** Speculative operations that don't cause exceptions



*Can't move load above branch  
because might cause spurious  
exception*



*Speculative load  
never causes  
exception, but sets  
"poison" bit on  
destination register*

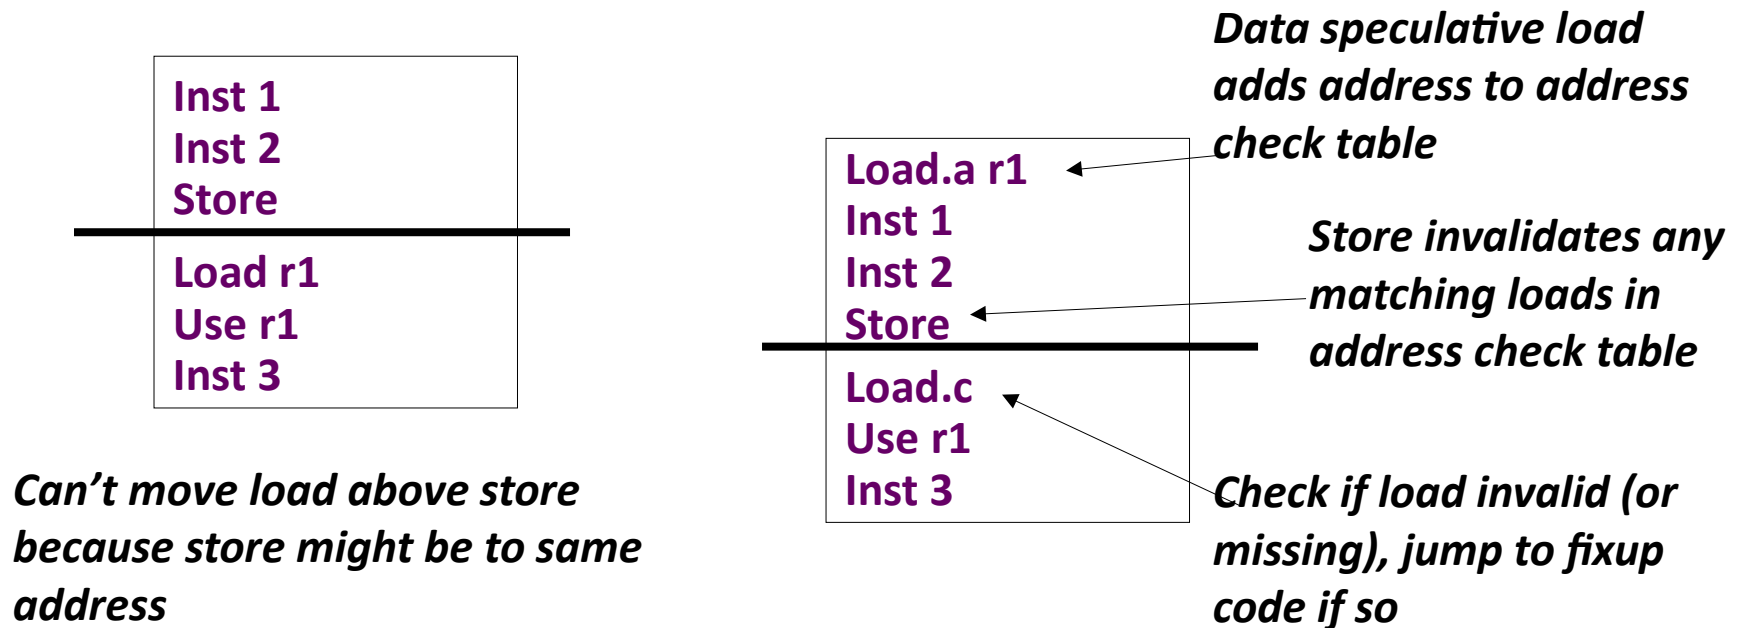
*Check for exception in  
original home block  
jumps to fixup code if  
exception detected*

Particularly useful for scheduling long latency loads early

# IA-64 Data Speculation

**Problem:** Possible memory hazards limit code scheduling

**Solution:** Hardware to check pointer hazards



Requires associative hardware in address check table

# Limits of Static Scheduling

---

- Unpredictable branches
- Variable memory latency (unpredictable cache misses)
- Code size explosion
- Compiler complexity
- Despite several attempts, VLIW has failed in general-purpose computing arena (so far).
  - More complex VLIW architectures are close to in-order superscalar in complexity, no real advantage on large complex apps.
- Successful in embedded DSP market
  - Simpler VLIWs with more constrained environment, friendlier code.

# Intel Kills Itanium

---

- Donald Knuth “ ... *Itanium approach that was supposed to be so terrific—until it turned out that the wished-for compilers were basically impossible to write.*”
- “*Intel officially announced the end of life and product discontinuance of the Itanium CPU family on January 30th, 2019*”, Wikipedia