

# 计算机组成与系统结构

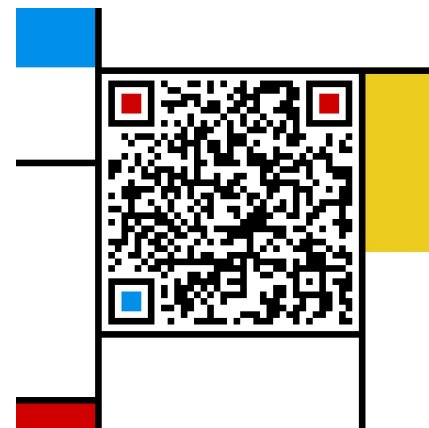
## Computer Organization & System Architecture

Huang Kejie (黄科杰) 百人计划研究员

Office: 玉泉校区老生仪楼 304

Email address: [huangkejie@zju.edu.cn](mailto:huangkejie@zju.edu.cn)

HP: 17706443800

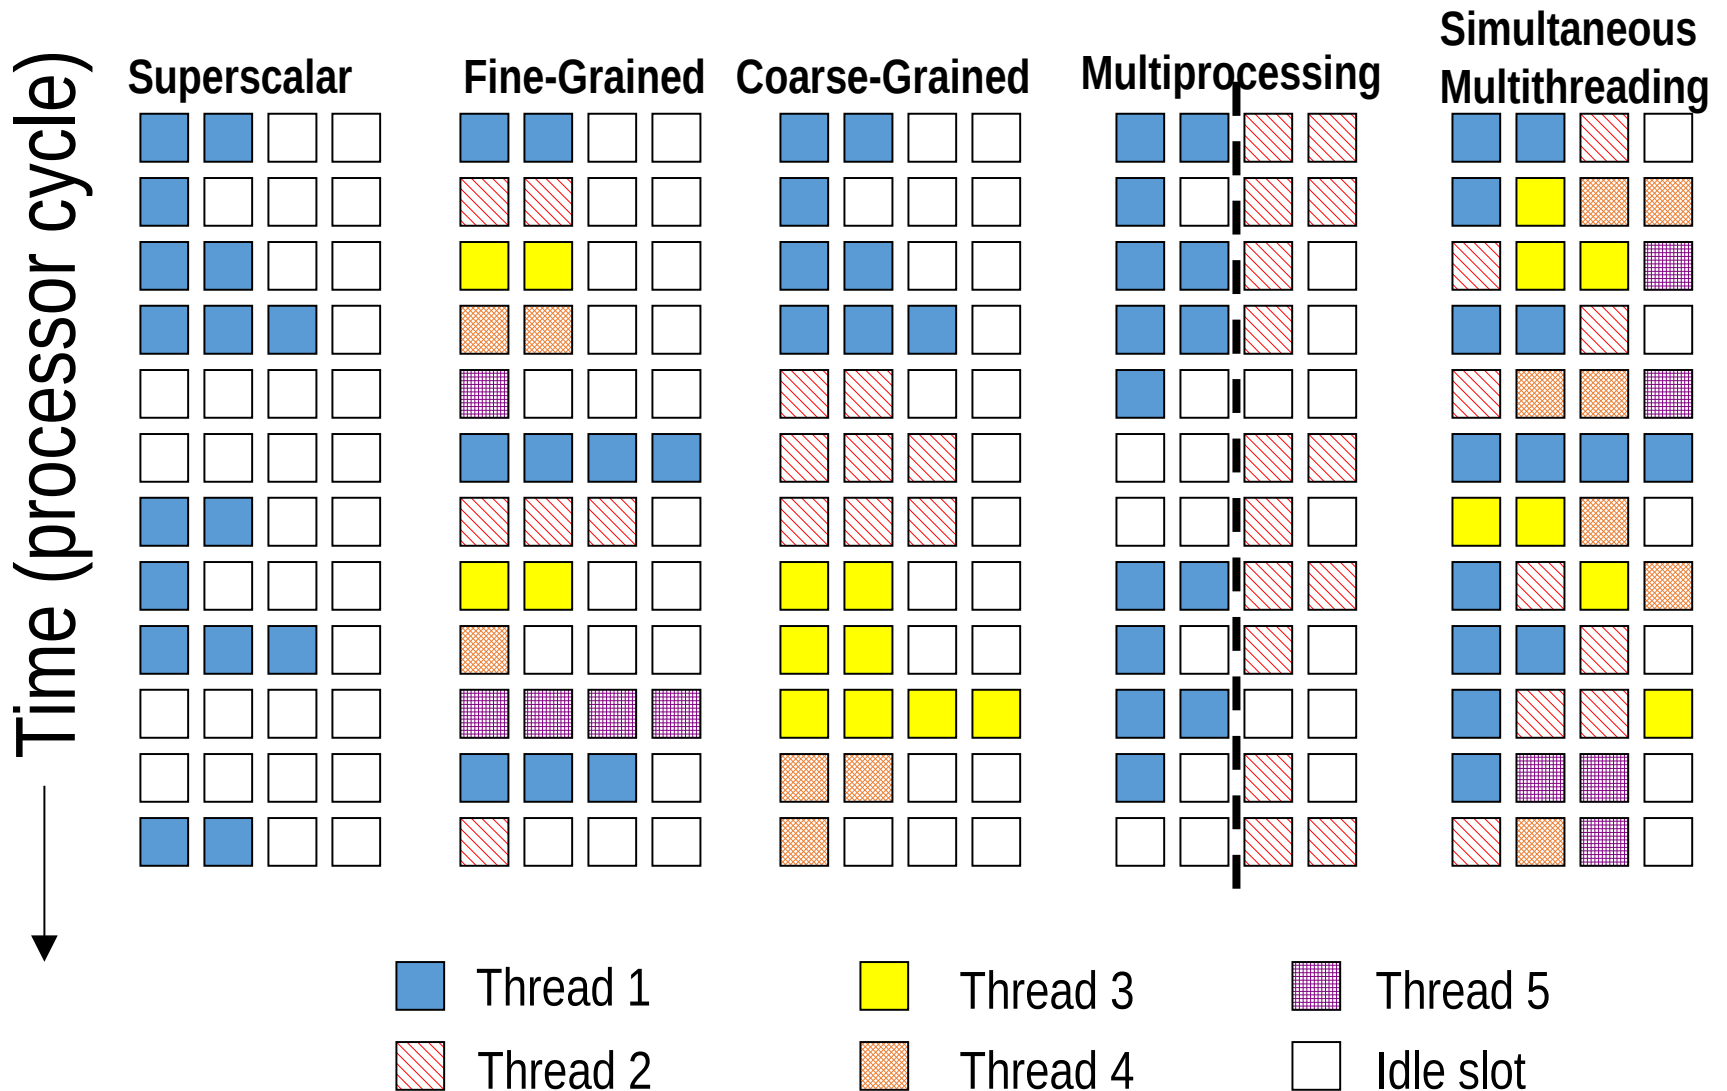


# Project Presentation

---

- The project presentation. Contents may cover but not limit to
  - The contribution of each member
  - Design review of MP2
    - The implementation of MP2
    - Test benches and simulation results
    - FPGA implementation and timing analysis
  - Proposal for MP3 (no need to write SV code)
    - The design plan to add 5-stage pipeline in MP2
    - The conflicts may meet in this design
    - Your solutions to address issues
    - Optimization in terms of area and speed
- Presentation date: Dec 16 and Dec 18
- 2 person in group, each group will have 12 minutes to present the work
- 1<sup>st</sup> project: 10 points, 2<sup>nd</sup> project: 20 points, presentation: 10 points

# ReCap: Multithreading



# Supercomputer Applications

---

- Typical application areas
  - Military research (nuclear weapons, cryptography)
  - Scientific research
  - Weather forecasting
  - Oil exploration
  - Industrial design (car crash simulation)
  - Bioinformatics
  - Cryptography
- All involve huge computations on large data set
- Supercomputers: CDC6600, CDC7600, Cray-1, ...
- In 70s-80s, Supercomputer  $\equiv$  Vector Machine

# Vector Supercomputers




- Epitomized by Cray-1, 1976:
- Scalar Unit
  - Load/Store Architecture
  - Vector Extension
  - Vector Registers
  - Vector Instructions
- Implementation
  - Hardwired Control
  - Highly Pipelined Functional Units
  - Interleaved Memory System
  - No Data Caches
  - No Virtual Memory


[©Cray Research, 1976]

# Vector Programming Model

*Scalar Registers*




*Vector Registers*



*Vector Length Register* **v1**


Vector Arithmetic  
Instructions

**vadd v3, v1, v2**



Vector Load and Store  
Instructions

**vldst v1, x1, x2**



# Vector Code Example

---

```
# C code
```

```
for (i=0; i<64; i++)  
    C[i] = A[i] + B[i];
```

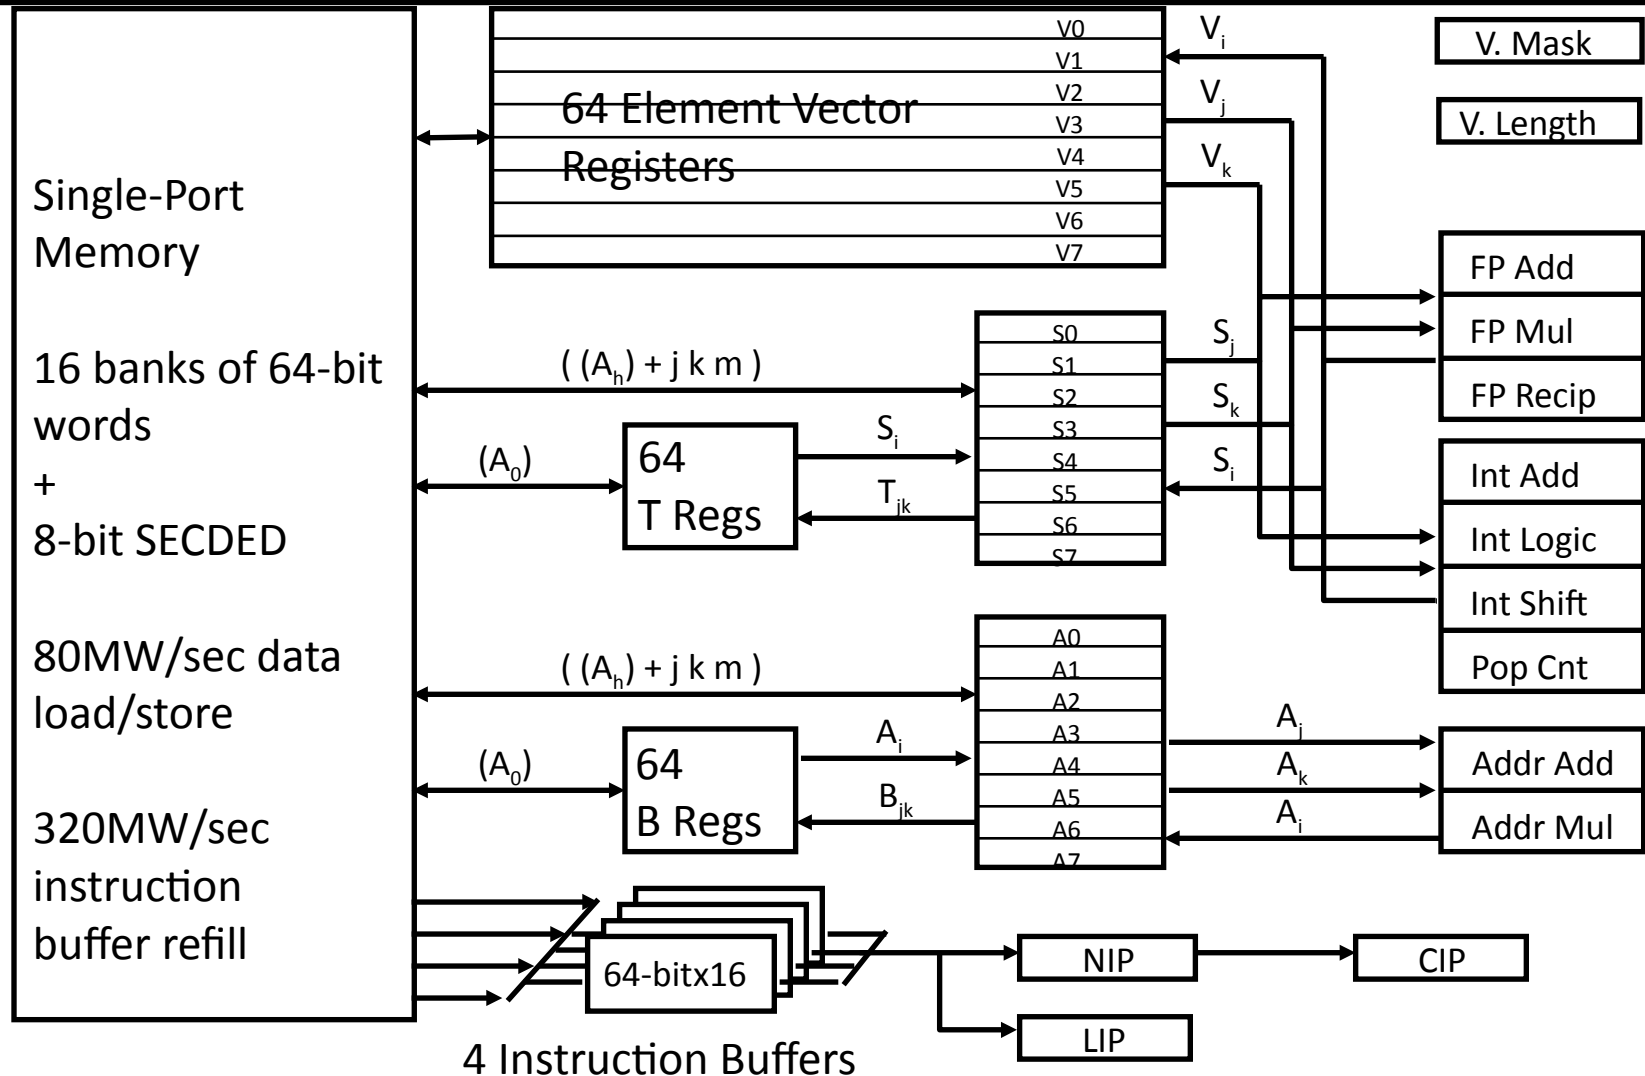
```
# Scalar Code
```

```
    li x4, 64  
loop:  
    fld f1, 0(x1)  
    fld f2, 0(x2)  
    fadd.d f3, f1, f2  
    fsd f3, 0(x3)  
    addi x1, 8  
    addi x2, 8  
    addi x3, 8  
    subi x4, 1  
    bnez x4, loop
```

```
# Vector Code
```

```
li x4, 64  
setvl x4  
vld v1, x1  
vld v2, x2  
vadd v3, v1, v2  
vst v3, x3
```

# Cray-1 (1976)



*memory bank cycle 50 ns    processor cycle 12.5 ns (80MHz)*




# Vector Instruction Set Advantages

---

- Compact
  - one short instruction encodes N operations
- Expressive, tells hardware that these N operations:
  - are independent
  - use the same functional unit
  - access disjoint registers
  - access registers in same pattern as previous instructions
  - access a contiguous block of memory (unit-stride load/store)
  - access memory in a known pattern (strided load/store)
- Scalable
  - can run same code on more parallel pipelines (lanes)

# Vector Arithmetic Execution

- Use deep pipeline ( $\Rightarrow$  fast clock) to execute element operations
- Simplifies control of deep pipeline because elements in vector are independent ( $\Rightarrow$  no hazards!)




$v3 \leftarrow v1 * v2$

# Vector Instruction Execution

vadd vc, va, vb





*Execution using  
one pipelined  
functional unit*

A[6] B[6]  
A[5] B[5]  
A[4] B[4]  
A[3] B[3]



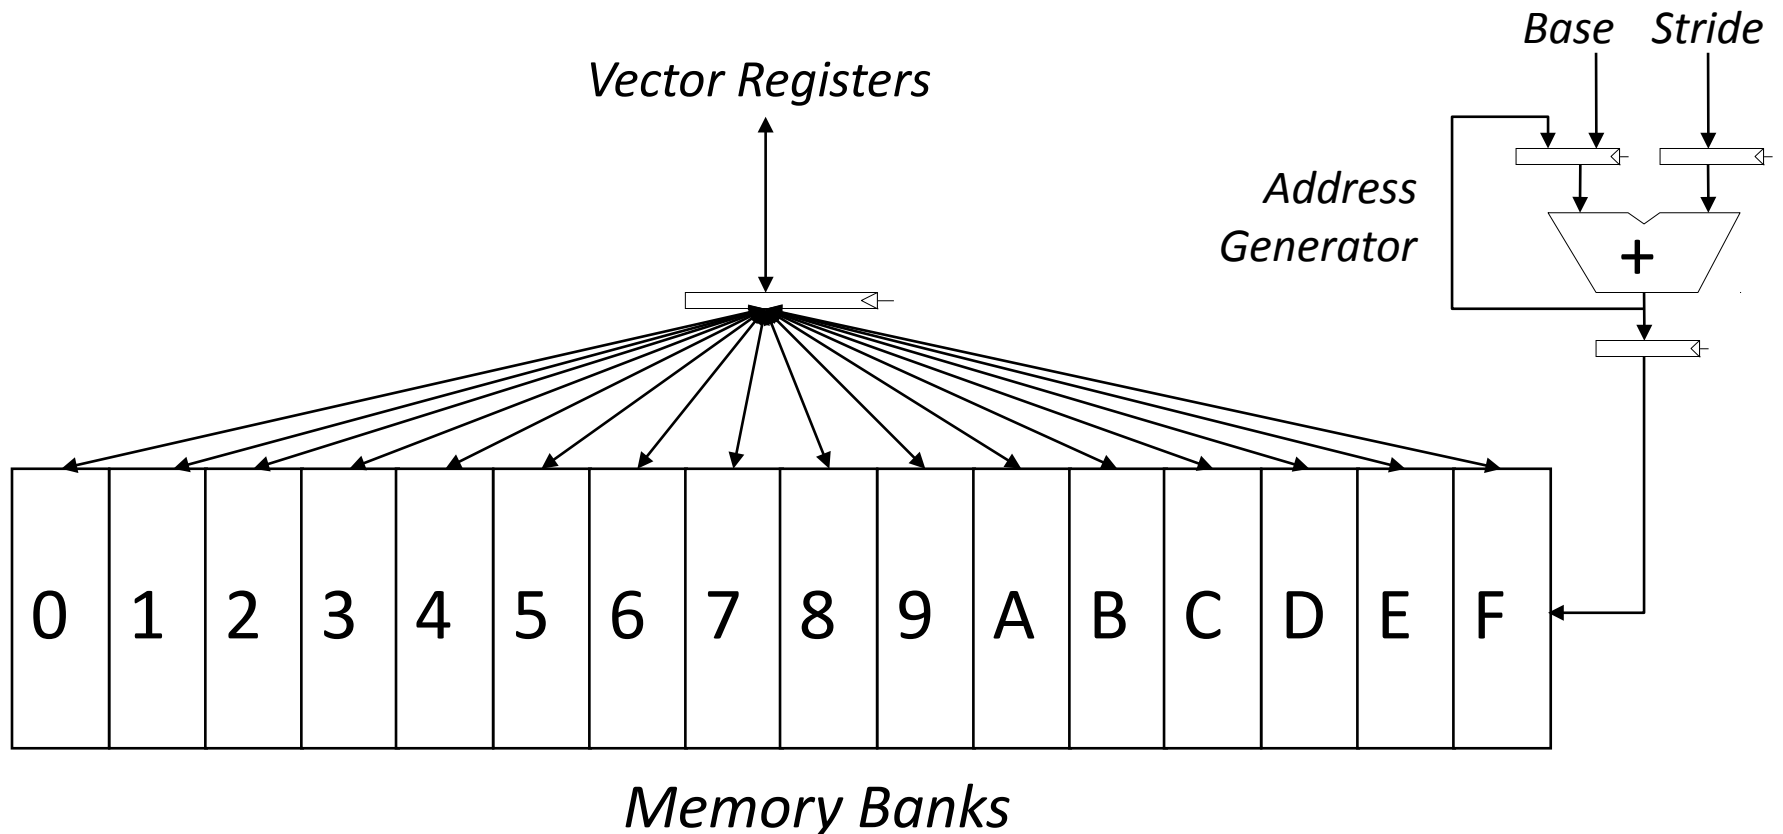
*Execution using four  
pipelined functional units*

A[24] B[24] A[25] B[25] A[26] B[26] A[27] B[27]  
A[20] B[20] A[21] B[21] A[22] B[22] A[23] B[23]  
A[16] B[16] A[17] B[17] A[18] B[18] A[19] B[19]  
A[12] B[12] A[13] B[13] A[14] B[14] A[15] B[15]




# Interleaved Vector Memory System

- Bank busy time: Time before bank ready to accept next request
- Cray-1, 16 banks, 4 cycle bank busy time, 12 cycle latency




# Vector Unit Structure



# T0 Vector Microprocessor (UCB/ICSI, 1995)

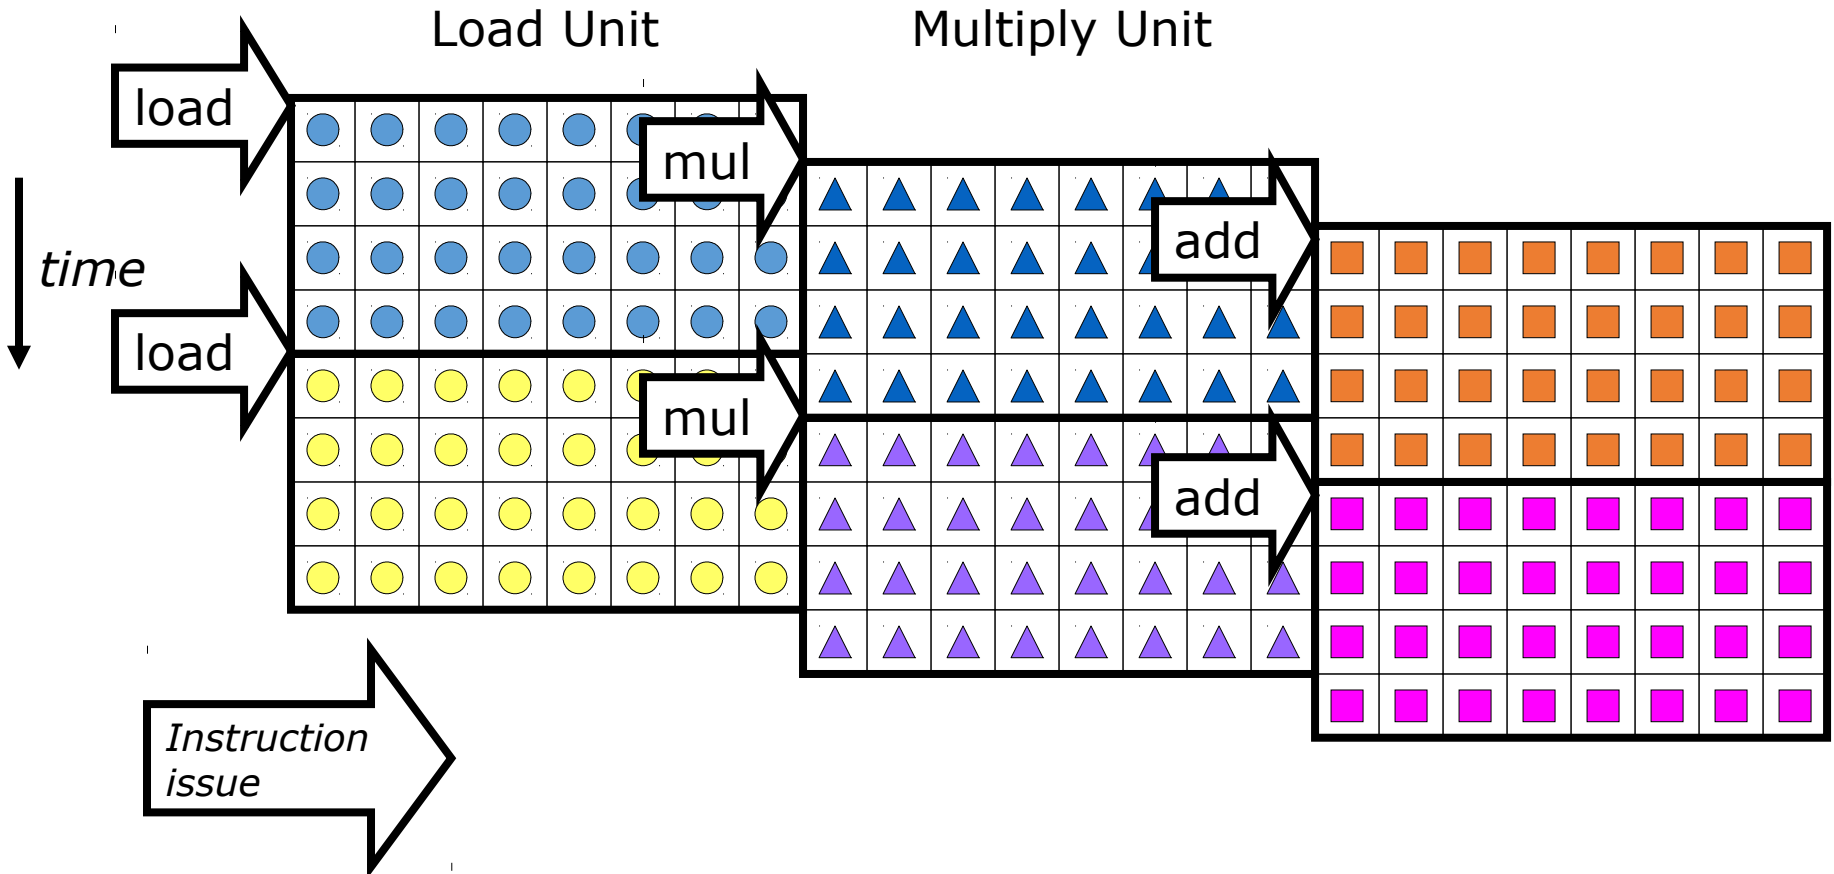
*Vector register  
elements striped  
over lanes*



*Lane*

# Vector Instruction Parallelism

- Can overlap execution of multiple vector instructions
  - example machine has 32 elements per vector register and 8 lanes




Complete 24 operations/cycle while issuing 1 short instruction/cycle

# Vector Chaining

- Vector version of register bypassing
  - introduced with Cray-1

```
vld  v1  
vmul v3, v1, v2  
vadd v5, v3, v4
```






# Vector Chaining Advantage

---

- Without chaining, must wait for last element of result to be written before starting dependent instruction



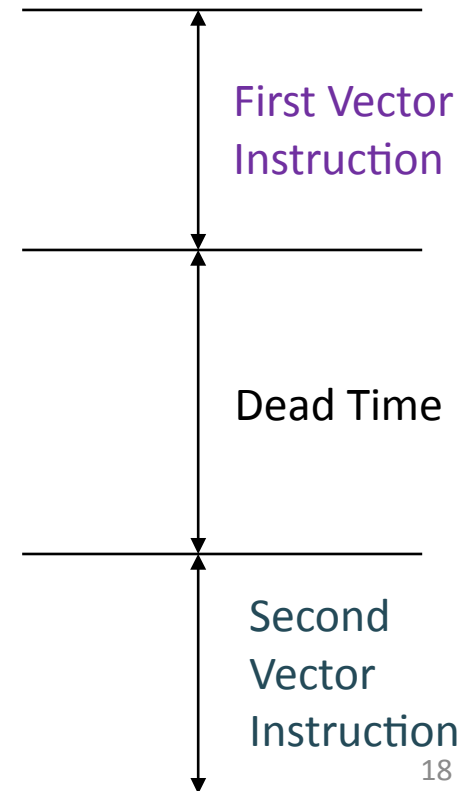
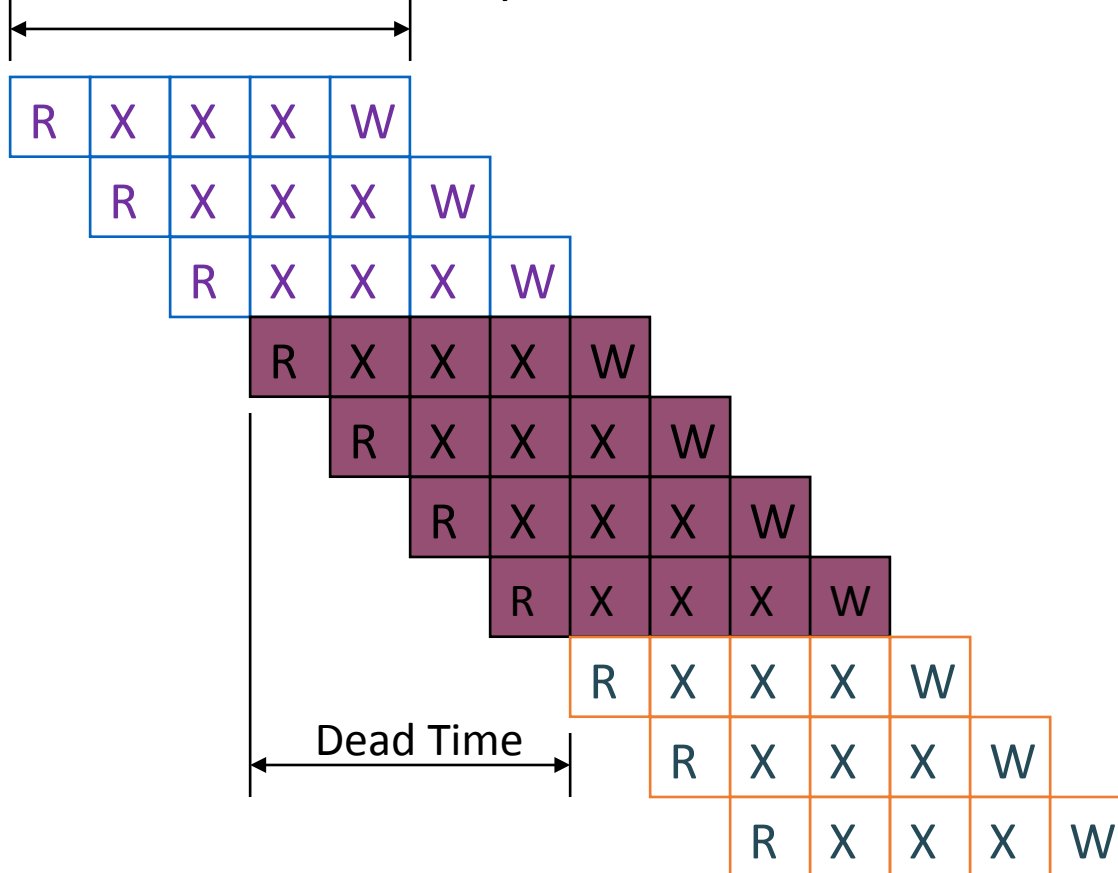
- With chaining, can start dependent instruction as soon as first result appears



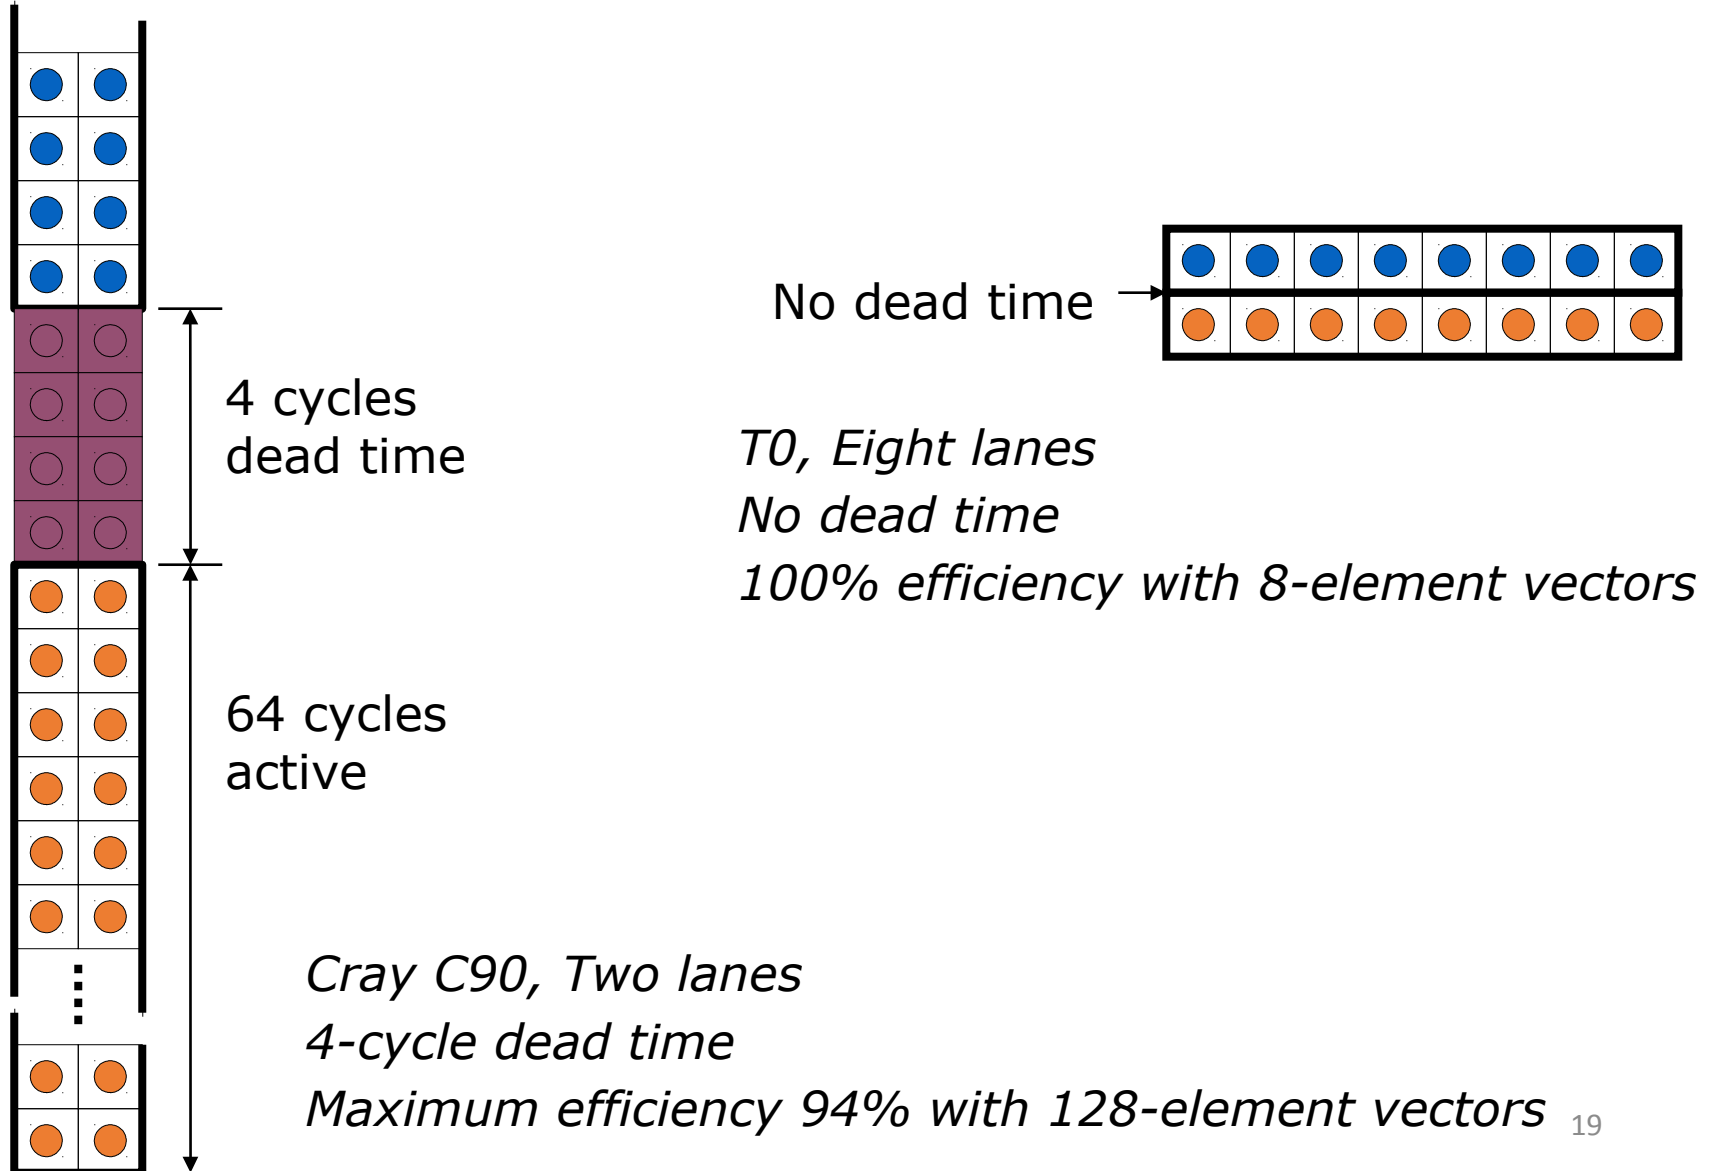
# Vector Startup

- Two components of vector startup penalty
  - functional unit latency (time through pipeline)
  - dead time or recovery time (time before another vector instruction can start down pipeline)

Functional Unit Latency



# Dead Time and Short Vectors



# Vector Memory-Memory versus Vector Register Machines

- Vector memory-memory instructions hold all vector operands in main memory
- The first vector machines, CDC Star-100 ('73) and TI ASC ('71), were memory-memory machines
- Cray-1 ('76) was first vector register machine

## Example Source Code

```
for (i=0; i<N; i++)  
{  
    C[i] = A[i] + B[i];  
    D[i] = A[i] - B[i];  
}
```

## Vector Memory-Memory Code

```
vadd C, A, B  
vsub D, A, B
```

## Vector Register Code

```
vld V1, A  
vld V2, B  
vadd V3, V1, V2  
vst V3, C  
vsub V4, V1, V2  
vst V4, D
```

# Vector Memory-Memory vs. Vector Register Machines

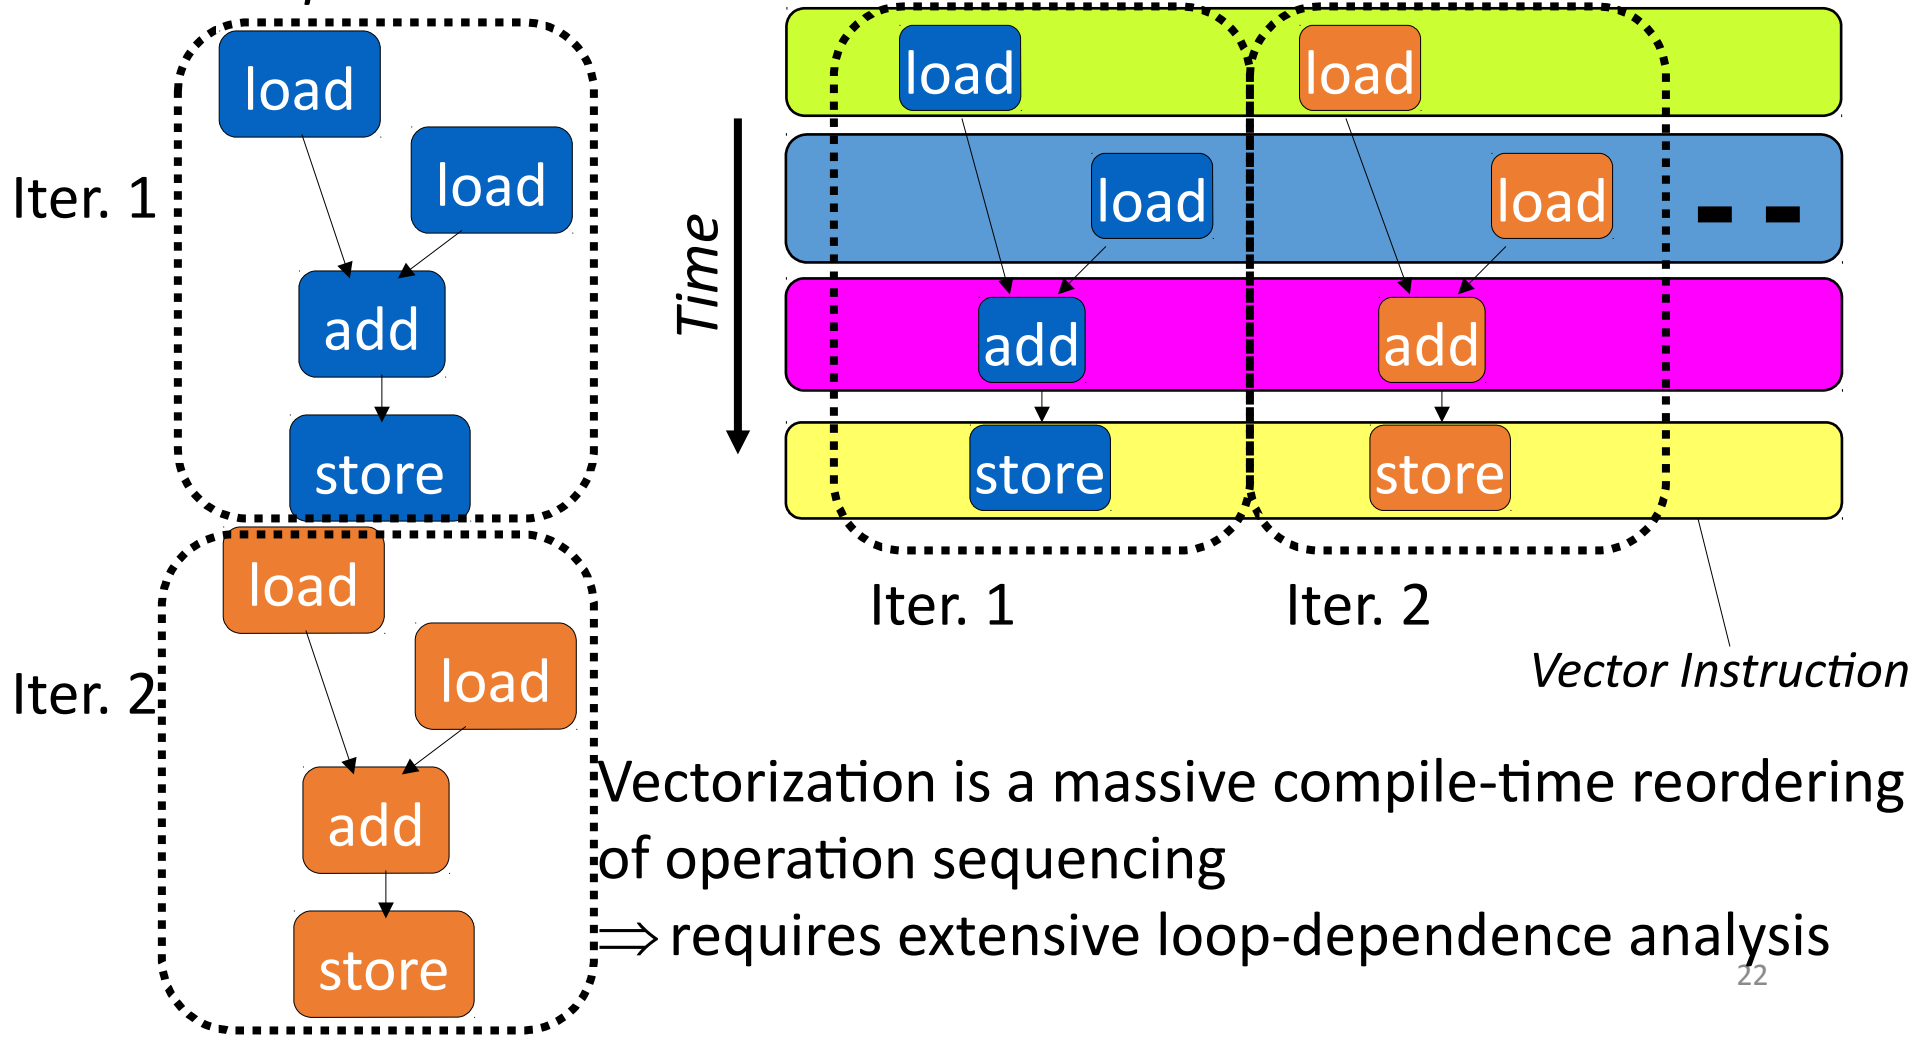
---

- Vector memory-memory architectures (VMMA) require greater main memory bandwidth, why?
  - All operands must be read in and out of memory
- VMMA make it difficult to overlap execution of multiple vector operations, why?
  - Must check dependencies on memory addresses
- VMMA incur greater startup latency
  - Scalar code was faster on CDC Star-100 for vectors < 100 elements
  - For Cray-1, vector/scalar breakeven point was around 2-4 elements
- Apart from CDC follow-ons (Cyber-205, ETA-10) all major vector machines since Cray-1 have had vector register architectures
  - (we ignore vector memory-memory from now on)

# Automatic Code Vectorization

```
for (i=0; i < N; i++)
```

*Scalar Sequential Code*    `C[i] = A[i] + B[i];`    *Vectorized Code*




# Vector Stripmining

**Problem:** Vector registers have finite length

**Solution:** Break loops into pieces that fit in registers, “*Stripmining*”

```
for (i=0; i<N; i++)  
    C[i] = A[i]+B[i];
```



```
    andi x1, xN, 63    # N mod 64  
    setv1 x1           # Do remainder  
loop:  
    vld v1, xA  
    sll x2, x1, 3      # Multiply by 8  
    add xA, x2         # Bump pointer  
    vld v2, xB  
    add xB, x2  
    vadd v3, v1, v2  
    vst v3, xC  
    add xC, x2  
    sub xN, x1         # Subtract elements  
    li x1, 64  
    setv1 x1           # Reset full length  
    bgtz xN, loop     # Any more to do?
```

# Vector Conditional Execution

---

Problem: Want to vectorize loops with conditional code:

```
for (i=0; i<N; i++)  
    if (A[i]>0) then  
        A[i] = B[i];
```

Solution: Add vector *mask* (or *flag*) registers

- vector version of predicate registers, 1 bit per element

...and *maskable* vector instructions

- vector operation becomes bubble (“NOP”) at elements where mask bit is clear

Code example:

```
cvm          # Turn on all elements  
vld vA, xA   # Load entire A vector  
vgt vA, f0   # Set bits in mask register where A>0  
vld vA, xB   # Load B vector into A under mask  
vst vA, xA   # Store A back to memory under mask
```



# Masked Vector Instructions

## Simple Implementation

- execute all N operations, turn off result writeback according to mask


M[7]=1 A[7] B[7]

M[6]=0 A[6] B[6]

M[5]=1 A[5] B[5]

M[4]=1 A[4] B[4]

M[3]=0 A[3] B[3]



M[0]=0

Write Enable

Write data port

## Density-Time Implementation

- scan mask vector and only execute elements with non-zero masks

M[7]=1

M[6]=0

M[5]=1

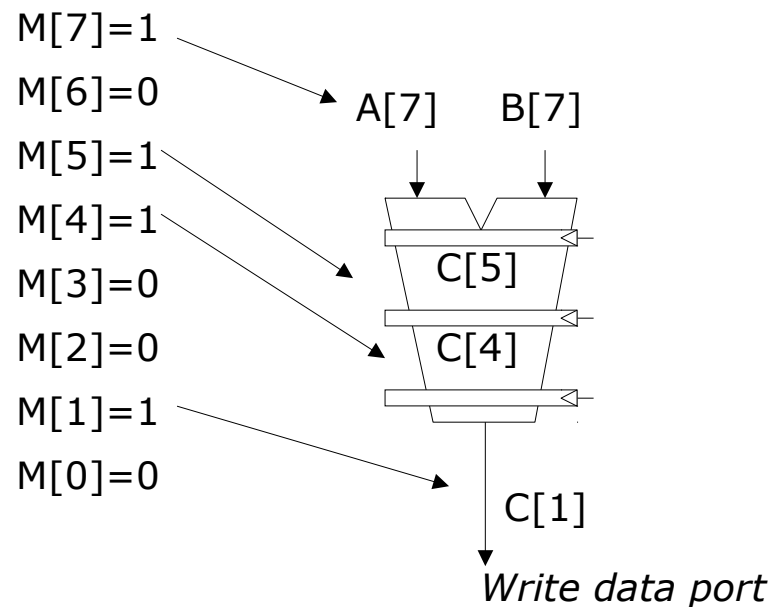
M[4]=1

M[3]=0

M[2]=0

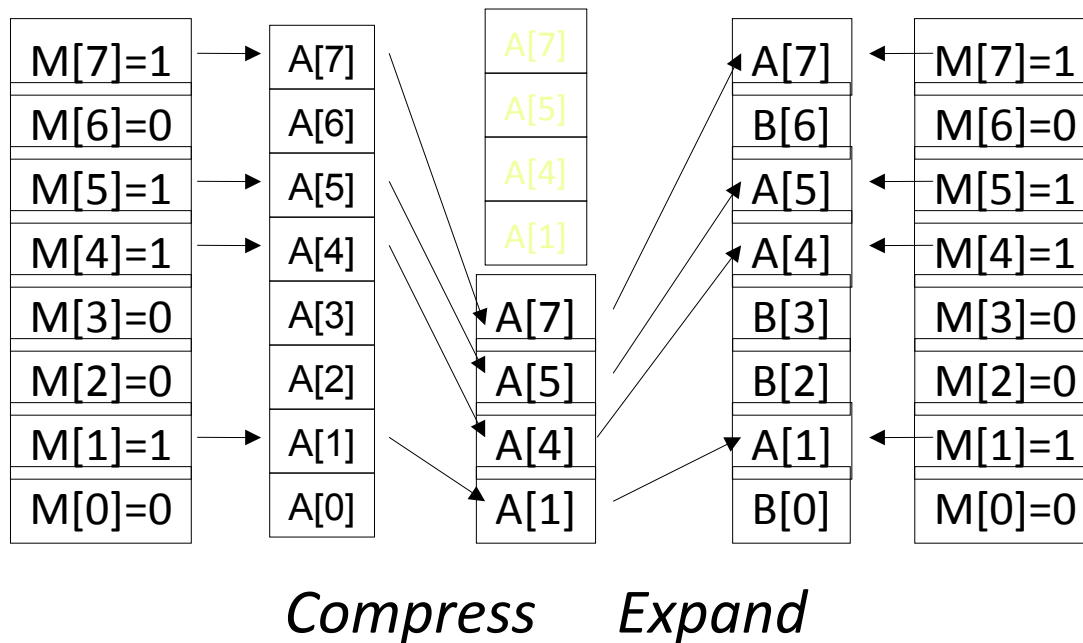
M[1]=1

M[0]=0



# Compress/Expand Operations

- Compress packs non-masked elements from one vector register contiguously at start of destination vector register
  - population count of mask vector gives packed vector length
- Expand performs inverse operation



Used for density-time conditionals and also for general selection operations

# Vector Reductions

---

**Problem:** Loop-carried dependence on reduction variables

```
sum = 0;
for (i=0; i<N; i++)
    sum += A[i];  # Loop-carried dependence on
sum
```

**Solution:** Re-associate operations if possible, use binary tree to perform reduction

```
# Rearrange as:
sum[0:VL-1] = 0          # Vector of VL partial sums
for(i=0; i<N; i+=VL)    # Stripmine VL-sized chunks
    sum[0:VL-1] += A[i:i+VL-1]; # Vector sum
# Now have VL partial sums in one vector register
do {
    VL = VL/2;          # Halve vector length
    sum[0:VL-1] += sum[VL:2*VL-1] # Halve no. of
    partials
} while (VL>1)
```

# Vector Scatter/Gather

---

Want to vectorize loops with indirect accesses:

```
for (i=0; i<N; i++)  
    A[i] = B[i] + C[D[i]]
```

Indexed load instruction (*Gather*)

```
vld vD, xD          # Load indices in D vector  
vldx vC, xC, vD      # Load indirect from xC base  
vld vB, xB           # Load B vector  
vadd vA, vB, vC      # Do add  
vst vA, xA           # Store result
```

# Vector Instruction Set Advantages

---

- Many vector machines have a very relaxed memory model, e.g.

```
vst v1, x1    # Store vector to x1  
vld v2, x1    # Load vector from x1
```

- No guarantee that elements of v2 will have value of elements of v1 even when store and load execute by same processor!


- Requires explicit memory barrier or fence

```
vst v1, x1    # Store vector to x1  
fence         # Enforce ordering s->l  
vld v2, x1    # Load vector from x1
```


- Vector machines support highly parallel memory systems (multiple lanes and multiple load and store units) with long latency (100+ clock cycles)

- hardware coherence checks would be prohibitively expensive
- vectorizing compiler can eliminate most dependencies

# Packed SIMD Extensions



- Very short vectors added to existing ISAs for microprocessors
- Use existing 64-bit registers split into 2x32b or 4x16b or 8x8b
  - Lincoln Labs TX-2 from 1957 had 36b datapath split into 2x18b or 4x9b
  - Newer designs have wider registers
    - 128b for PowerPC AltiVec, Intel SSE2/3/4
    - 256b/512b for Intel AVX
- Single instruction operates on all elements within register



# Packed SIMD versus Vectors

---

- Limited instruction set:
  - no vector length control
  - no strided load/store or scatter/gather
  - unit-stride loads must be aligned to 64/128-bit boundary
- Limited vector register length:
  - requires superscalar dispatch to keep multiply/add/load units busy
  - loop unrolling to hide latencies increases register pressure
- Trend towards fuller vector support in microprocessors
  - Better support for misaligned memory accesses
  - Support of double-precision (64-bit floating-point)
  - New Intel AVX spec (announced April 2008), 256b vector registers (expandable up to 1024b), gather added, scatter to follow
  - ARM Scalable Vector Extensions (SVE)

# New RISC-V “V” Vector Extension

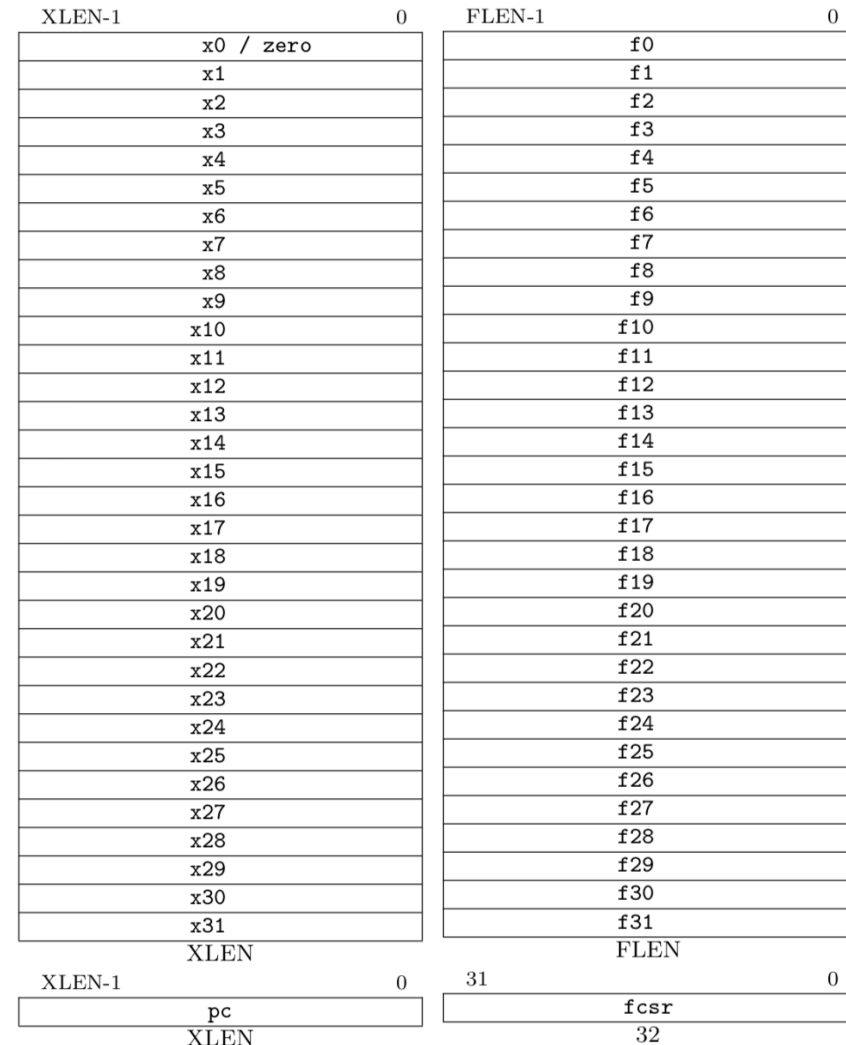
---

- Being added as a standard extension to the RISC-V ISA
  - An updated form of Cray-style vectors for modern microprocessors
- Today, a short tutorial on current draft standard, v0.7
  - v0.7 is intended to be close to final version of RISC-V vector extension
  - Still a work in progress, so details might change before standardization
  - <https://github.com/riscv/riscv-v-spec>



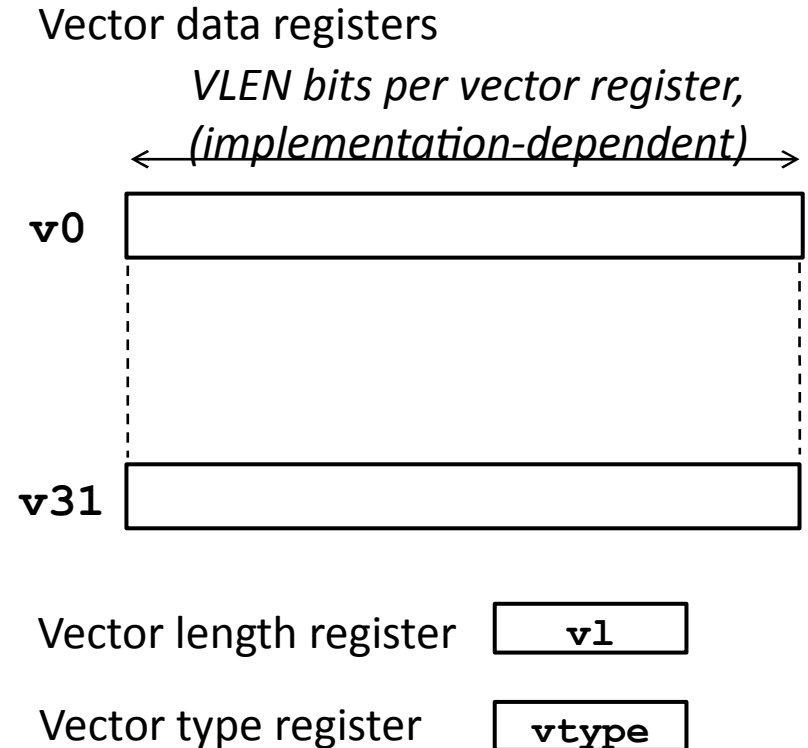
# RISC-V Scalar State

- Program counter (**pc**)
- 32x32/64-bit integer registers (**x0–x31**)
  - **x0** always contains a 0
- Floating-point (FP), adds 32 registers (**f0–f31**)
  - each can contain a single- or double-precision FP value (32-bit or 64-bit IEEE FP)
- FP status register (**fcsr**), used for FP rounding mode & exception reporting
- ISA string options:
  - RV32I (XLEN=32, no FP)
  - RV32IF (XLEN=32, FLEN=32)
  - RV32ID (XLEN=32, FLEN=64)
  - RV64I (XLEN=64, no FP)
  - RV64IF (XLEN=64, FLEN=32)
  - RV64ID (XLEN=64, FLEN=64)




# Vector Extension Additional State

- 32 vector data registers, **v0–v31**, each VLEN bits long
- Vector length register **v1**
- Vector type register **vtype**
- Other control registers:
  - **vstart**
  - For trap handling
  - **vrn/vxsat**
  - Fixed-point rounding mode/saturation
  - Also appear in **fcsr**



# Vector Type Register

vtype register layout



Bits	Contents
1:0	vlmul[1:0]
4:2	vsew[2:0]
6:5	vediv[1:0]
XLEN-1:7	Reserved (write 0)

vsew[2:0]			SEW
0	0	0	8
0	0	1	16
0	1	0	32
0	1	1	64
1	0	0	128
1	0	1	256
1	1	0	512
1	1	1	1024


**vsew[2:0]** field encodes standard element width (SEW) in bits of elements in vector register ( $SEW = 8 * 2^{vsew}$ )

**vlmul[1:0]** encodes vector register length multiplier ( $LMUL = 2^{vlmul} = 1-8$ )


**vediv[1:0]** encodes how vector elements are divided into equal sub-elements ( $EDIV = 2^{vediv} = 1-8$ )

# Example Vector Register Data Layouts (LMUL=1)


VLEN=32b



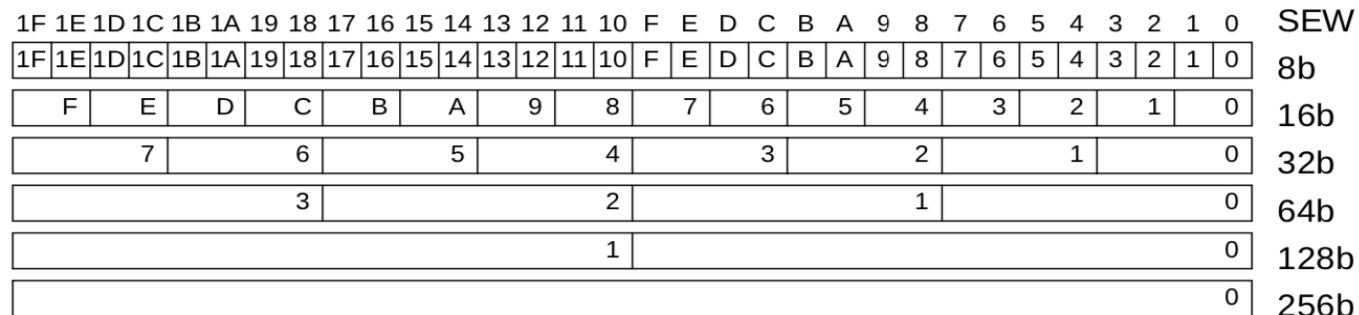
VLEN=64b



VLEN=128b



VLEN = 256b



# Setting vector configuration, **vsetvli/vsetvl**


The **vsetvli{i}** configuration instructions set the **vtype** register, and also set the **vl** register, returning the **vl** value in a scalar register

**vsetvli rd, rs1, e8 # Set SEW=8, vl=min(VLEN/SEW, rs1), rd=vl**

Resulting machine vector length setting      Requested application vector length

**vtype** parameters (SEW,LMUL,EDIV) encoded as immediate in instruction

Instruction encoding



Usually use immediate form, **vsetvli**, to set **vtype** parameters.  
The register version **vsetvl** is usually used for context save/restore

## `vsetvl{i}` operation

---

- The first scalar register argument, `rs1`, is the requested application vector length (AVL)
- The type argument (either immediate or second register) indicates how the vector registers should be configured
  - Configuration includes size of each element
- The vector length is set to the minimum of requested AVL and the maximum supported vector length (VLMAX) in the new configuration
  - $\text{VLMAX} = \text{LMUL} * \text{VLEN} / \text{SEW}$
  - $\text{vl} = \min(\text{AVL}, \text{VLMAX})$
- The value placed in `vl` is also written to the scalar destination register `rd`

# Simple stripmined vector memcpy example

```
# void *memcpy(void* dest, const void* src, size_t n)
# a0=dest, a1=src, a2=n
#
```

Set configuration,  
calculate vector strip  
length

Unit-stride  
vector load  
bytes


Unit-stride  
vector store  
bytes

```
memcpy:
    mv a3, a0 # Copy destination
loop:
    vsetvli t0, a2, e8          # Vectors of 8b
    vlb.v v0, (a1)             # Load bytes
    add a1, a1, t0              # Bump pointer
    sub a2, a2, t0              # Decrement count
    vsb.v v0, (a3)              # Store bytes
    add a3, a3, t0              # Bump pointer
    bnez a2, loop               # Any more?
    ret                          # Return
```

Binary machine code can run on machines with any VLEN!


# Vector Load Instructions

## unit-stride




## strided

*Scalar stride (bytes)*



## indexed

*Vector of offsets (bytes)*




*Scalar base address*




# Vector Store Instructions


unit-stride



strided



indexed



*Vector store data*

# Vector Unit-Stride Loads/Stores

---

```
# vd destination, rs1 base address, vm is mask encoding (v0.t or <missing>)
vlb.v          vd, (rs1), vm      # 8b signed
vlh.v          vd, (rs1), vm      # 16b signed
vlw.v          vd, (rs1), vm      # 32b signed

vlbu.v         vd, (rs1), vm      # 8b unsigned
vlhu.v         vd, (rs1), vm      # 16b unsigned
vlwu.v         vd, (rs1), vm      # 32b unsigned

vle.v          vd, (rs1), vm      # SEW

# vs3 store data, rs1 base address, vm is mask encoding (v0.t or <missing>)
vsb.v          vs3, (rs1), vm     # 8b store
vsh.v          vs3, (rs1), vm     # 16b store
vsw.v          vs3, (rs1), vm     # 32b store
vse.v          vs3, (rs1), vm     # SEW store
```

# Vector Strided Load/Store Instructions

---

# vd destination, rs1 base address, rs2 byte stride

**vlsb.v**            **vd**, (**rs1**), **rs2**, **vm** # 8b

**vlsh.v**            **vd**, (**rs1**), **rs2**, **vm** # 16b

**vlsw.v**            **vd**, (**rs1**), **rs2**, **vm** # 32b

**vlsbu.v**           **vd**, (**rs1**), **rs2**, **vm** # unsigned 8b

**vlshu.v**           **vd**, (**rs1**), **rs2**, **vm** # unsigned 16b

**vlswu.v**           **vd**, (**rs1**), **rs2**, **vm** # unsigned 32b

**vlse.v**            **vd**, (**rs1**), **rs2**, **vm** # SEW

# vs3 store data, rs1 base address, rs2 byte stride

**vssb.v**            **vs3**, (**rs1**), **rs2**, **vm** # 8b

**vssh.v**            **vs3**, (**rs1**), **rs2**, **vm** # 16b

**vssw.v**            **vs3**, (**rs1**), **rs2**, **vm** # 32b

**vsse.v**            **vs3**, (**rs1**), **rs2**, **vm** # SEW

# Vector Indexed Loads/Stores

---

```
# vd destination, rs1 base address, vs2 indices
vldb.v      vd, (rs1), vs2, vm # 8b
vlxh.v      vd, (rs1), vs2, vm # 16b
vlxw.v      vd, (rs1), vs2, vm # 32b

vlxbu.v     vd, (rs1), vs2, vm # 8b unsigned
vlxhu.v     vd, (rs1), vs2, vm # 16b unsigned
vlxwu.v     vd, (rs1), vs2, vm # 32b unsigned

vlxe.v      vd, (rs1), vs2, vm # SEW

# Vector ordered-indexed store instructions
# vs3 store data, rs1 base address, vs2 indices
vsxb.v      vs3, (rs1), vs2, vm # 8b
vsxh.v      vs3, (rs1), vs2, vm # 16b
vsxw.v      vs3, (rs1), vs2, vm # 32b
vsxe.v      vs3, (rs1), vs2, vm # SEW

# Vector unordered-indexed store instructions
vsuxb.v     vs3, (rs1), vs2, vm # 8b
vsuxh.v     vs3, (rs1), vs2, vm # 16b
vsuxw.v     vs3, (rs1), vs2, vm # 32b
vsuxe.v     vs3, (rs1), vs2, vm # SEW
```

# Vector Length Multiplier, LMUL

- Gives fewer but longer vector registers
- Set by `vlmul[1:0]` field in `vtype` during `setvli`

LMUL=2

F	E	D	C	B	A	9	8	7	6	5	4	3	2	1	0	Byte
			3				2				1				0	$v2 * n + 0$
			7				6				5				4	$v2 * n + 1$

LMUL=4

F	E	D	C	B	A	9	8	7	6	5	4	3	2	1	0	Byte
			9				8				1				0	$v4 * n + 0$
			B				A				3				2	$v4 * n + 1$
			D				C				5				4	$v4 * n + 2$
			F				E				7				6	$v4 * n + 3$

# LMUL=8 stripmined vector memcpy example

```
# void *memcpy(void* dest, const void* src, size_t n)
# a0=dest, a1=src, a2=n
#
memcpy:
    mv a3, a0 # Copy destination
loop:
    vsetvli t0, a2, e8,m8 # Vectors of 8b
    vlb.v v0, (a1)        # Load bytes
    add a1, a1, t0         # Bump pointer
    sub a2, a2, t0         # Decrement count
    vsb.v v0, (a3)         # Store bytes
    add a3, a3, t0         # Bump pointer
    bnez a2, loop
    ret                  # Any more?
                        # Return
```

Set configuration, calculate vector strip length

Unit-stride vector load bytes

Unit-stride vector store bytes

Combine eight vector registers into group (v0,v1,...,v7)

Binary machine code can run on machines with any VLEN!

# Mixed-Width Loops

---

- Have different element widths in one loop, even in one instruction
- Want same number of elements in each vector register, even if different bits/element
- Solution: Keep SEW/LMUL constant

VLEN=256b, SLEN=128b

SEW=8b, LMUL=1, VLMAX=32

1F	1E	1D	1C	1B	1A	19	18	17	16	15	14	13	12	11	10	F	E	D	C	B	A	9	8	7	6	5	4	3	2	1	0	Byte
1F	1E	1D	1C	1B	1A	19	18	17	16	15	14	13	12	11	10	F	E	D	C	B	A	9	8	7	6	5	4	3	2	1	0	v1 * n + 0

SEW=16b, LMUL=2, VLMAX=32

1F	1E	1D	1C	1B	1A	19	18	17	16	15	14	13	12	11	10	F	E	D	C	B	A	9	8	7	6	5	4	3	2	1	0	Byte
17	16	15	14	13	12	11	10	7	6	5	4	3	2	1	0																	v2 * n + 0
1F	1E	1D	1C	1B	1A	19	18	F	E	D	C	B	A	9	8																	v2 * n + 1

SEW=32b, LMUL=4, VLMAX=32

1F	1E	1D	1C	1B	1A	19	18	17	16	15	14	13	12	11	10	F	E	D	C	B	A	9	8	7	6	5	4	3	2	1	0	Byte
13						12									10					3							1				0	v4 * n + 0
17						16									14					7							5				4	v4 * n + 1
1B						1A									18					B							9				8	v4 * n + 2
1F						1E									1C					F							D				C	v4 * n + 3

SEW=64b, LMUL=8, VLMAX=32

1F	1E	1D	1C	1B	1A	19	18	17	16	15	14	13	12	11	10	F	E	D	C	B	A	9	8	7	6	5	4	3	2	1	0	Byte
						11									10																0	v8 * n + 0
						13									12																2	v8 * n + 1
						15									14																4	v8 * n + 2
						17									16																6	v8 * n + 3
						19									18																8	v8 * n + 4
						1B									1A																A	v8 * n + 5
						1D									1C																C	v8 * n + 6
						1F									1E																E	v8 * n + 7