

# 计算机组成与系统结构

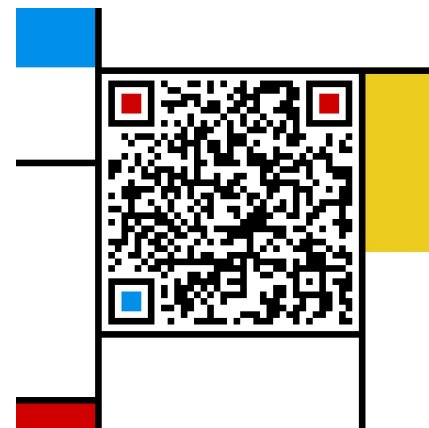
## Computer Organization & System Architecture

Huang Kejie (黄科杰) 百人计划研究员

Office: 玉泉校区老生仪楼 304

Email address: [huangkejie@zju.edu.cn](mailto:huangkejie@zju.edu.cn)

HP: 17706443800

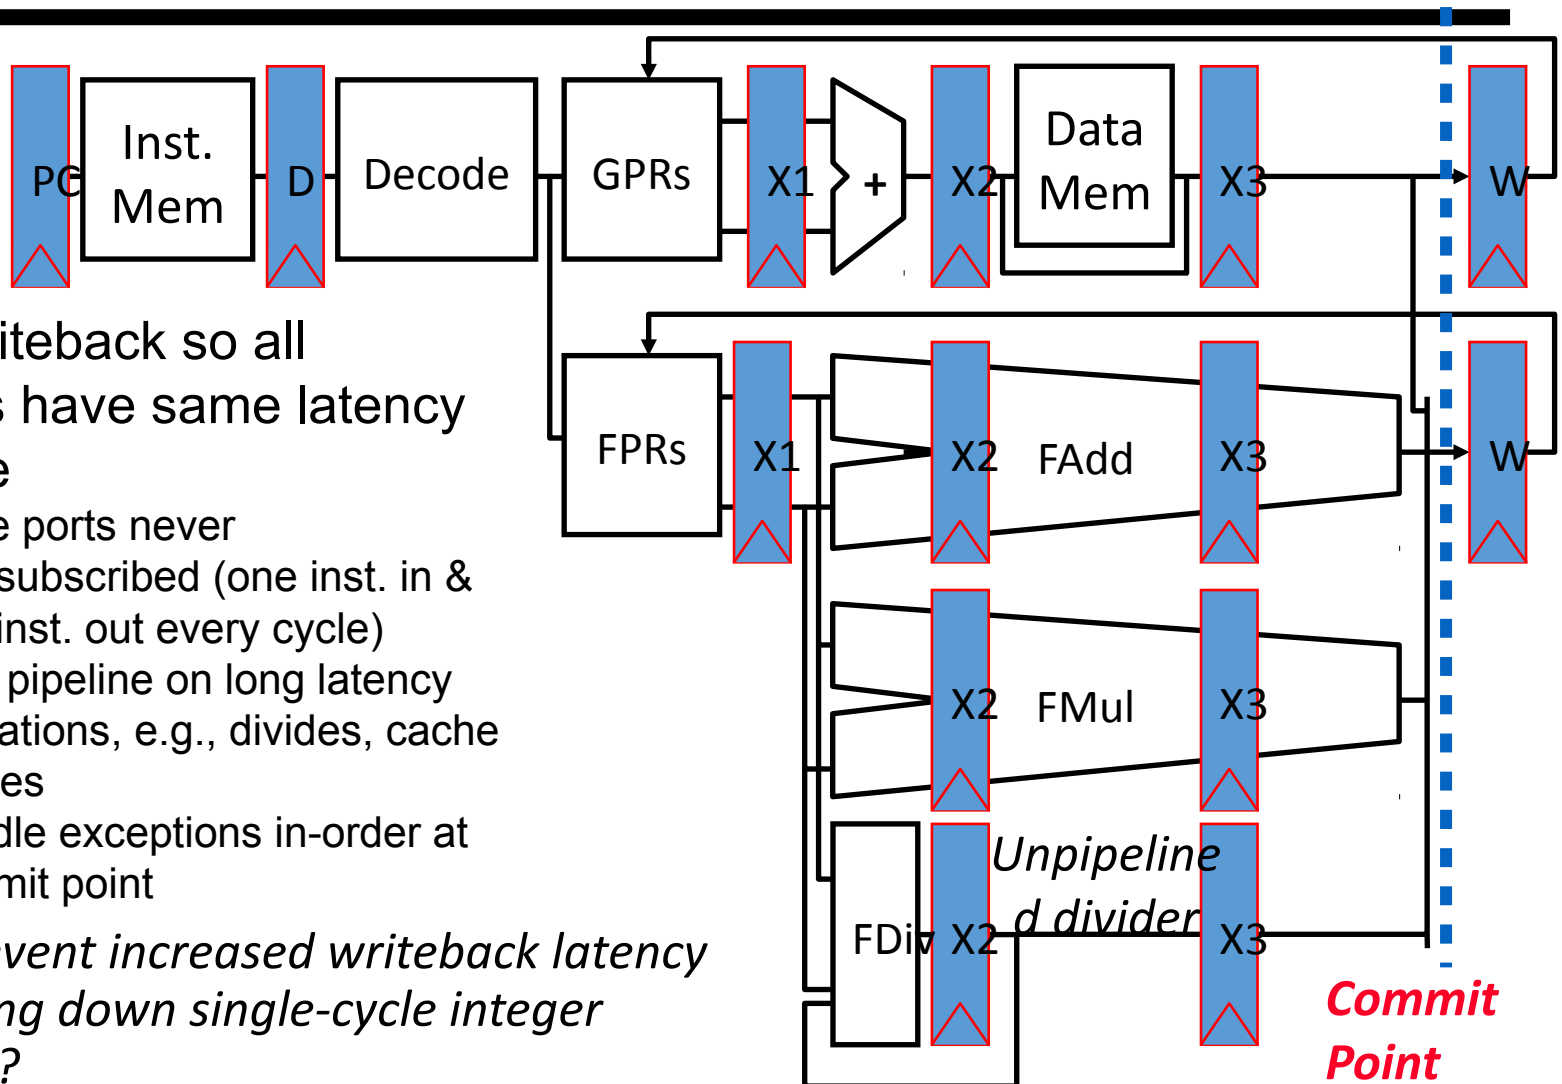


# Recap: Complex In-Order Pipeline

- Delay writeback so all operations have same latency to W stage
  - Write ports never oversubscribed (one inst. in & one inst. out every cycle)
  - Stall pipeline on long latency operations, e.g., divides, cache misses
  - Handle exceptions in-order at commit point

*How to prevent increased writeback latency from slowing down single-cycle integer operations?*


**Bypassing**



# Scoreboard Dynamics

	Functional Unit Status								Registers Reserved for Writes	
	Int(1)	Add(1)	Mult(3)			Div(4)		WB		
t0	$I_1$		f6				f6			
t1	$I_2$	f2		f6			f6	f2		
t2			f6		f2		f6	f2		$I_2$
t3	$I_3$	f0				f6		f6, f0		
t4		f0				f6	f6	f0		$I_1$
t5	$I_4$		f0	f8				f0, f8		
t6			f8			f0	f0	f8		$I_3$
t7	$I_5$	f10			f8		f8	f10		
t8			f8	f10		f8	f10		$I_5$	
t9				f8		f8		$I_4$		
t10	$I_6$	f6					f6			
t11	$I_1$	FDIV.D	f6, f6		f4	f6			$I_6$	
	$I_2$	FLD	f2, 45(x3)							
	$I_3$	FMULT.D	f0, f2, f4							
	$I_4$	FDIV.D	f8, f6, f2							
	$I_5$	FSUB.D	f10, f0, f6							
	$I_6$	FADD.D	f6, f0, f2							

# Reorder Buffer Management



## ROB managed circularly

- “exec” bit is set when instruction begins execution
- When an instruction completes its “use” bit is marked free
- $ptr_2$  is incremented only if the “use” bit is marked free

## Instruction slot is candidate for execution when:

- It holds a valid instruction (“use” bit is set)
- It has not already started execution (“exec” bit is clear)
- Both operands are available ( $p1$  and  $p2$  are set)

# Renaming & Out-of-order Issue

## An example

Renaming table

	p	data
f1		
f2		v1
f3		
f4		t5
f5		
f6		t3
f7		
f8		v4

v1

data /  $t_i$

Reorder buffer

Ins#	use	exec	op	p1	src1	p2	src2
1	0	0	LD				
2	0	0	LD				
3	1	0	MUL	0	v2	1	v1
4	0	0	SUB	1	v1	1	v1
5	1	0	DIV	1	v1	0	v4

$t_1$

$t_2$

$t_3$

$t_4$

$t_5$

.

.

1 FLD	f2, 34(x2)
2 FLD	f4, 45(x3)
3 FMULT.D	f6, f4, f2
4 FSUB.D	f8, f2, f2
5 FDIV.D	f4, f2, f8
6 FADD.D	f10, f6, f4

- When are tags in sources replaced by data?  
*Whenever an FU produces data*
- When can a name be reused?  
*Whenever an instruction completes*

# Physical Register Management

Rename Table		Physical Regs		Free List
x0		P0		P0
x1	P8	P1		P1
x2		P2		P3
x3	P7	P3		P2
x4		P4		P4
x5		P5	<x6> p	
x6	P5	P6	<x7> p	
x7	P6	P7	<x3> p	
		P8	<x1> p	
		Pn		

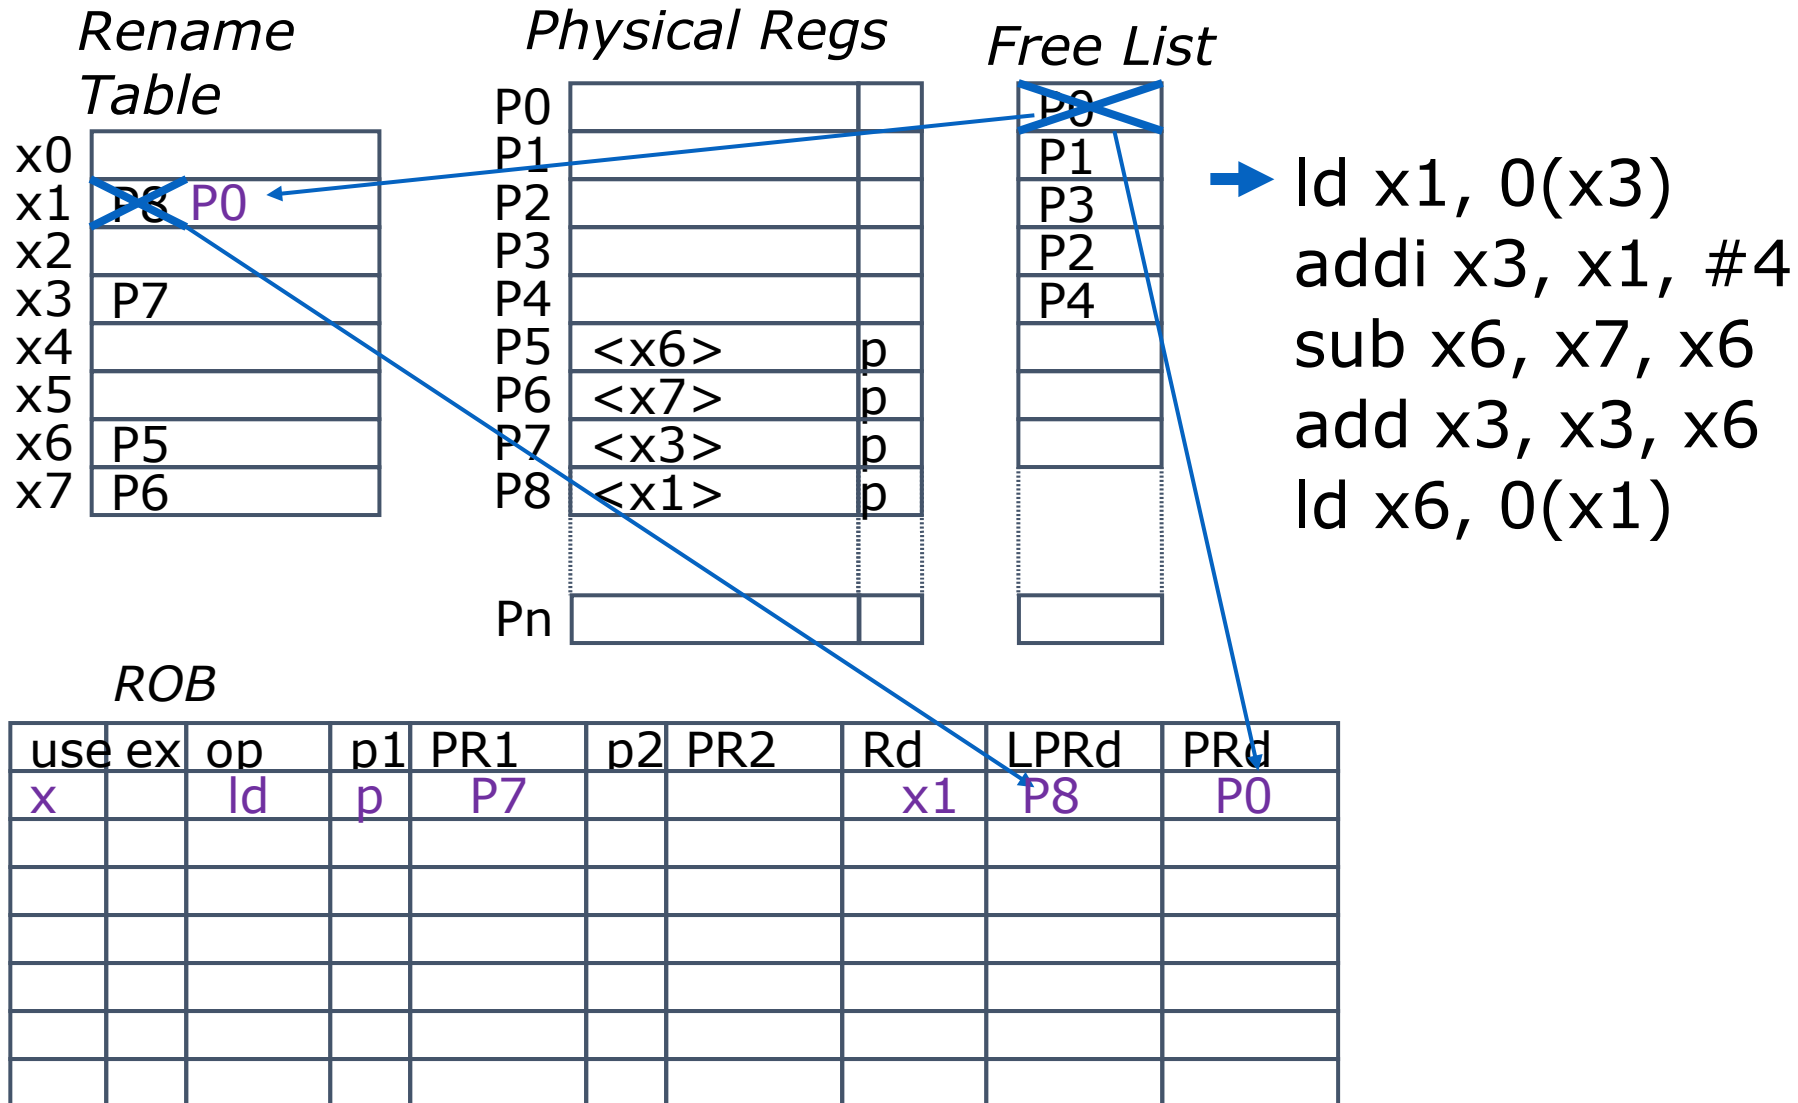
ld x1, 0(x3)  
 addi x3, x1, #4  
 sub x6, x7, x6  
 add x3, x3, x6  
 ld x6, 0(x1)

ROB

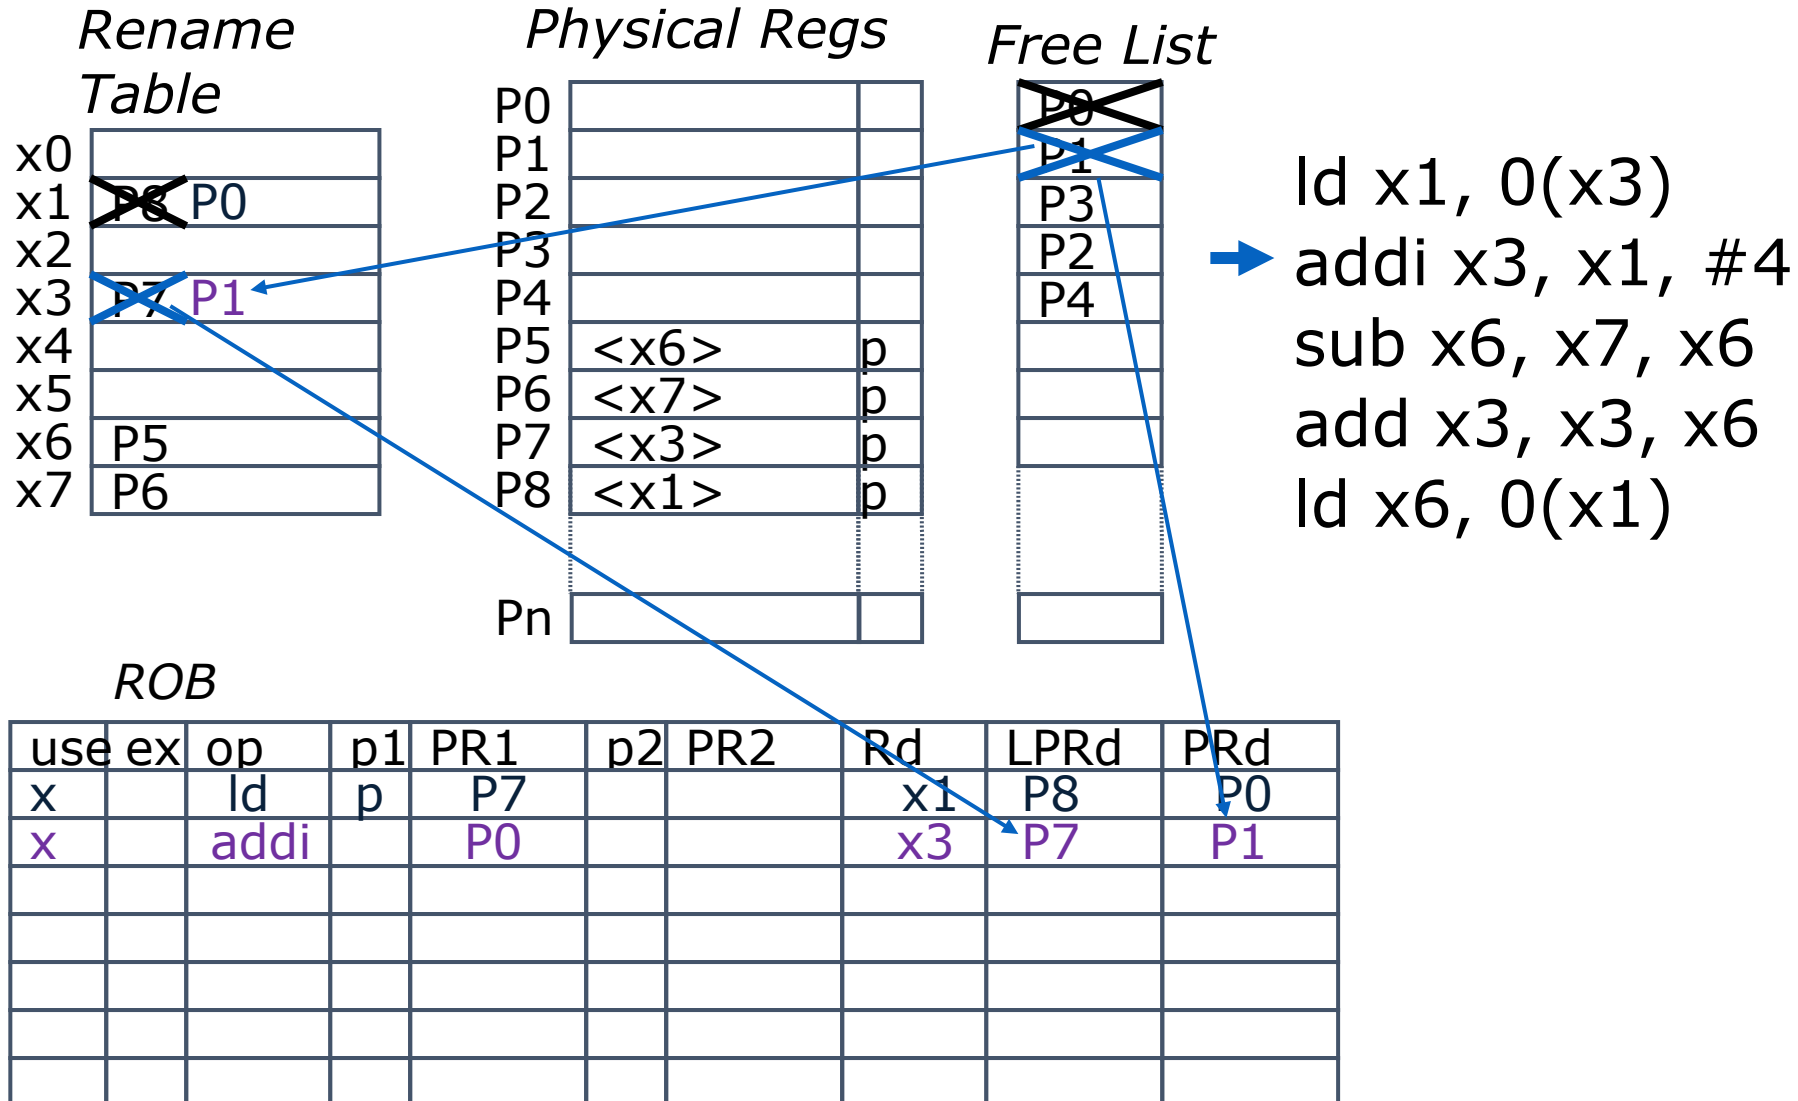
use	ex	op	p1	PR1	p2	PR2	Rd	LPRd	PRd

(LPRd requires third read port on Rename Table for each instruction)

# Physical Register Management




# Physical Register Management






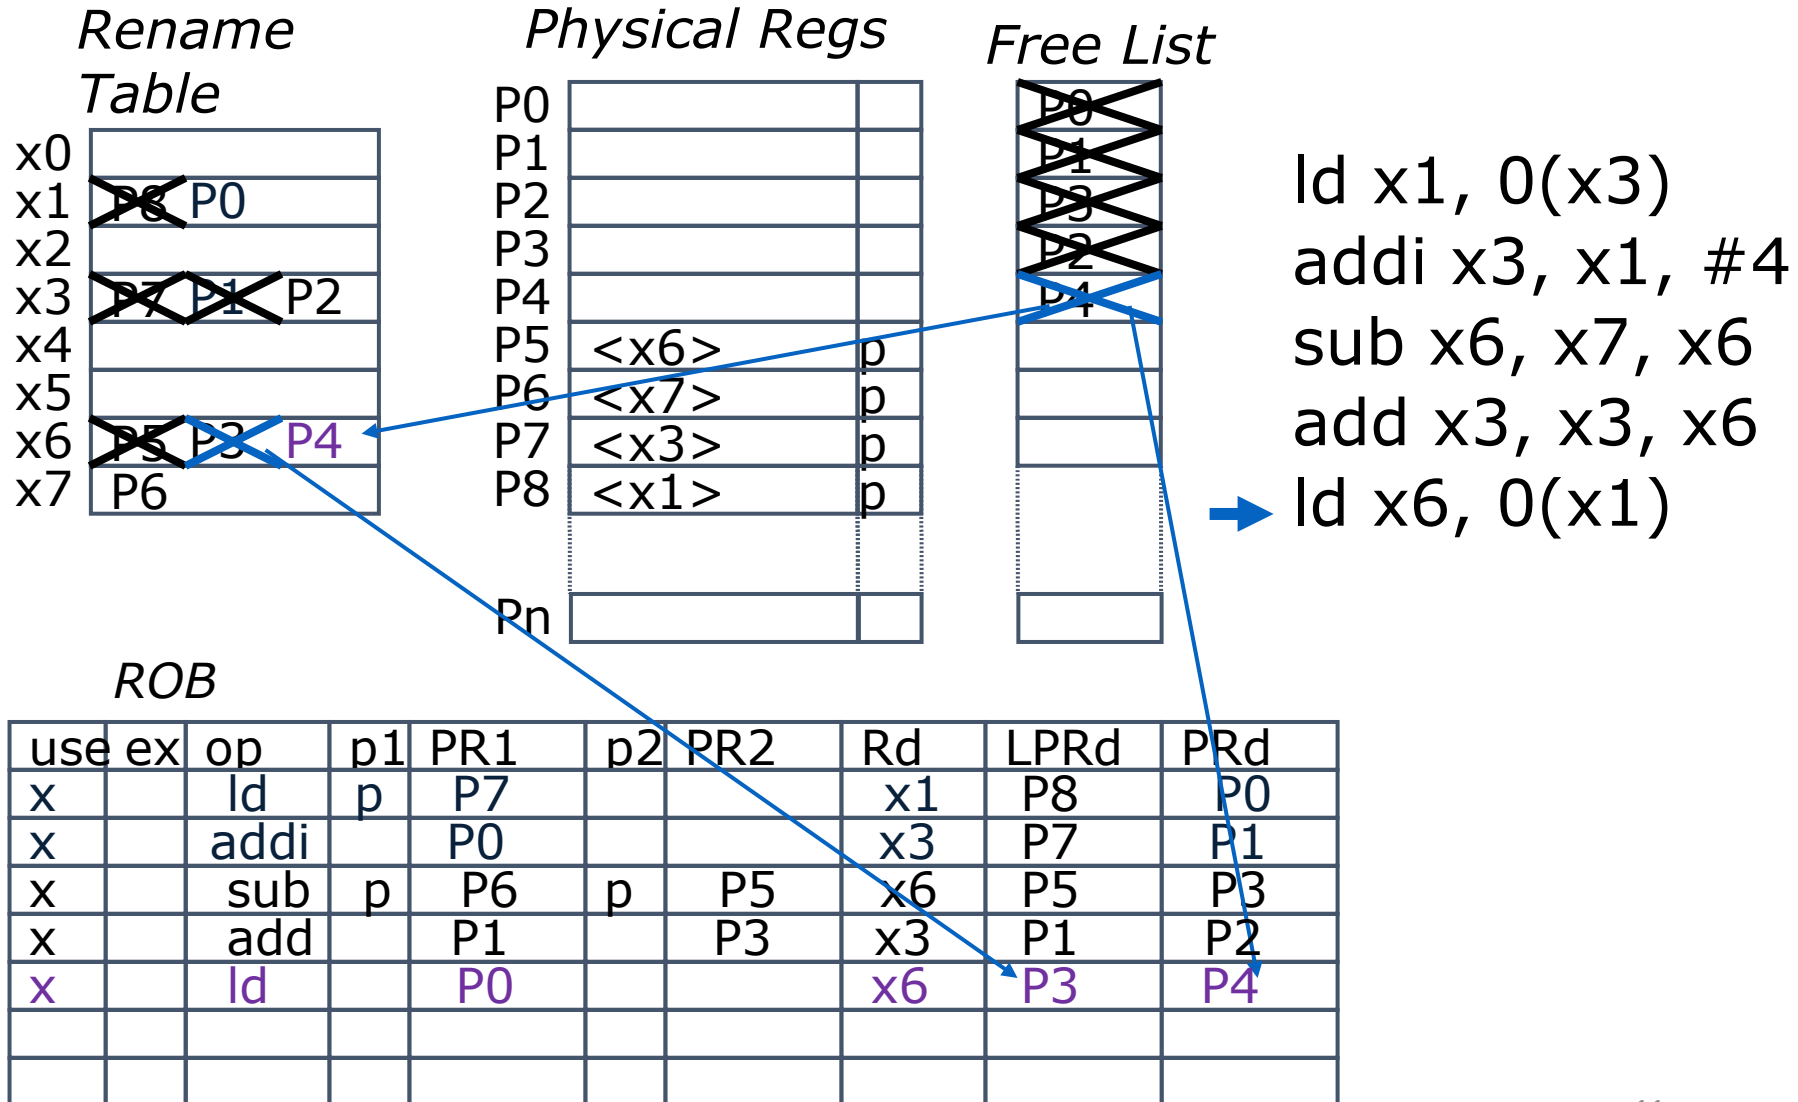
# Physical Register Management



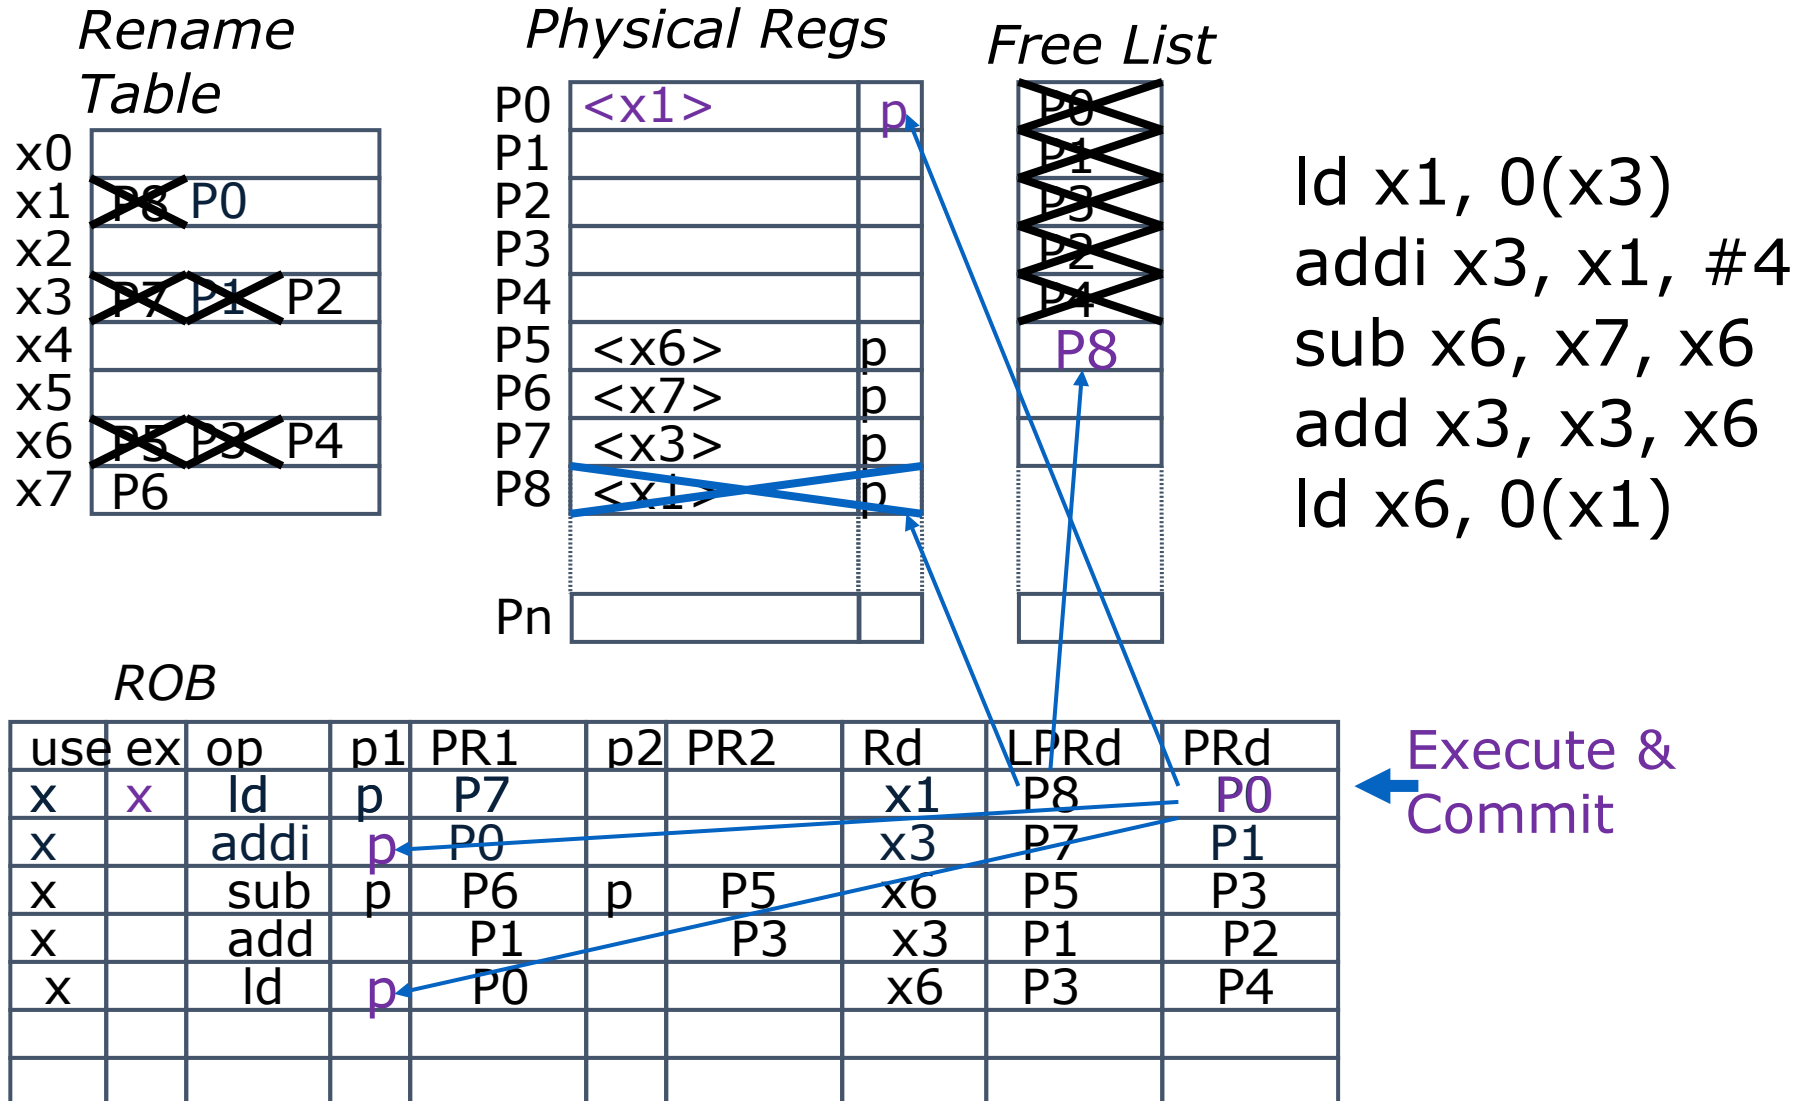
# Physical Register Management



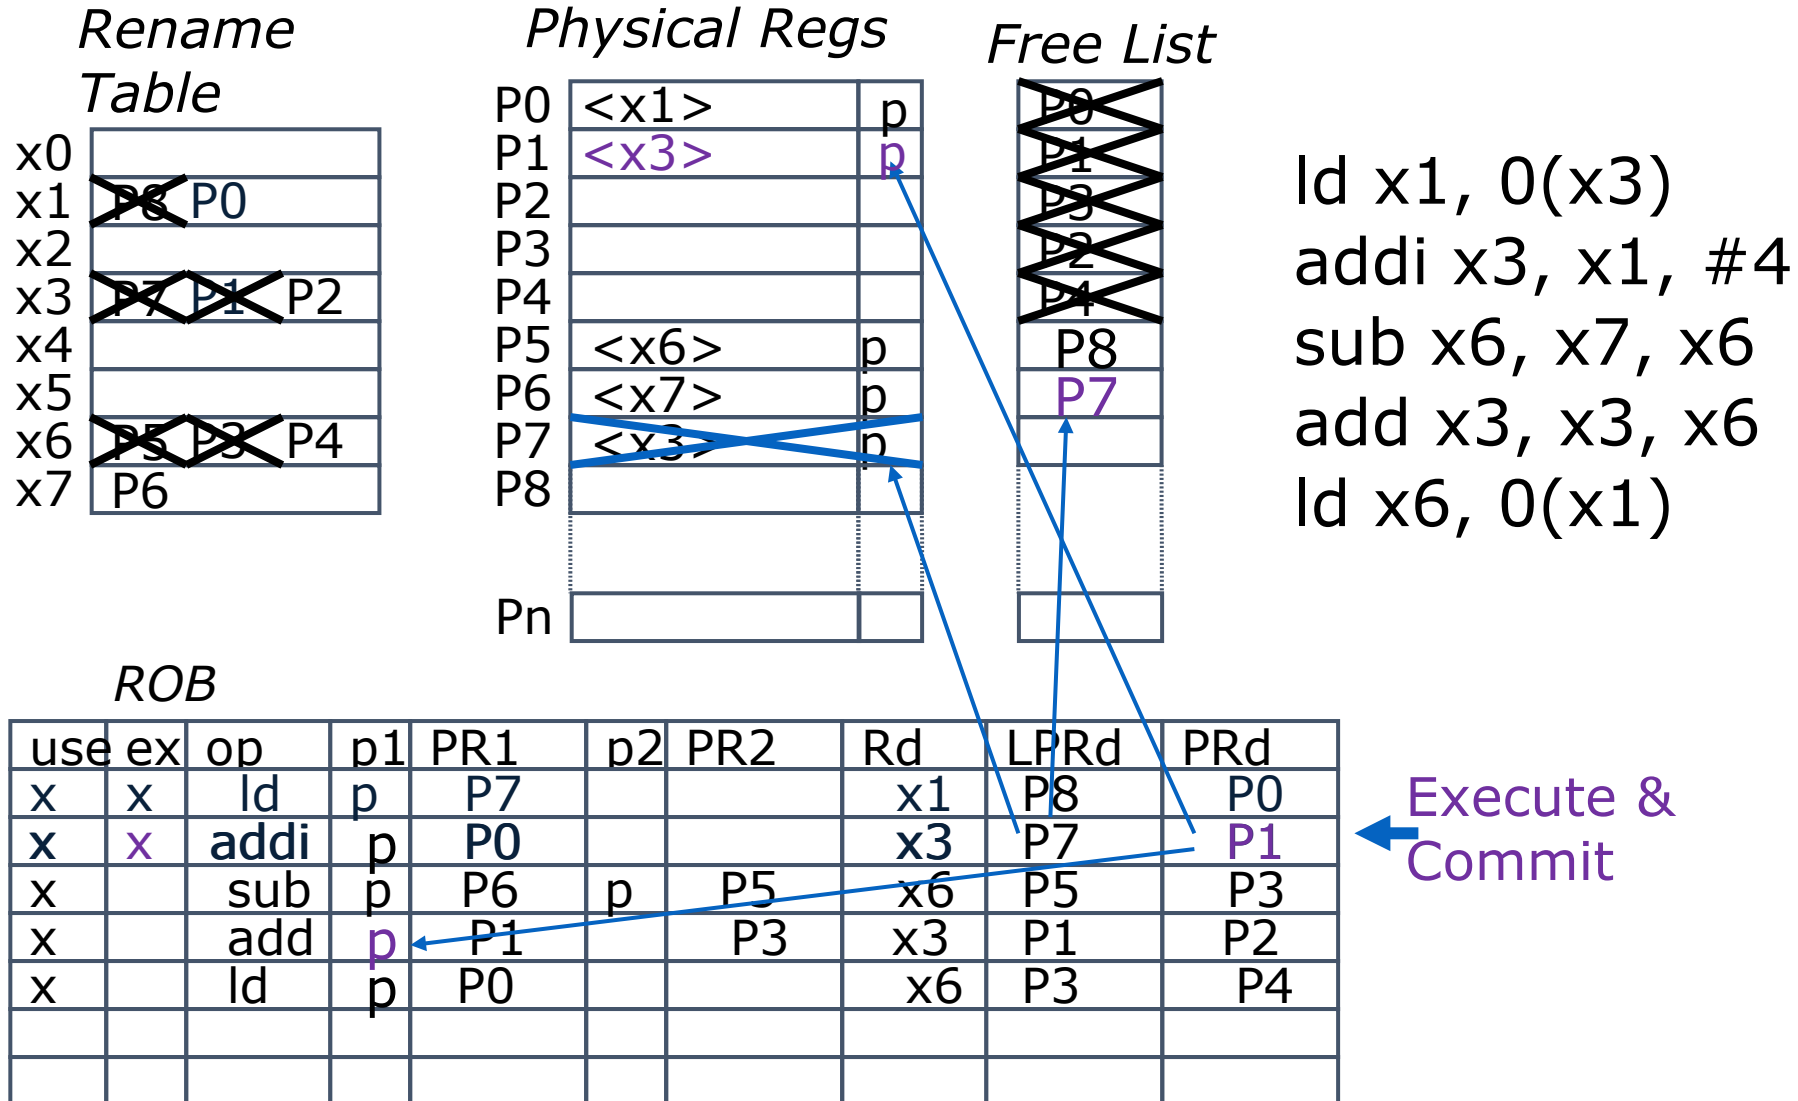
# Physical Register Management



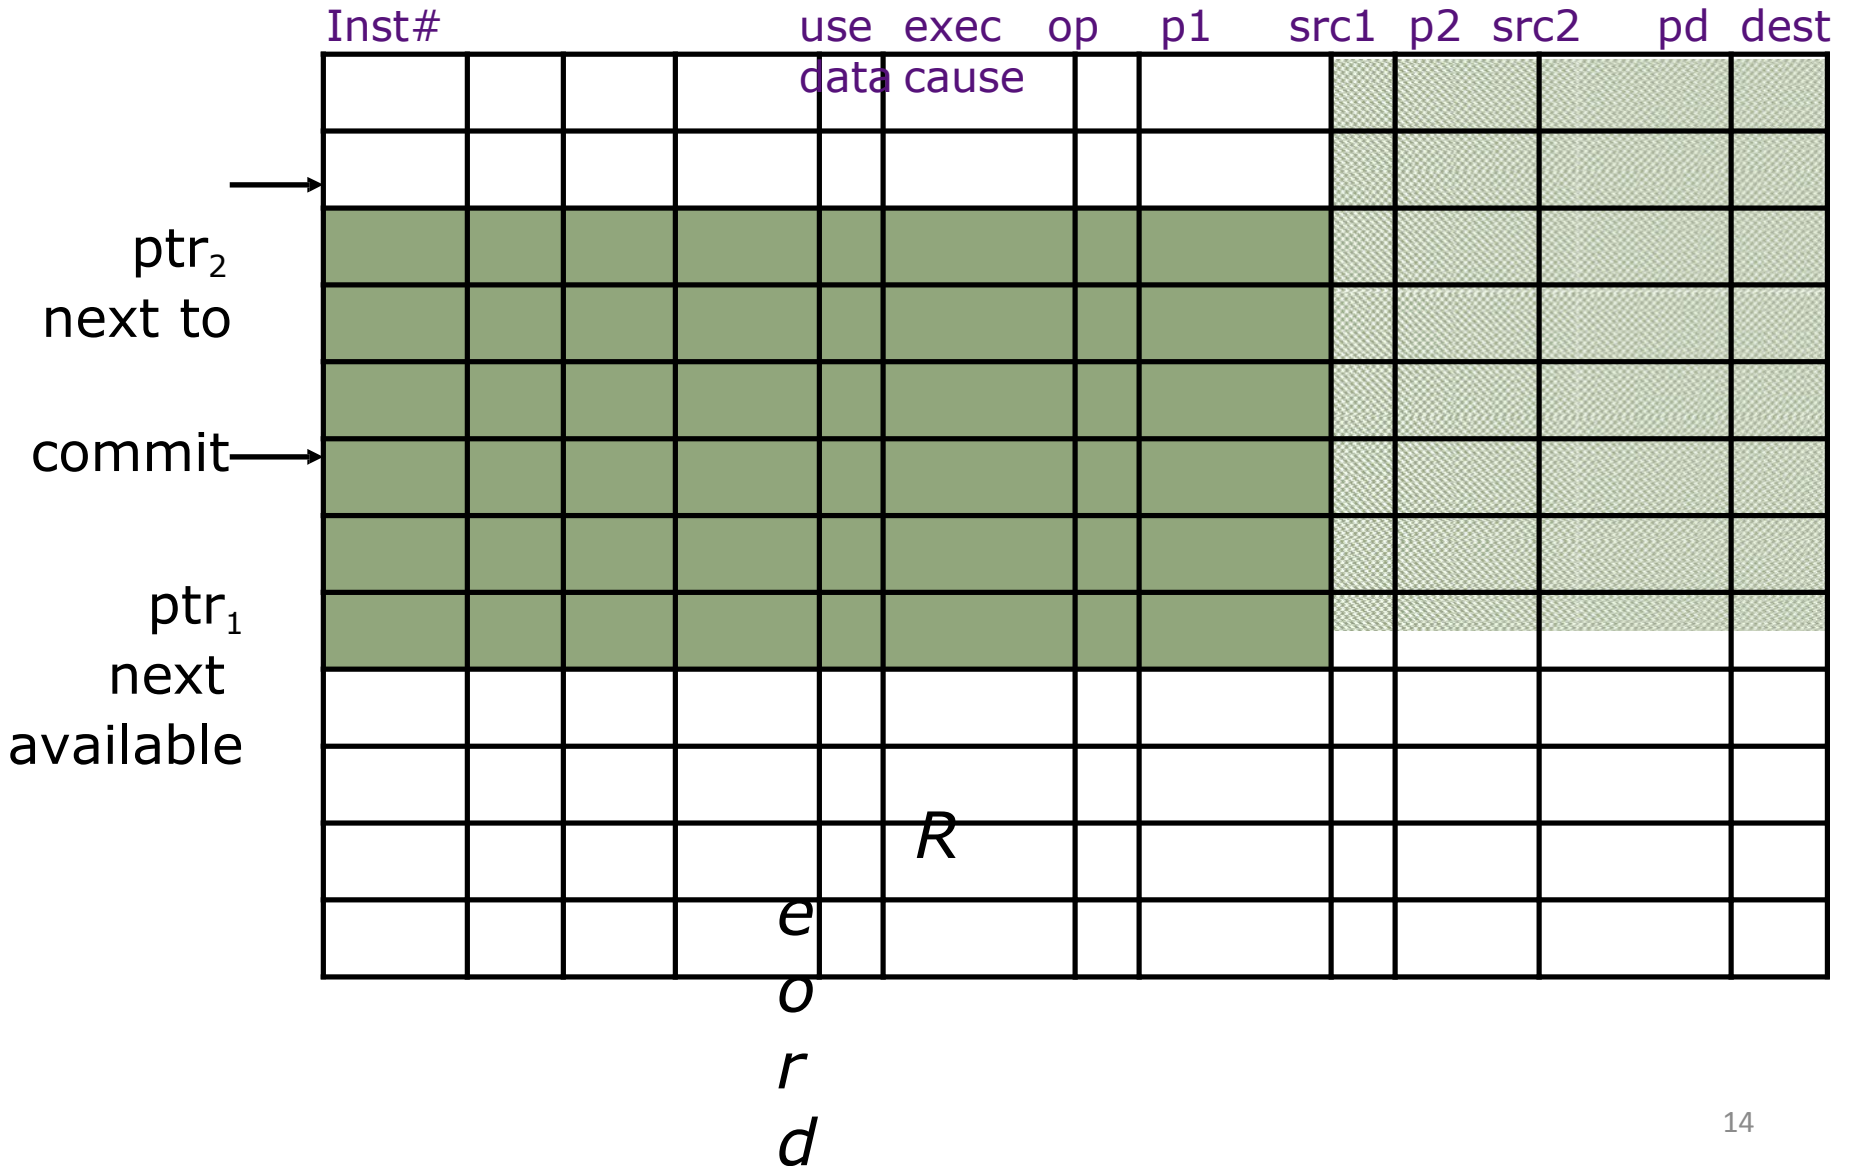
# Physical Register Management



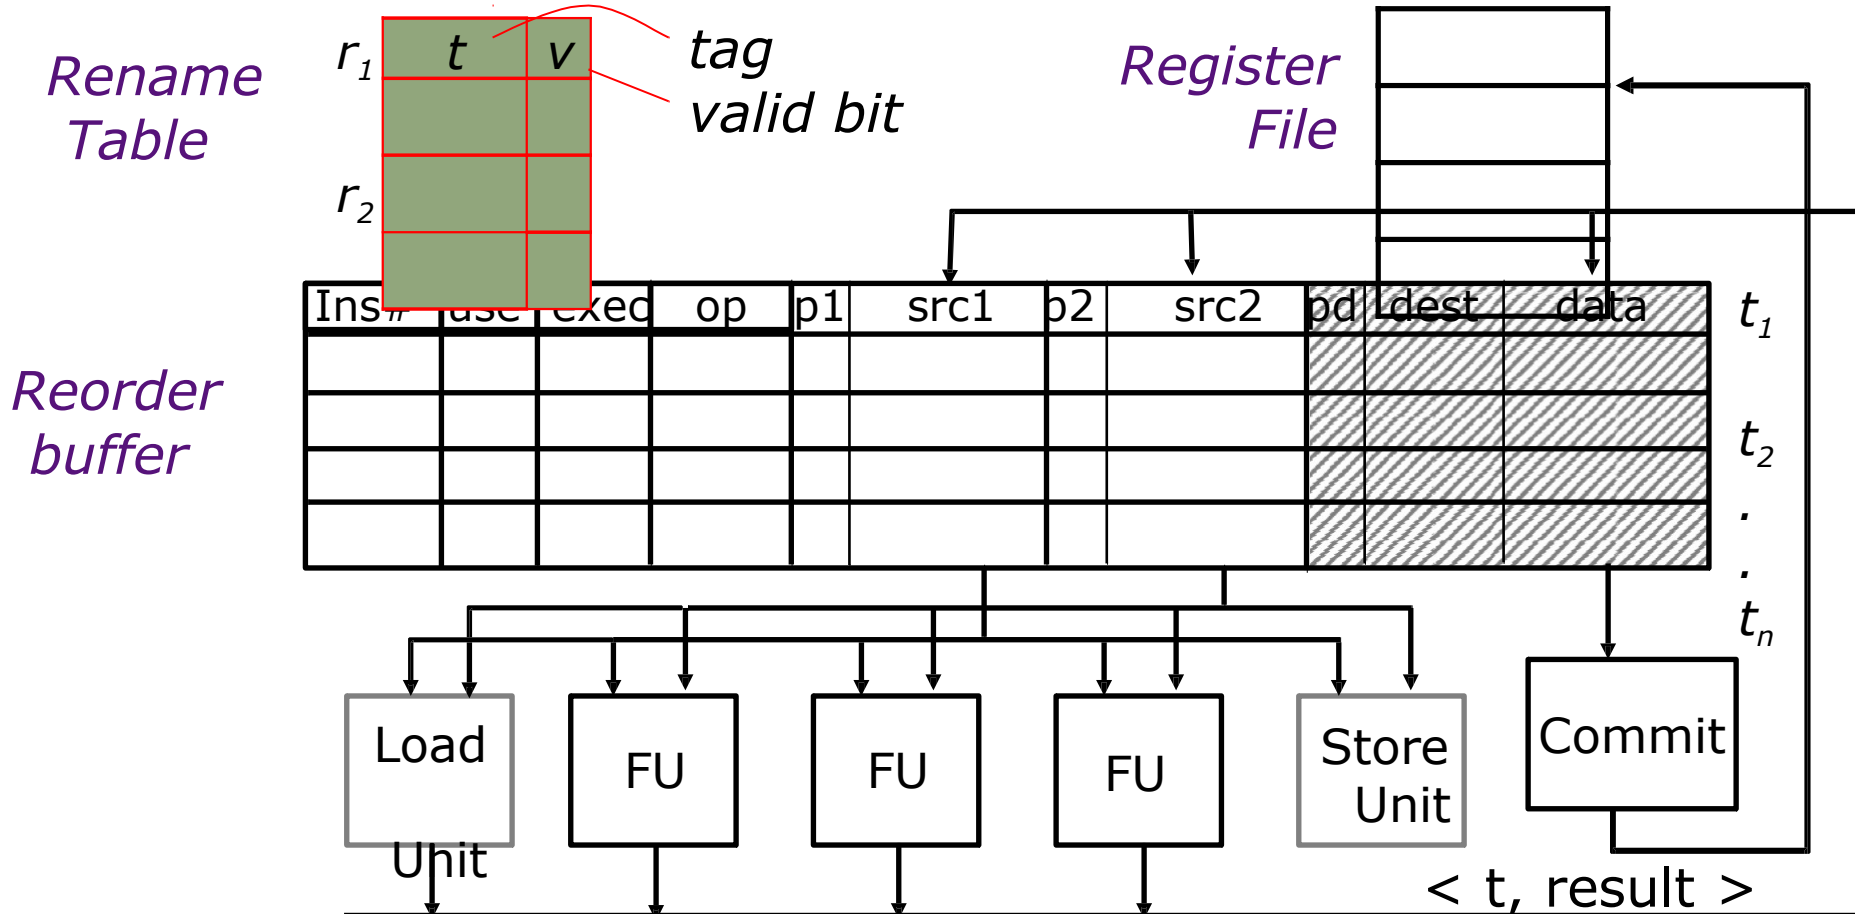
# Physical Register Management



# Extensions for Precise Exceptions



# Renaming Table




Renaming table is a cache to speed up register name look up. It needs to be cleared after each exception taken.

When else  
are valid bits cleared?  
*Control transfers*

# Branch Prediction Bits

- Assume 2 BP bits per instruction
- Change the prediction after two consecutive mistakes!

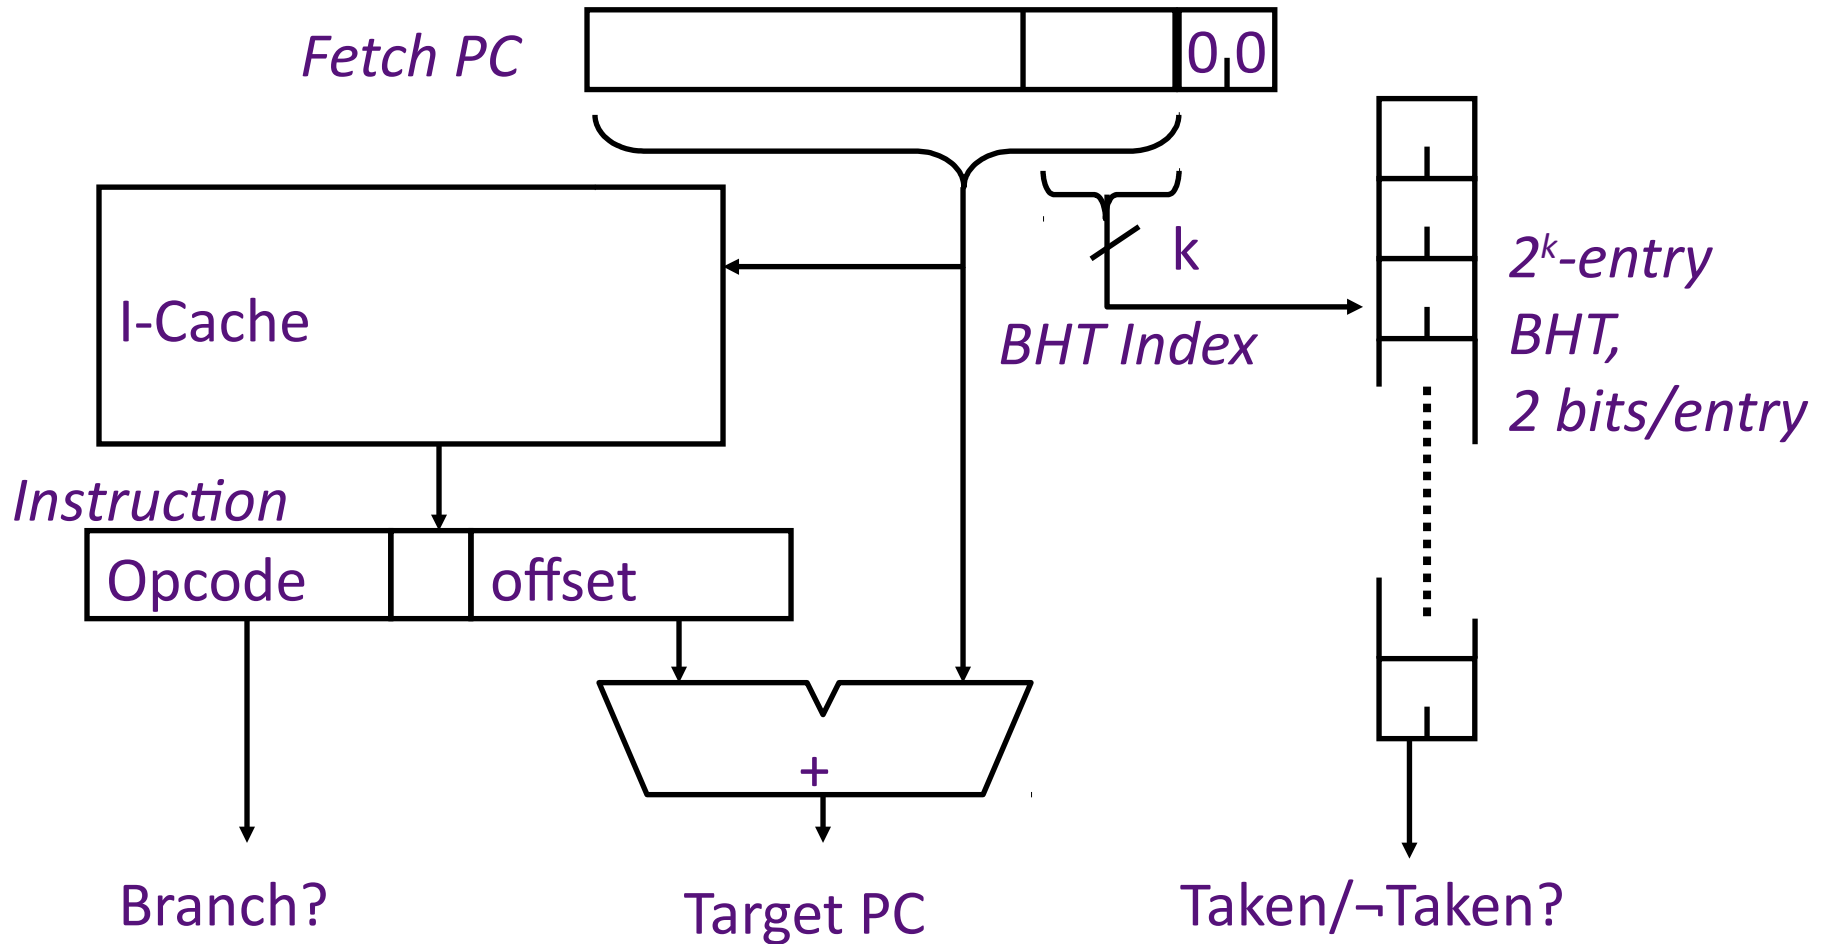


*BP state:*

*(predict take/¬take) x (last prediction right/wrong)*



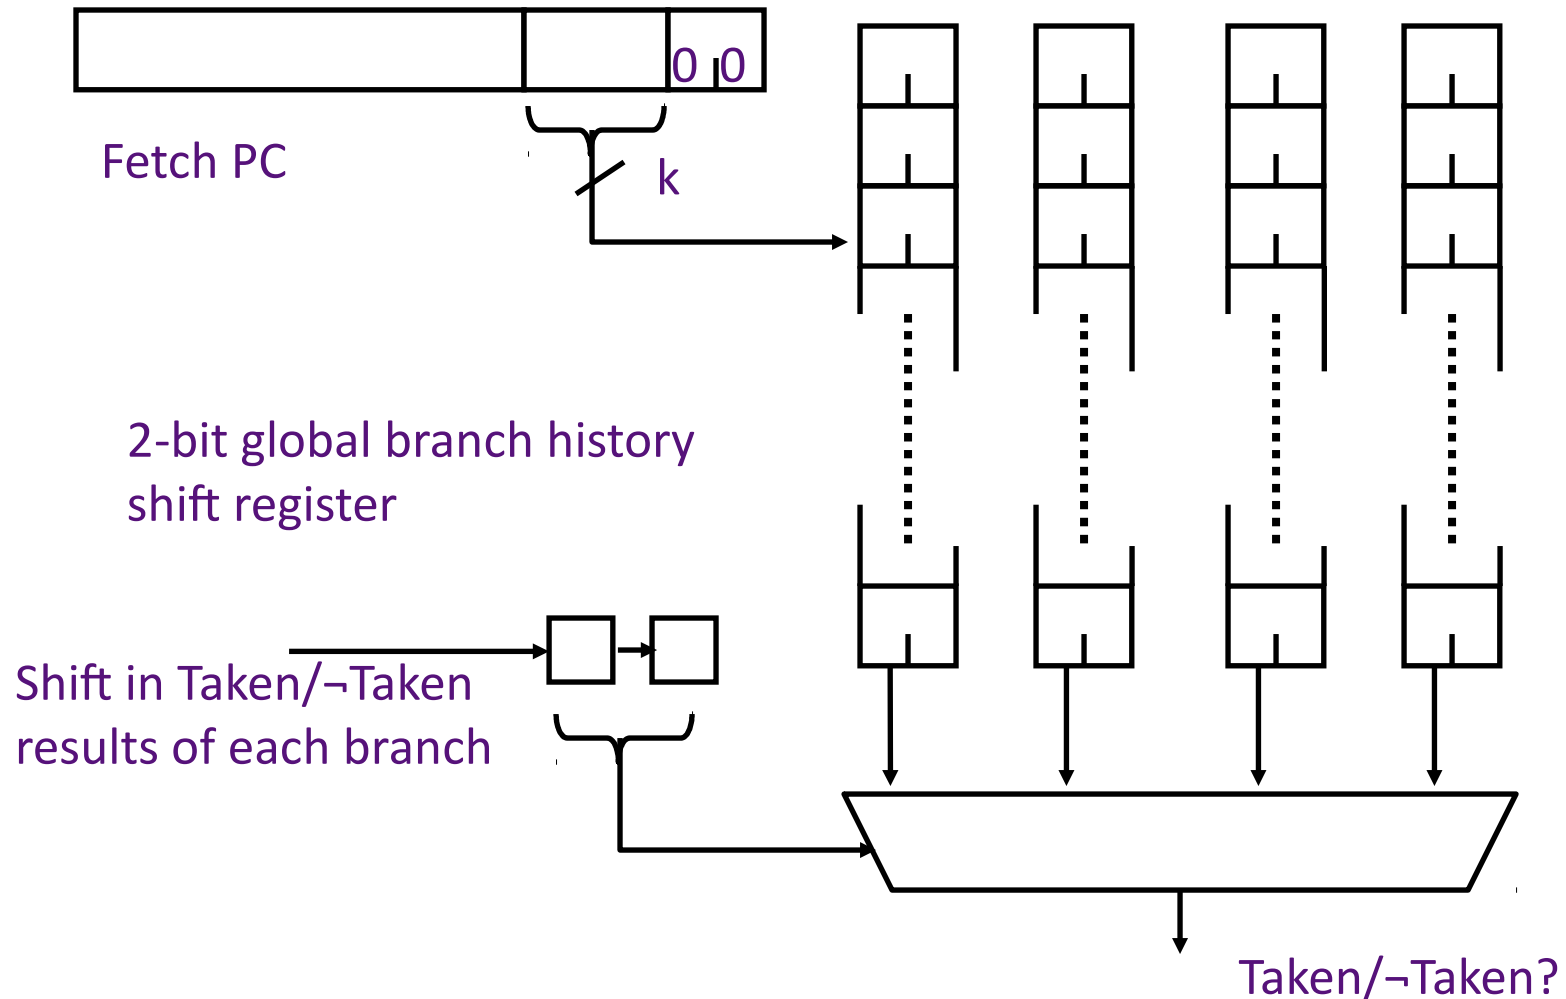
# Branch History Table (BHT)




4K-entry BHT, 2 bits/entry, ~80-90% correct predictions

# Two-Level Branch Predictor

*Pentium Pro uses the result from the last two branches to select one of the four sets of BHT bits (~95% correct)*



# Branch Target Buffer (BTB)



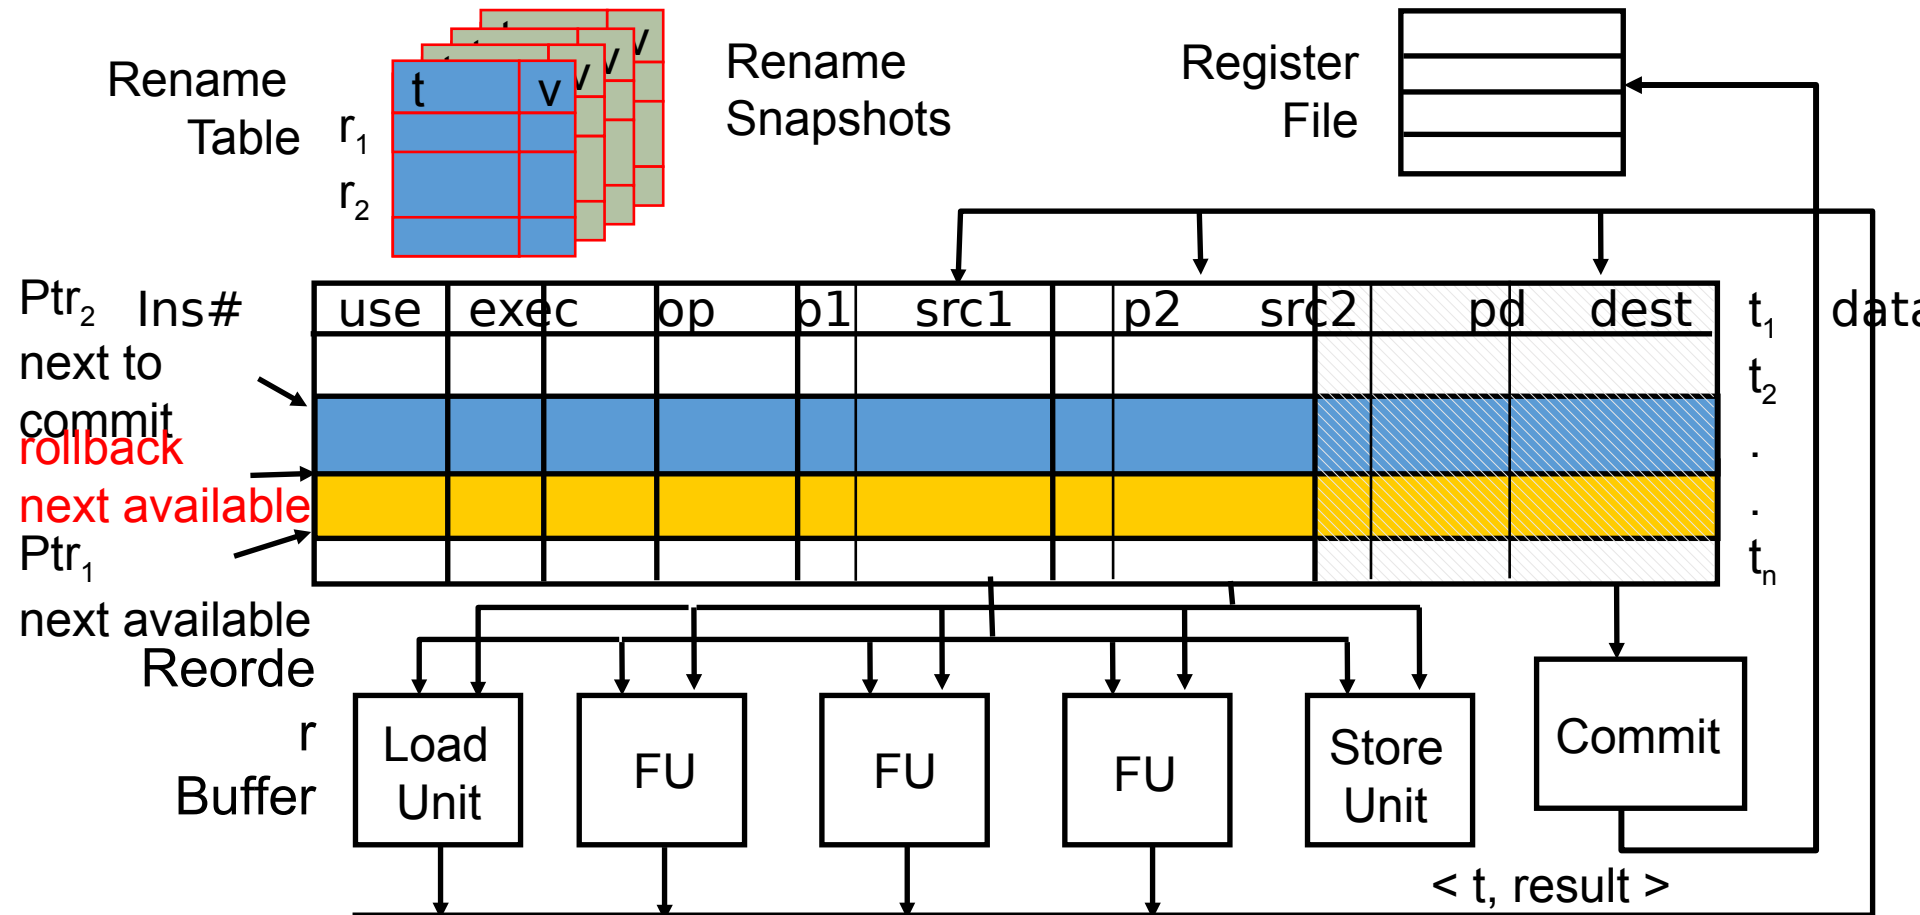
- Keep both the branch PC and target PC in the BTB
- PC+4 is fetched if match fails
- Only *taken* branches and jumps held in BTB
- Next PC determined *before* branch fetched and decoded

# Rename Table Recovery

---

- Have to quickly recover rename table on branch mispredicts
- MIPS R10K only has four snapshots for each of four outstanding speculative branches
- Alpha 21264 has 80 snapshots, one per ROB instruction

# Mispredict Recovery



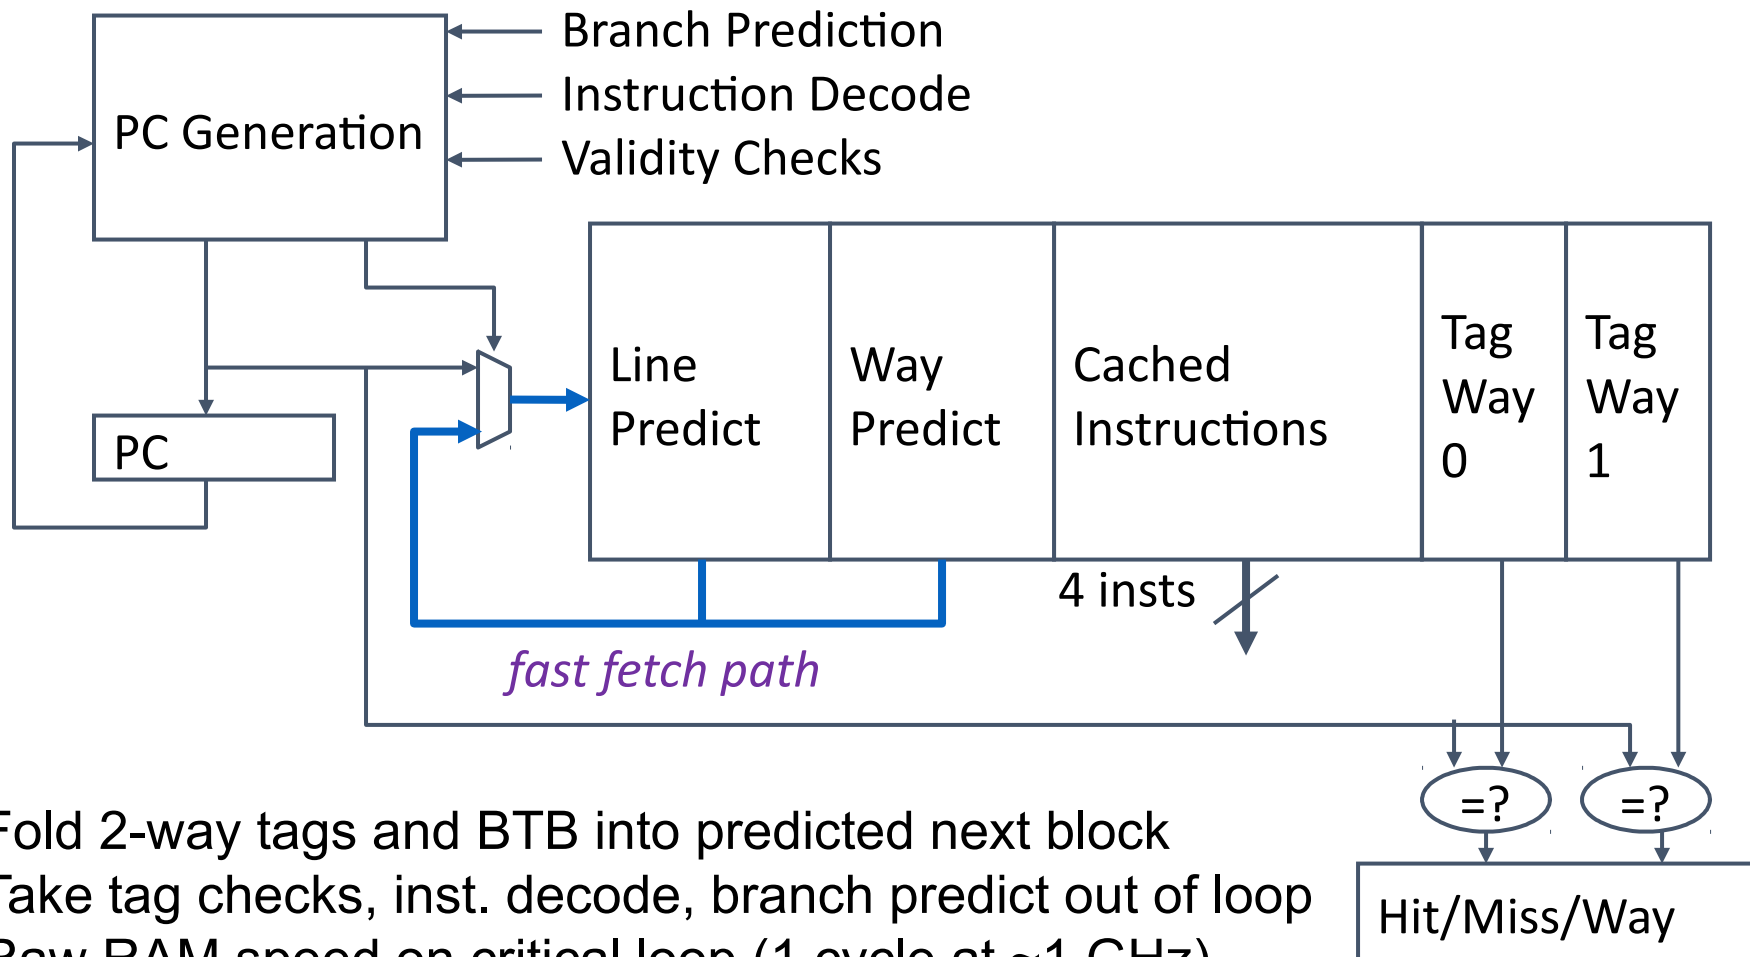
- Take snapshot of register rename table at each predicted branch, recover earlier snapshot if branch mispredicted

# Improving Instruction Fetch

---

- Performance of speculative out-of-order machines often limited by instruction fetch bandwidth
  - speculative execution can fetch 2-3x more instructions than are committed
  - mispredict penalties dominated by time to refill instruction window
  - taken branches are particularly troublesome

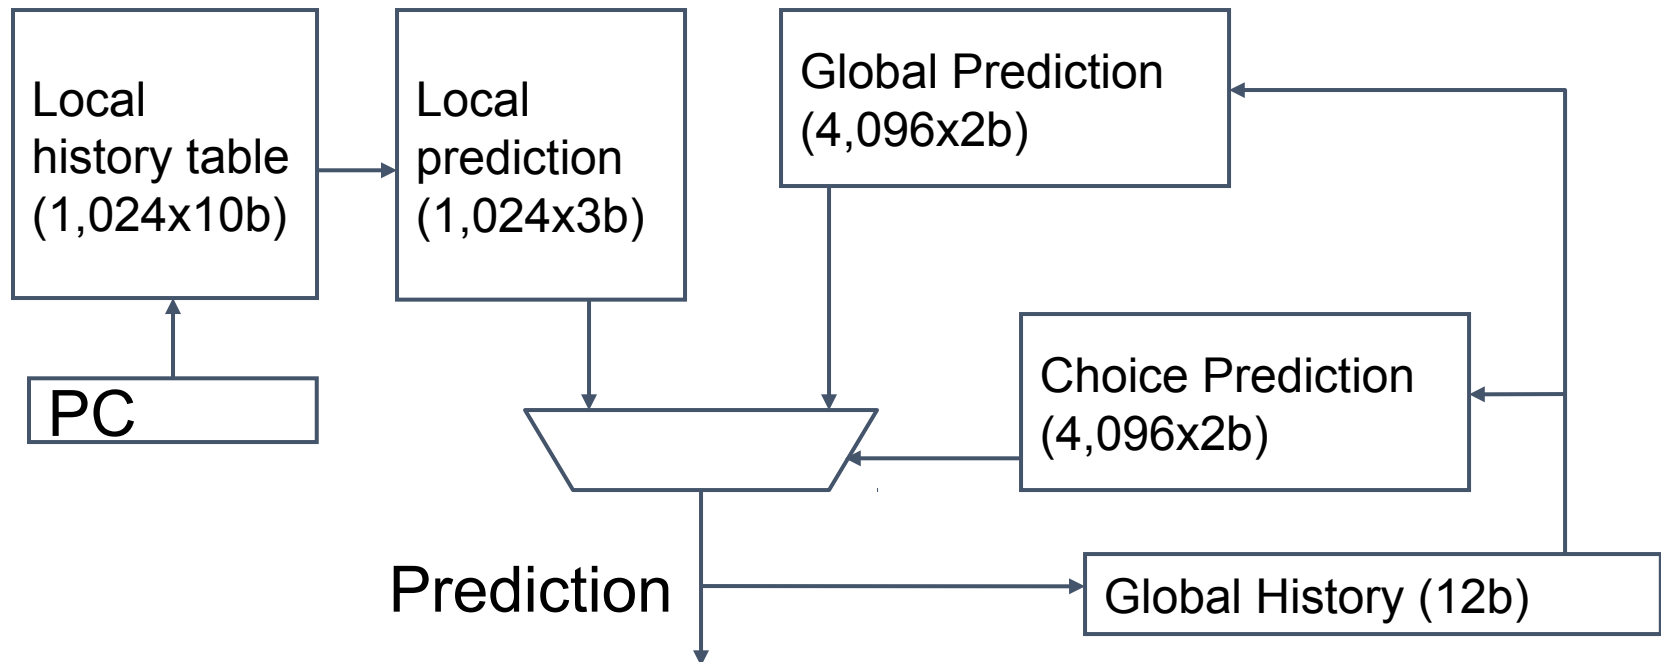
# Increasing Taken Branch Bandwidth (Alpha 21264 I-Cache)



- Fold 2-way tags and BTB into predicted next block
- Take tag checks, inst. decode, branch predict out of loop
- Raw RAM speed on critical loop (1 cycle at ~1 GHz)
- 2-bit hysteresis counter per block prevents overtraining

# Tournament Branch Predictor (Alpha 21264)

- Choice predictor learns whether best to use local or global branch history in predicting next branch
- Global history is speculatively updated but restored on mispredict
- Claim 90-100% success on range of applications



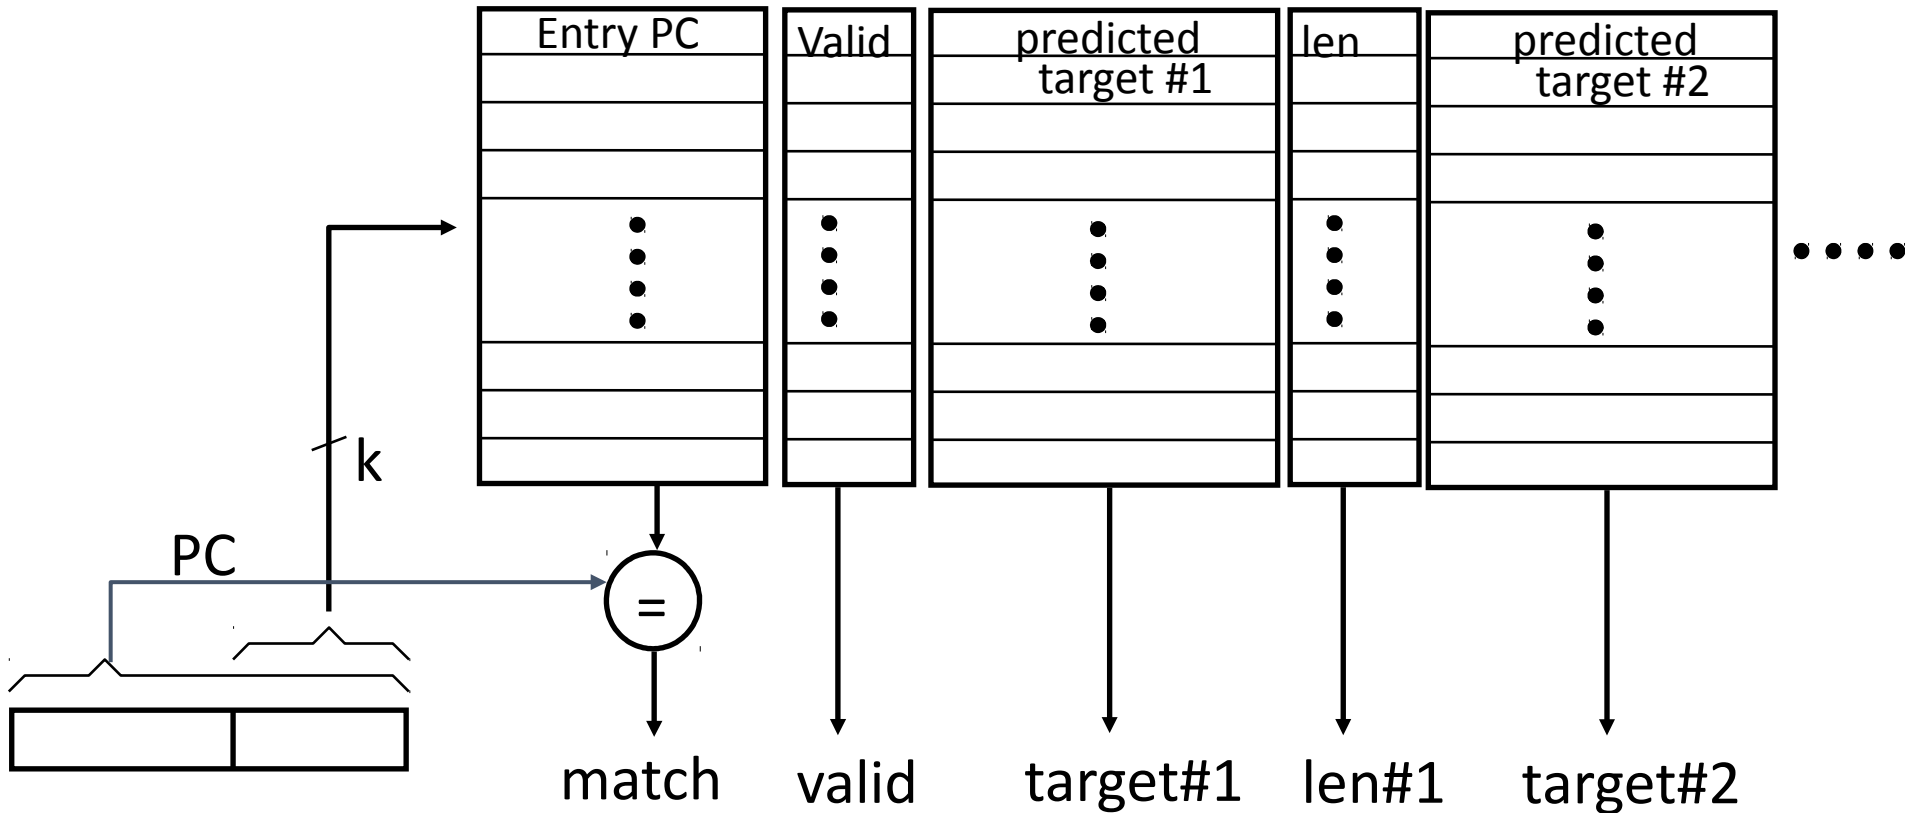


# Taken Branch Limit

---

- Integer codes have a taken branch every 6-9 instructions
- To avoid fetch bottleneck, must execute multiple taken branches per cycle when increasing performance
- This implies:
  - predicting multiple branches per cycle
  - fetching multiple non-contiguous blocks per cycle

# Branch Address Cache (Yeh, Marr, Patt)



Extend BTB to return multiple branch predictions per cycle

# Fetching Multiple Basic Blocks


---

- Requires either
  - multiported cache: expensive
  - interleaving: bank conflicts will occur
- Merging multiple blocks to feed to decoders adds latency, increasing mispredict penalty and reducing branch throughput

# Trace Cache

---

- Key Idea: Pack multiple non-contiguous basic blocks into one contiguous trace cache line



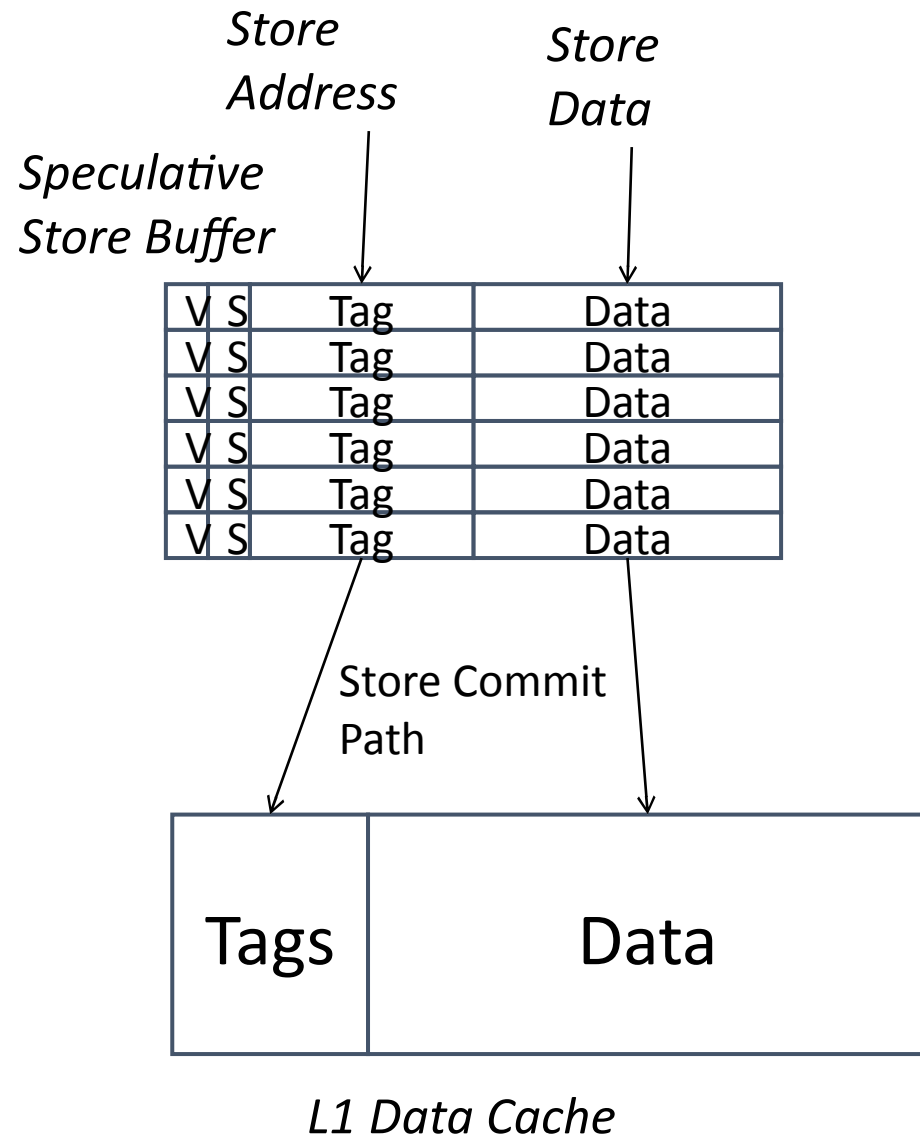
- Single fetch brings in multiple basic blocks
- Trace cache indexed by start address and next n branch predictions
- Used in Intel Pentium-4 processor to hold decoded uops

# Load-Store Queue Design

---

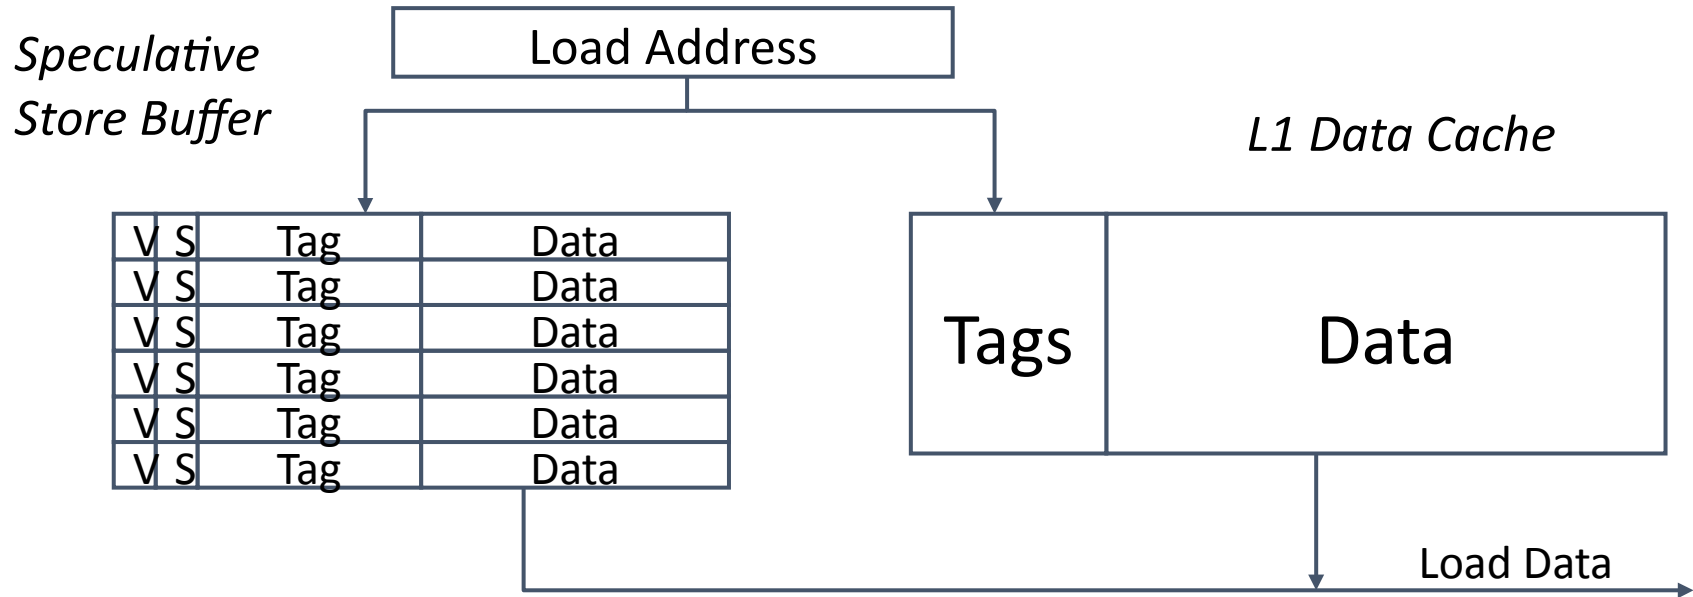
- After control hazards, data hazards through memory are probably next most important bottleneck to superscalar performance
- Modern superscalars use very sophisticated load-store reordering techniques to reduce effective memory latency by allowing loads to be speculatively issued

# Speculative Store Buffer



- Just like register updates, stores should not modify the memory until after the instruction is committed. A speculative store buffer is a structure introduced to hold speculative store data.
- During decode, store buffer slot allocated in program order
- Stores split into “store address” and “store data” micro-operations
- “Store address” execution writes tag
- “Store data” execution writes data
- Store commits when oldest instruction and both address and data available:
  - clear speculative bit and eventually move data to cache
- On store abort:
  - clear valid bit

# Load bypass from speculative store buffer



- If data in both store buffer and cache, which should we use?  
**Speculative store buffer**
- If same address in store buffer twice, which should we use?  
**Youngest store older than load**

# Memory Dependencies

---

```
sd x1, (x2)
ld x3, (x4)
```

- When can we execute the load?



# In-Order Memory Queue

---

- Execute all loads and stores in program order
  - => Load and store cannot leave ROB for execution until all previous loads and stores have completed execution
- Can still execute loads and stores speculatively, and out-of-order with respect to other instructions
- Need a structure to handle memory ordering...

# Conservative O-o-O Load Execution

---

```
sd x1, (x2)
ld x3, (x4)
```

- Can execute load before store, if addresses known and **x4 != x2**
- Each load address compared with addresses of all previous uncommitted stores
  - can use partial conservative check i.e., bottom 12 bits of address, to save hardware
- Don't execute load if any previous store address not known
- (MIPS R10K, 16-entry address queue)

# Address Speculation

---

```
sd x1, (x2)
ld x3, (x4)
```

- Guess that **x4**  $\neq$  **x2**
- Execute load before store address known
- Need to hold all completed but uncommitted load/store addresses in program order
- If subsequently find **x4==x2**, squash load and all following instructions
- => Large penalty for inaccurate address speculation

# Memory Dependence Prediction (Alpha 21264)

---

```
sd x1, (x2)
ld x3, (x4)
```

- Guess that  $x4 \neq x2$  and execute load before store
- If later find  $x4 == x2$ , squash load and all following instructions, but mark load instruction as store-wait
- Subsequent executions of the same load instruction will wait for all previous stores to complete
- Periodically clear store-wait bits

# Open-source Berkeley RISC-V Processors

---

- Sodor Collection
  - RV32I - Entry, educational, not synthesizable
- Rocket-chip SoC generator
- Z-scale
  - RV32IM - micro-controller
- Rocket
  - RV64G - in-order, single-issue application core
- BOOM
  - RV64G - out-of-order, superscalar application core

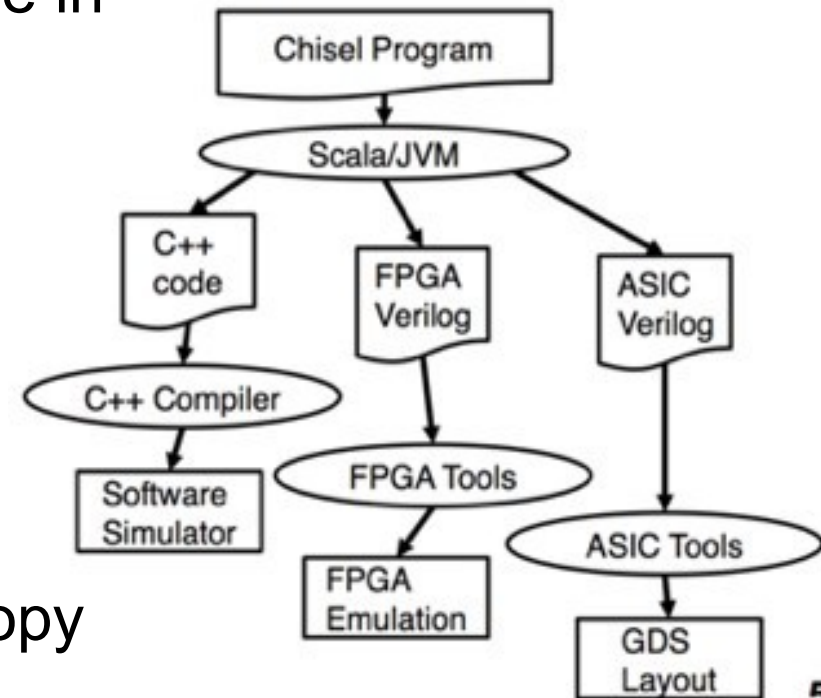
# What is BOOM?

---

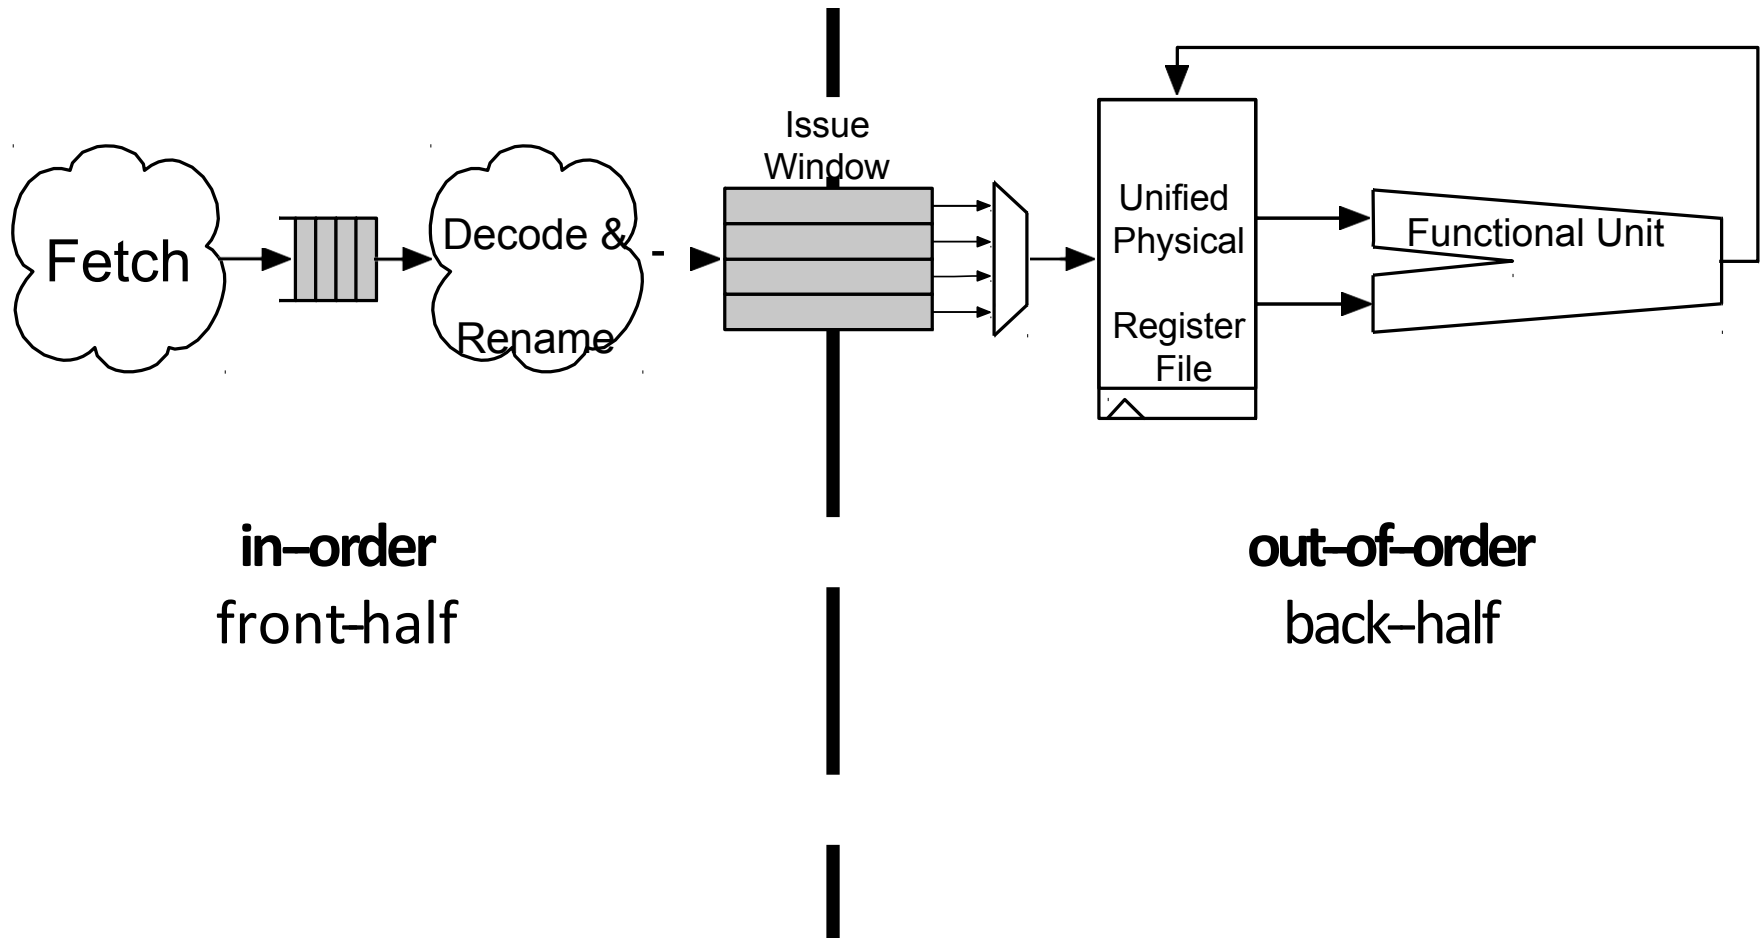
- superscalar, out-of-order processor written in Berkeley's Chisel RTL
- It is synthesizable
- It is parameterizable
- We hope to use it as a platform for architecture research
- It is open-source!
  - BOOM is a work-in-progress

# Chisel

- Hardware Construction Language embedded in Scala
- **not** a high-level synthesis language  
hardware module is a data structure in Scala
- Full power of Scala for writing generators
  - object-oriented programming
  - factory objects, traits, overloading
  - functional programming
  - high-order funs, anonymous funcs, currying
- generated C++ simulator is 1:1 copy of Verilog designs
- version 3.0 coming soon!
  - uses an IR called FIRRTL

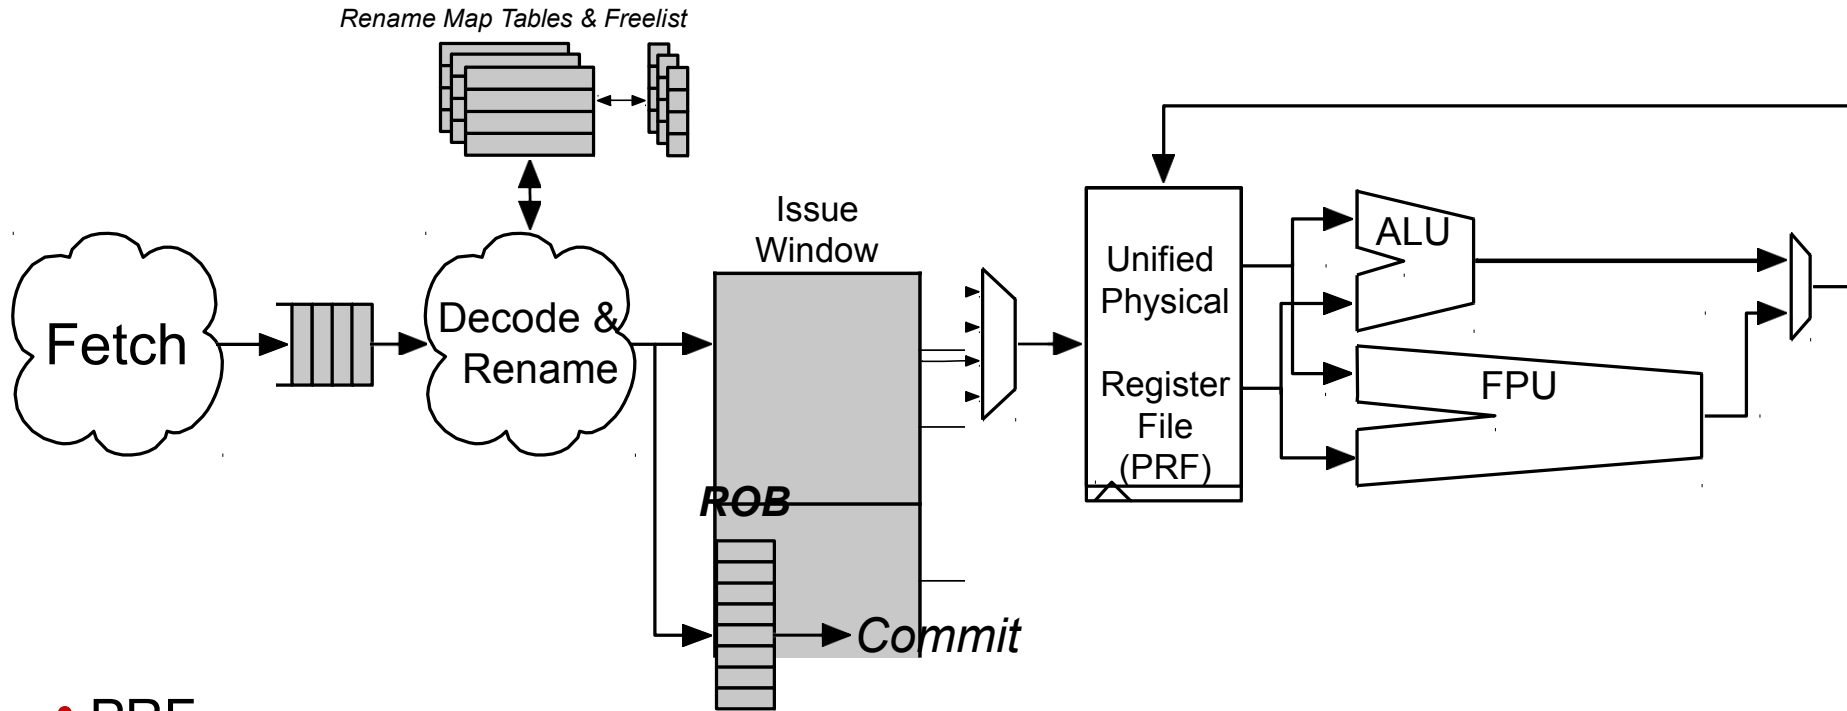


# BOOM Pipeline





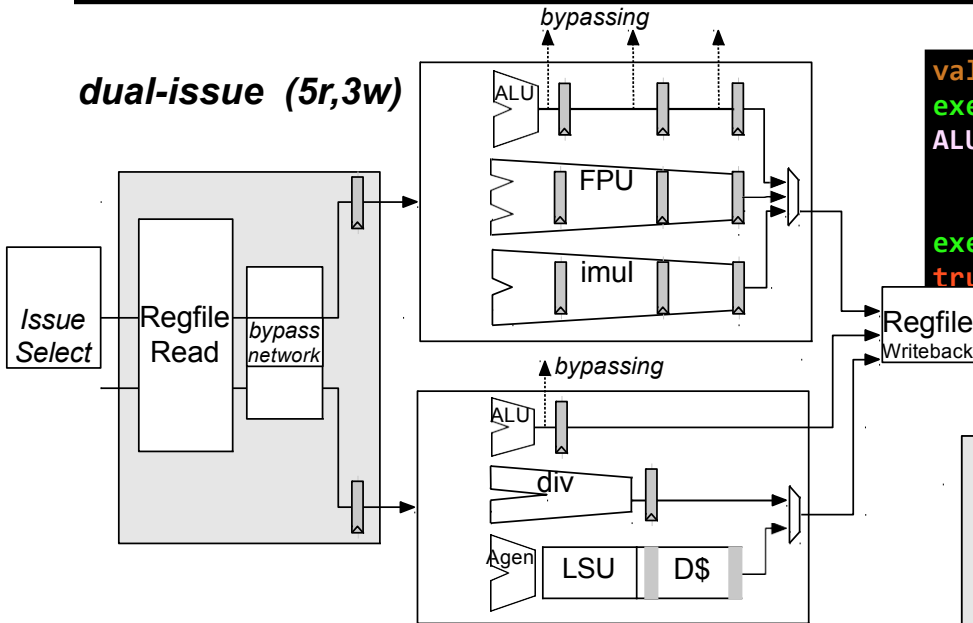
# BOOM Pipeline



- PRF
  - explicit renaming
  - holds speculative and committed data
  - holds both x-regs, f-regs
- split ROB/issue window design
- Unified Issue Window
  - holds all instructions

# Parameterized Superscalar

*dual-issue (5r,3w)*

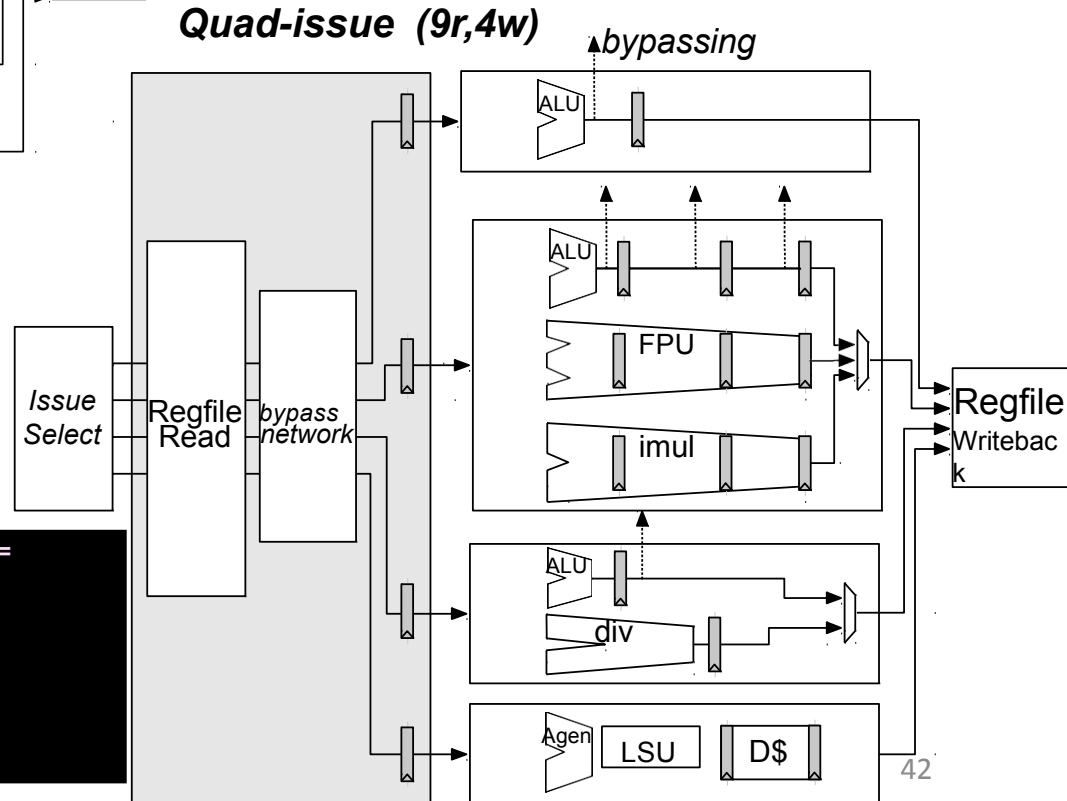


```
val exe_units = ArrayBuffer[ExecutionUnit]()
exe_units += Module(new ALUExeUnit(is_branch_unit = true
                                   , has_fpu = true
                                   , has_mul = true
                                   , has_div = true
                                   , fp_mem_support = true))
exe_units += Module(new ALUMemExeUnit(is_branch_unit = true
                                       , has_fpu = true
                                       , has_mul = true
                                       , has_div = true
                                       , fp_mem_support = true))

```

OR

*Quad-issue (9r,4w)*



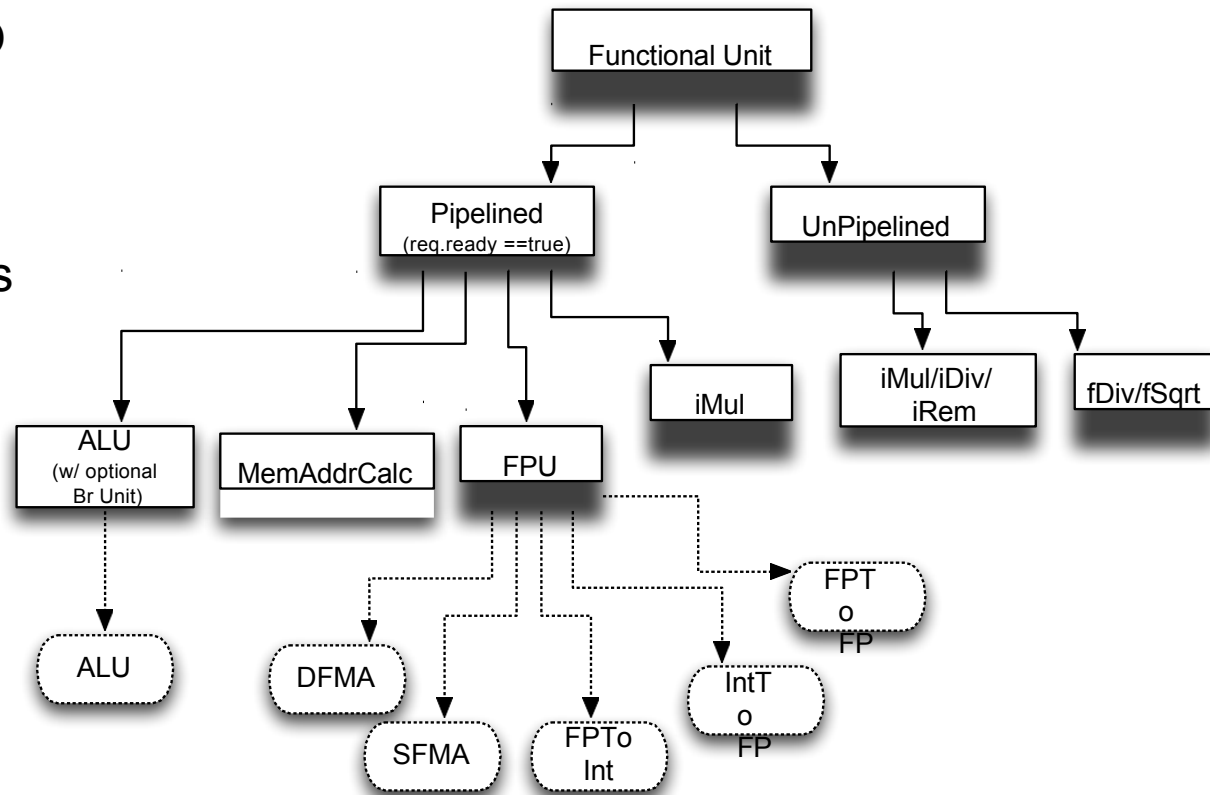
```
exe_units += Module(new ALUExeUnit(is_branch_unit = true
                                   , has_fpu = true
                                   , has_mul = true
                                   , has_div = true
                                   , fp_mem_support = true))
exe_units += Module(new ALUExeUnit(is_branch_unit = true
                                   , has_fpu = true
                                   , has_mul = true
                                   , has_div = true
                                   , fp_mem_support = true))
exe_units += Module(new ALUExeUnit(is_branch_unit = true
                                   , has_fpu = true
                                   , has_mul = true
                                   , has_div = true
                                   , fp_mem_support = true))
exe_units += Module(new ALUExeUnit(is_branch_unit = true
                                   , has_fpu = true
                                   , has_mul = true
                                   , has_div = true
                                   , fp_mem_support = true))
exe_units += Module(new MemExeUnit())

```

# Open-source Berkeley RISC-V

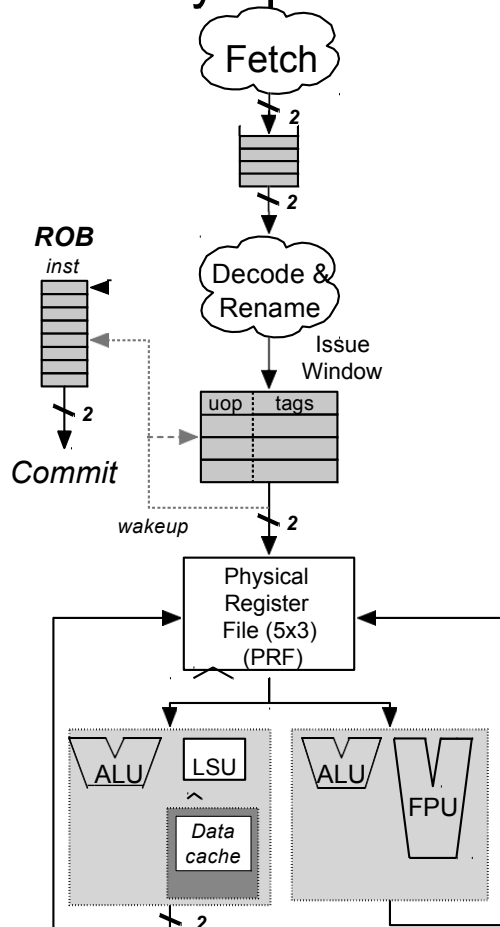
## Processors

- Abstract FunctionalUnit
  - describes common IO
- Pipelined/Unpipelined
  - handles storing uop metadata, branch resolution, branch kills
- Concrete Subclasses
  - instantiates the actual expert-written FU
  - no modifications required to get FU working with speculative OoO
  - allows easy “stealing” of external code



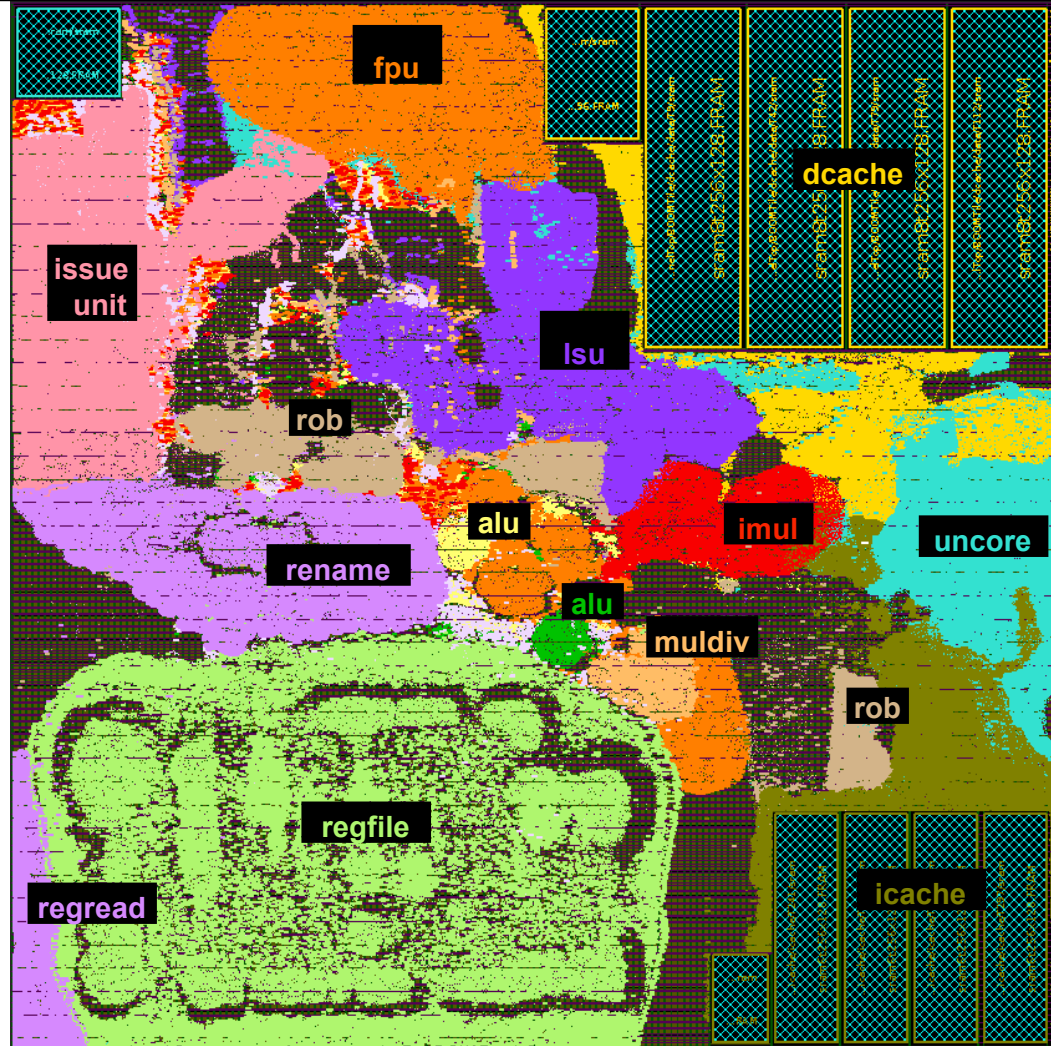
# Synthesizable

- Targeting ASIC
- Runs on FPGA
  - Zynq zc706



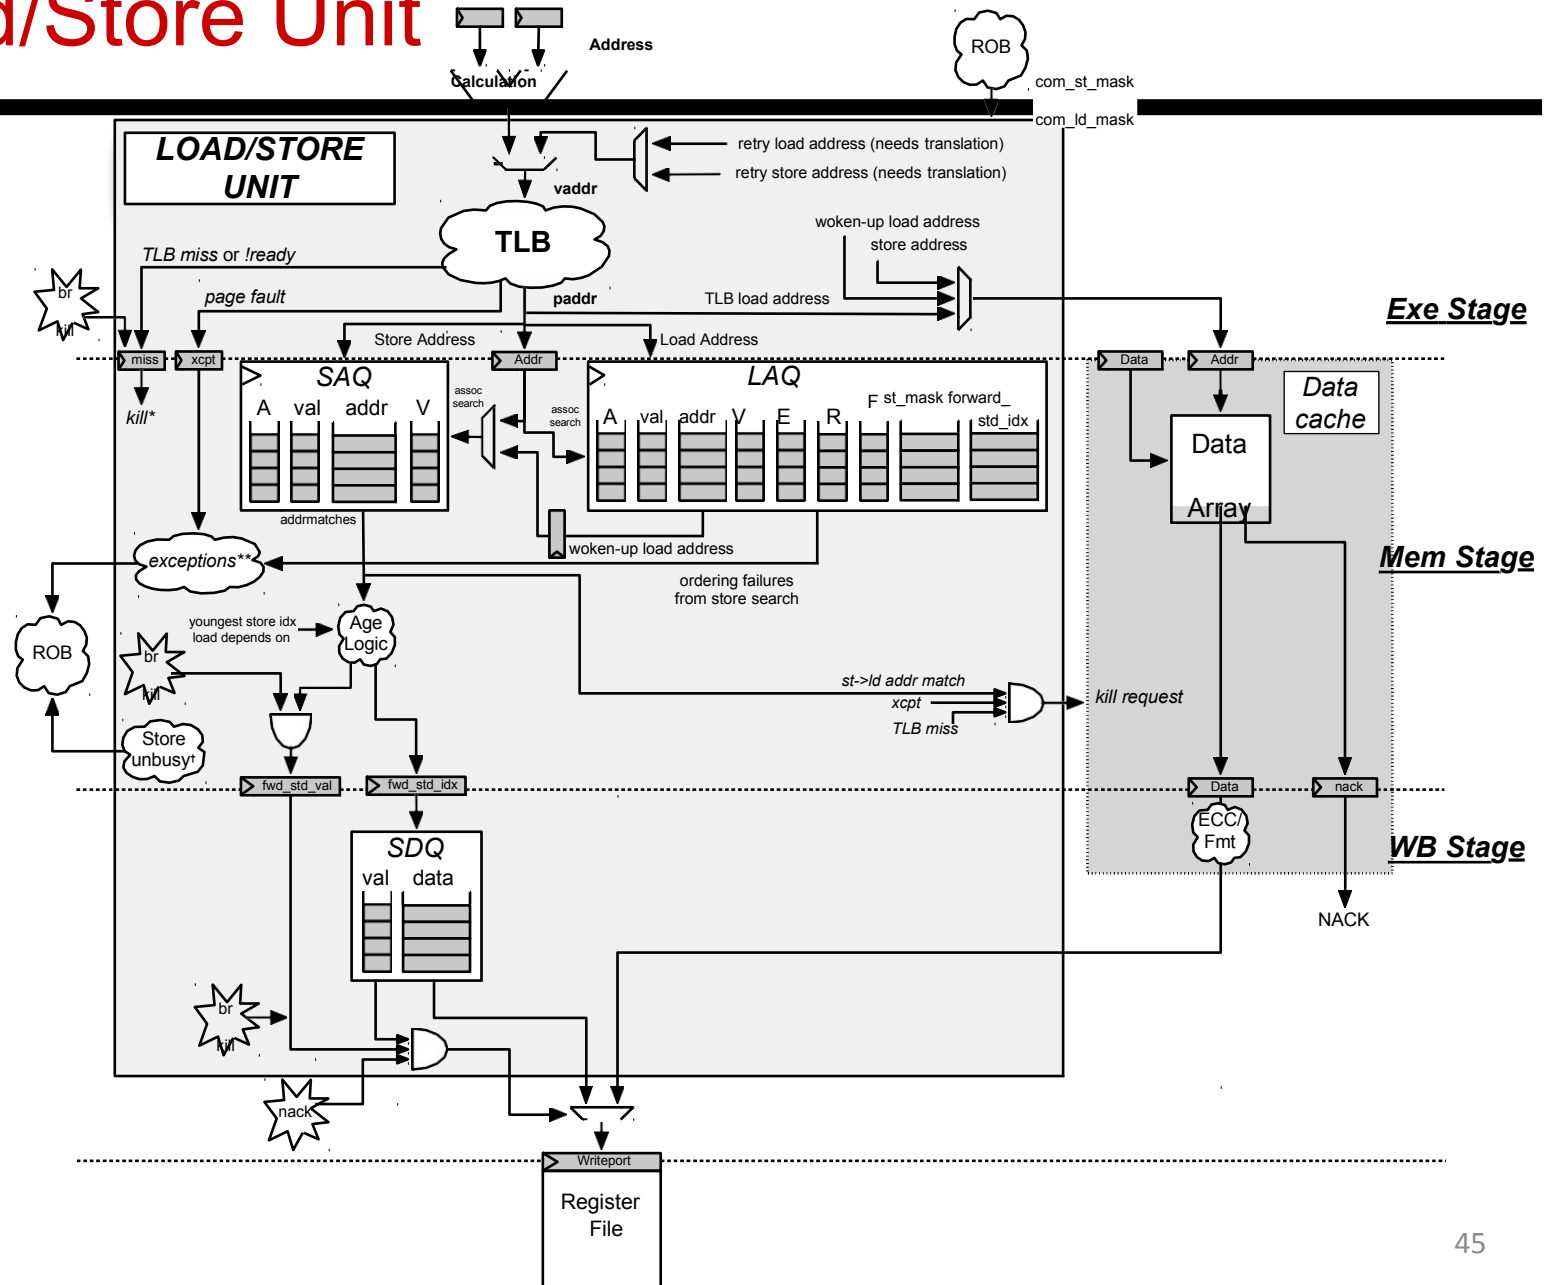
preliminary results

TSMC 45nm floorplan



2-wide BOOM, 16kB L1 caches  
**1.2 mm<sup>2</sup>**

# Load/Store Unit




# Comparison against ARM

Category	ARM Cortex-A9	RISC-V BOOM-2w
ISA	32-bit ARM v7	64-bit RISC-V v2 (RV64G)
Architecture	2 wide, 3+1 issue Out-of-Order 8-stage	2 wide, 3 issue Out-of-Order 6-stage
Performance	<b>3.59</b> CoreMarks/MHz	<b>3.91</b> CoreMarks/MHz
Process	TSMC 40GPLUS	TSMC 40GPLUS
Area with 32K caches	~2.5 mm <sup>2</sup>	~1.00 mm <sup>2</sup>
Area efficiency	<b>1.4</b> CoreMarks/MHz/mm <sup>2</sup>	<b>3.9</b> CoreMarks/MHz/mm <sup>2</sup>
Frequency	1.4 GHz	1.5 GHz
Power	<b>0.5-1.9 W</b> (2 cores + L2) @ TSMC 40nm, 0.8-2.0 GHz	<b>0.25 W</b> (1 core + L1) @ TSMC 45nm, 1 GHz

**+9%!**

note:  
not to scale



# Repositories

---

- BOOM
  - <https://github.com/ucb-bar/riscv-boom>
  - just the BOOM core code
- Rocket-chip
  - <https://github.com/ucb-bar/rocket-chip>
  - the rest of the SoC
  - chisel, rocket, uncore, junctions, riscv-tools, fpga-zynq, etc.
  - generates C++ emulator, verilog for FPGA
  - <http://riscv.org/tutorial-hpca2015/riscv-rocket-chip-tutorial-bootcamp-hpca2015.pdf> (for more information)

# BOOM Quick-start

```
$ export ROCKETCHIP_ADDONS="boom"
$ git clone https://github.com/ucb-bar/rocket-chip.git
$ cd rocket-chip
$ git checkout boom
$ git submodule update --init
$ cd riscv-tools
$ git submodule update --init --recursive riscv-tests
$ cd ../emulator; make run CONFIG=BOOMCPPConfig
```

- "BOOM-chip" is (currently) a branch of rocket-chip
- "make run" builds and runs riscv-tests suite
- there are many different CONFIGs available!
  - rocket-chip/src/main/scala/PrivateConfigs.scala
  - rocket-chip/boom/src/main/scala/configs.scala
  - slightly different configs for different targets (CPP, FPGA, VLSI)



# How do you verify and debug BOOM?

---

- parameterization makes this tough!
- tested with:
  - riscv-tests
    - (assembly functional tests + bare-metal micro-benchmarks)
  - CoreMark + riscv-pk
  - SPEC + Linux
  - riscv-torture

# riscv-torture

---

- now open-source!
  - <https://github.com/ucb-bar/riscv-torture>
- torture generates random tests to stress the core pipeline
- runs test on Spike and your processor
  - architectural register state dumped to memory on program termination
  - diff state
  - if error found, finds smallest version of test that exhibits an error

# riscv-torture quick-start

---

## Run Rocket-chip (BOOM-chip?)

```
$ git clone https://github.com/ucb-bar/rocket-chip.git
$ cd rocket-chip
$ git submodule update --init
$ cd riscv-tools
$ git submodule update --init --recursive riscv-tests
$ cd ../emulator; make run CONFIG=BOOMCPPConfig
```

## Run Torture

```
$ cd rocket-chip
$ git clone https://github.com/ucb-bar/riscv-torture.git
$ cd riscv-torture
$ git submodule update --init
$ vim Makefile # change RTL_CONFIG=BOOMCPPConfig
$ make igentest # test that torture works, gen a single test
$ make cnight# run C++ emulator overnight
```

# Commit Logging

---

- BOOM can generate commit logs
  - (priv level, PC, inst, wb raddr, wb data)
- Spike supports generating commit logs!
  - configure Spike with "--enable-commitlog" flag
  - outputs to stderr
- only semi-automated
  - valid differences from Spike and hw
  - e.g., spinning on LR+SC

**# modify the build.sh in riscv-tools**

```
build_project riscv-isa-sim --prefix=$RISCV --with-fesvr=$RISCV --enable-commitlog
```

# What documentation is available?

---

- A design document is in progress
  - <https://github.com/ccelio/riscv-boom-doc>
- Wiki
  - <https://github.com/ucb-bar/riscv-boom/wiki>

# Conclusion

- BOOM is RV64G, runs SPEC on Linux on an FPGA
- BOOM is ~10k loc and 4 person-years of work
- excellent platform for prototyping new ideas
- now open-source!

